

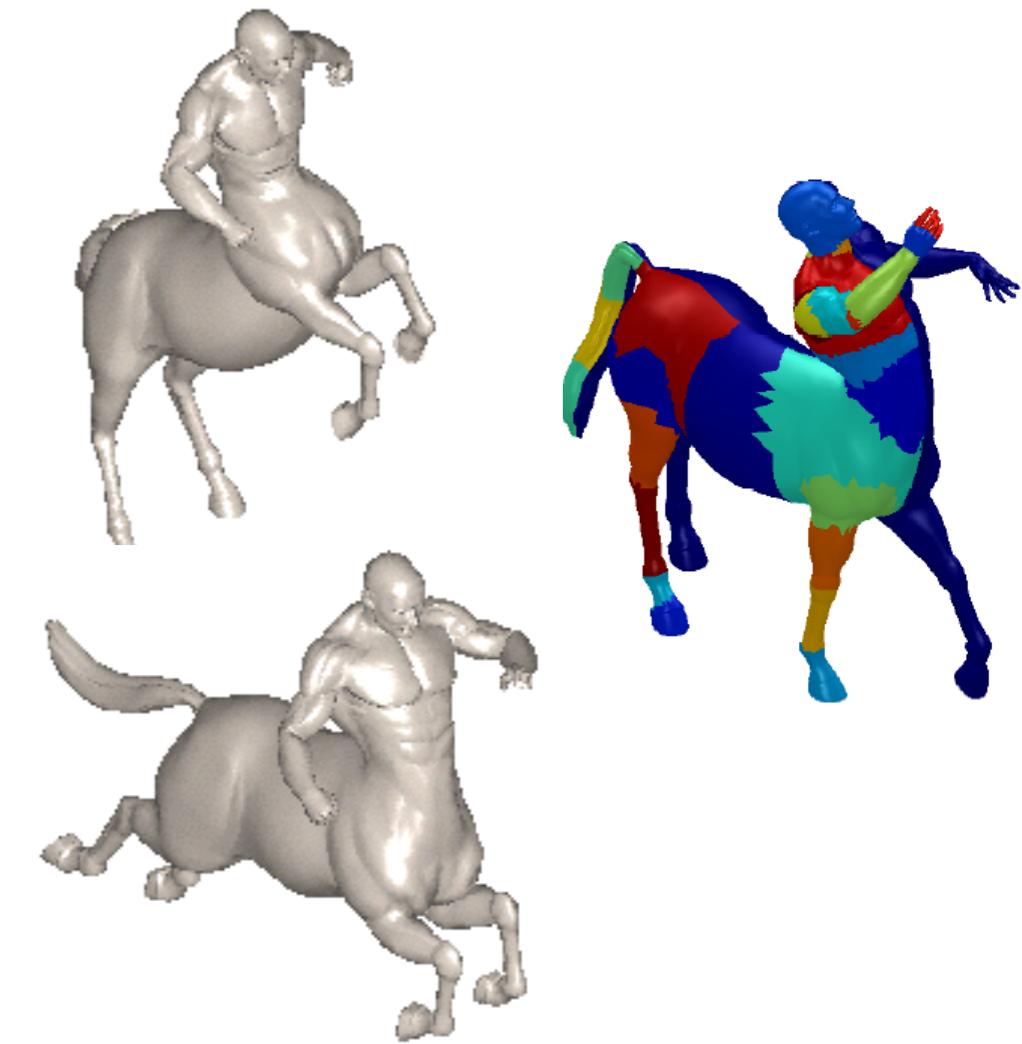
# Bayesian Nonparametric Discovery of Layers and Parts from Scenes and Objects

Soumya Ghosh

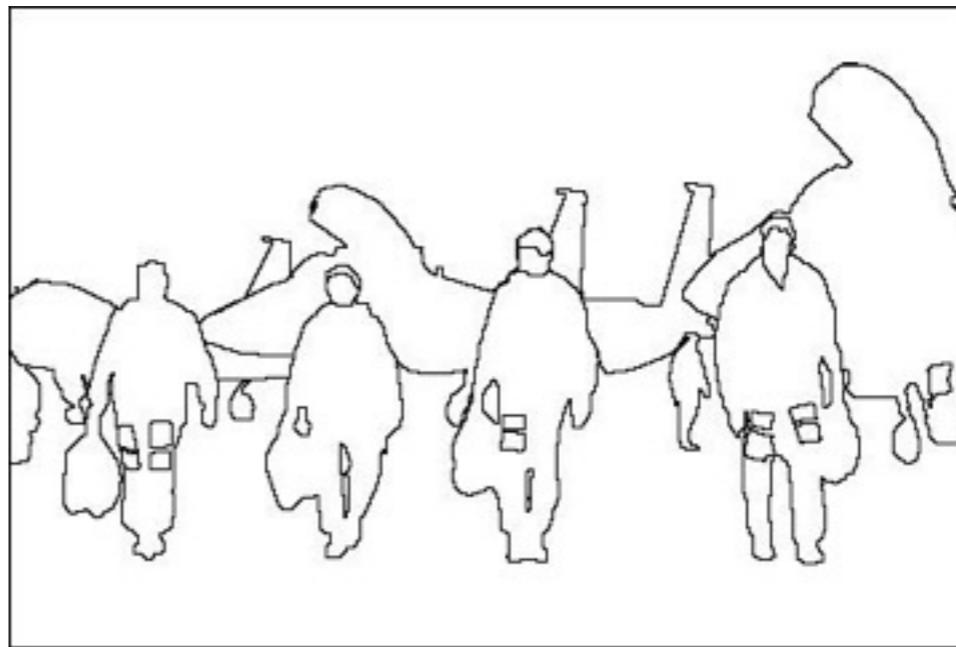
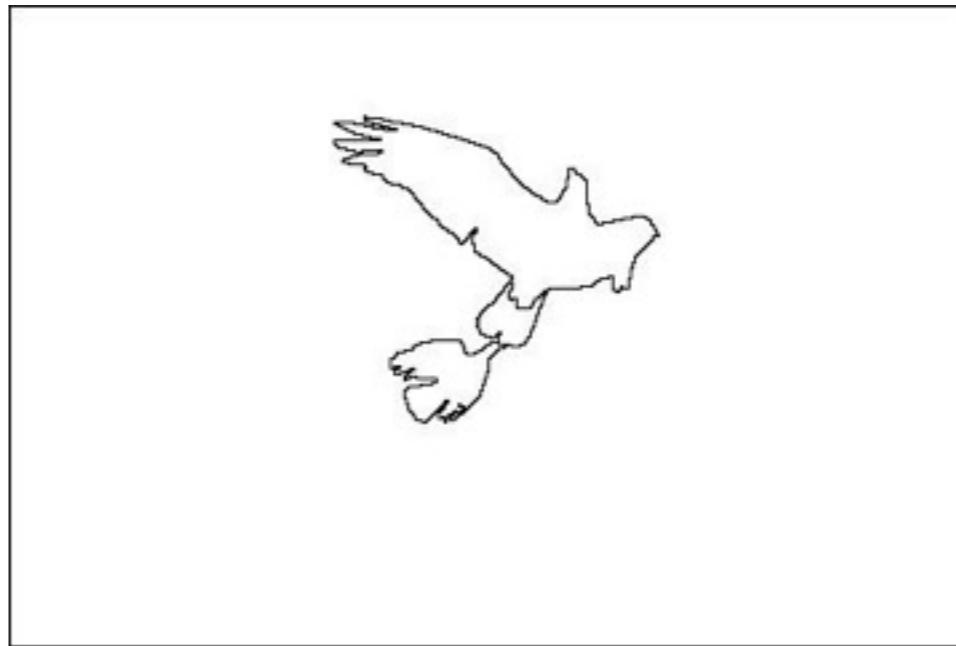
Advisor: Erik Sudderth

Committee: Michael Black and James Hays

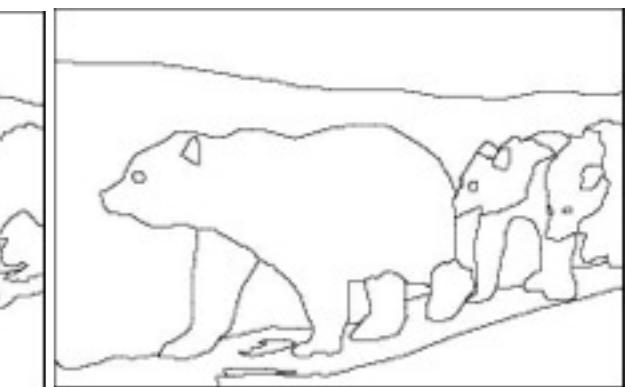
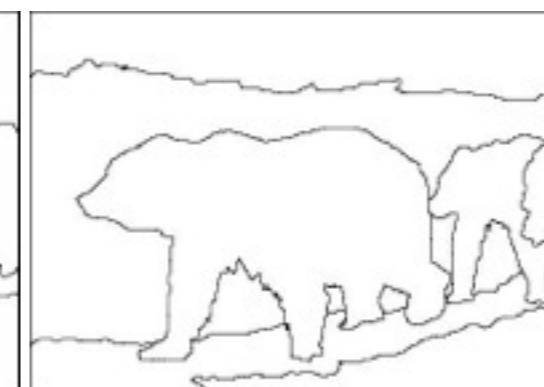
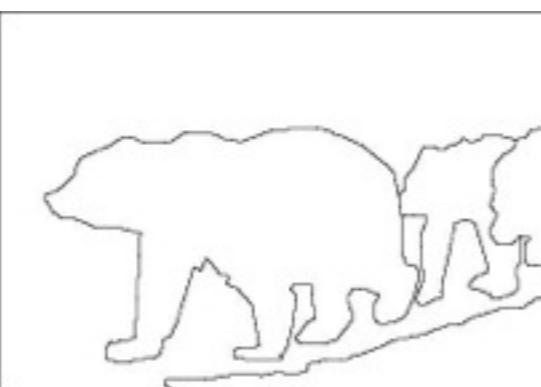
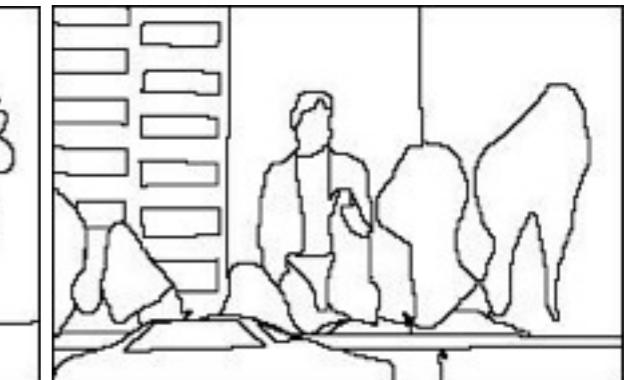
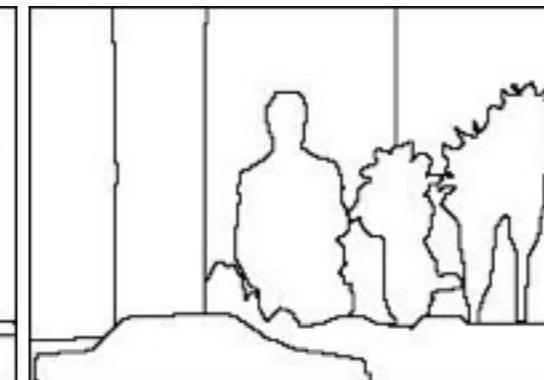
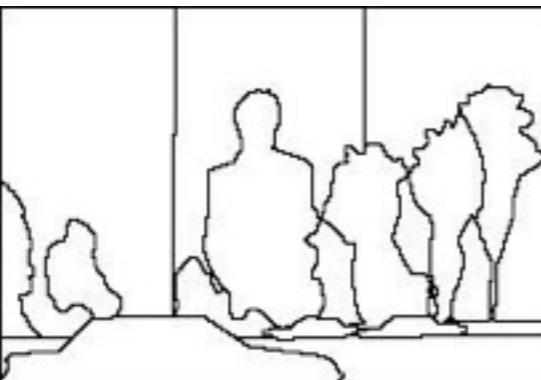
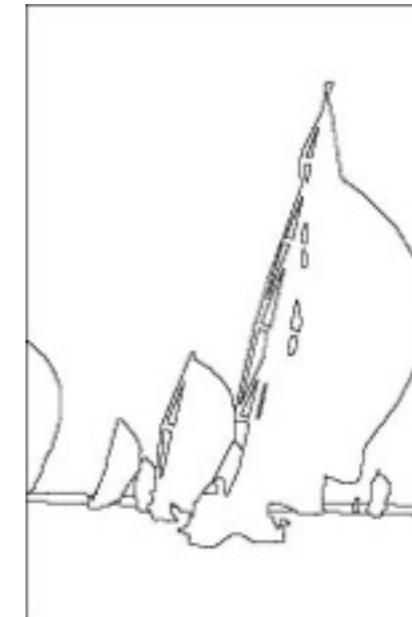
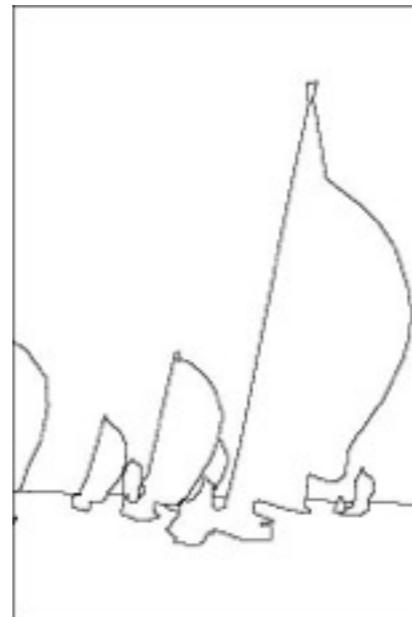
# Introduction



# World is Complex



# World is Complex

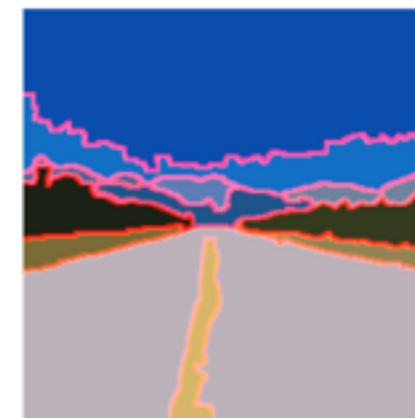
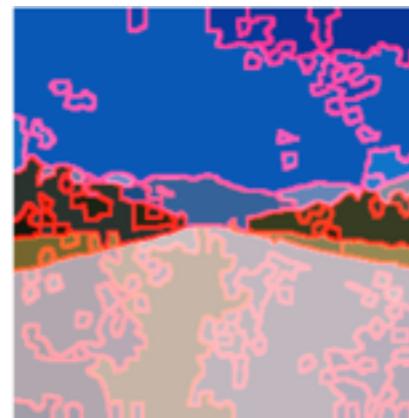
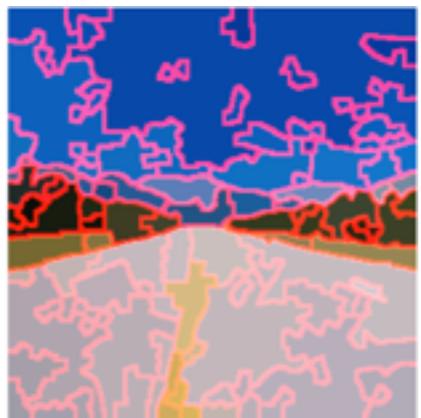


Human Segmentations

# Model Desiderata

- Automatic **model selection** - adapt to variability in image/video/object complexity
- Manage **uncertainty** - retain a distribution over possible segmentations
- Model **spatial** and **temporal** correlations
- Match **statistics** of human segmentations

# Adapting to complexity: Distributions over partitions



$$p(\mathbf{Z}_1)$$

$$p(\mathbf{X}|\mathbf{Z}_1)$$

Data

Partition

$$p(\mathbf{Z}_2)$$

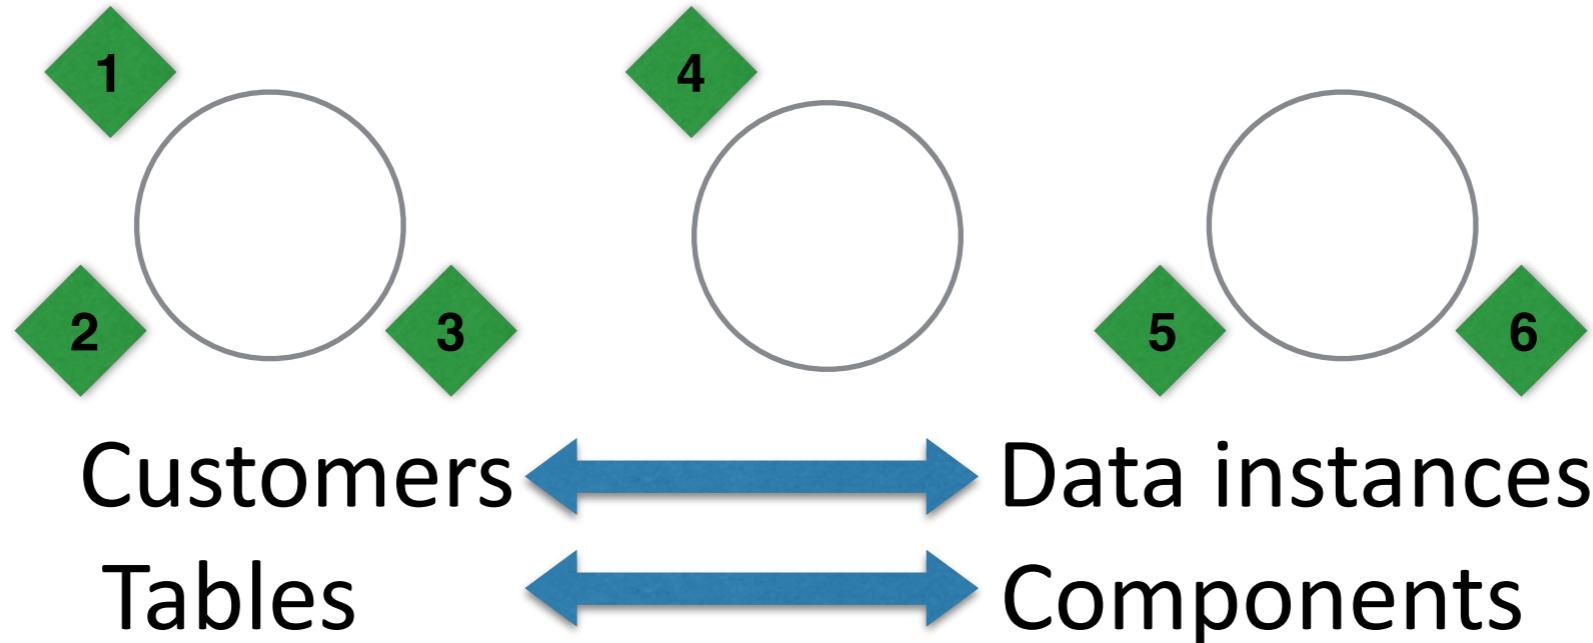
$$p(\mathbf{X}|\mathbf{Z}_2)$$

$$p(\mathbf{Z}_3)$$

$$p(\mathbf{X}|\mathbf{Z}_3)$$

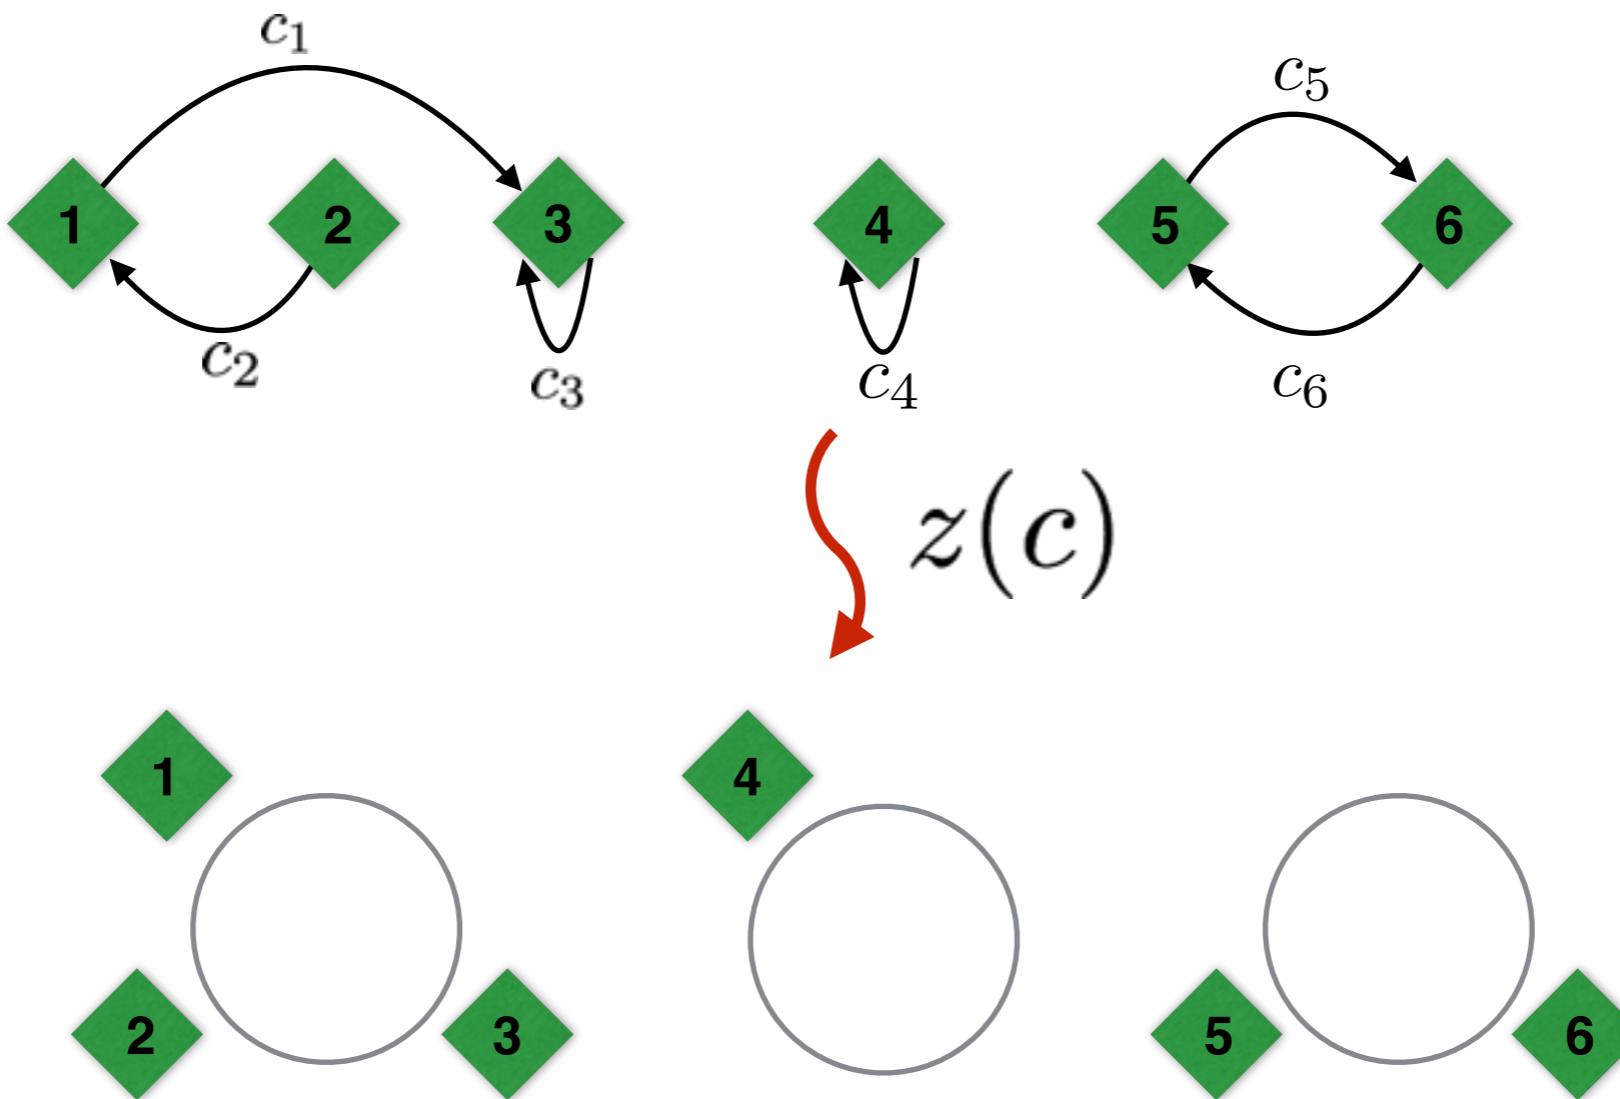
$$\mathbf{Z}^* = \underset{\mathbf{Z}_i}{\operatorname{arg\max}} p(\mathbf{Z}|\mathbf{X}) p(\mathbf{X}|\mathbf{Z}) p(\mathbf{Z}) (\mathbf{Z} = \mathbf{Z}_i)$$

# A distribution over partitions: Chinese Restaurant Process



- A restaurant with infinitely many tables, each table with infinite capacity.
- The first customer sits at a table.
- Subsequent customers pick a table  $k$  w.p.:  $\propto \begin{cases} n_k & \text{if } k \text{ is an existing table} \\ \alpha & k \text{ is a new table} \end{cases}$

# Modeling Correlations: Distance dependent Chinese Restaurant Process (ddCRP)



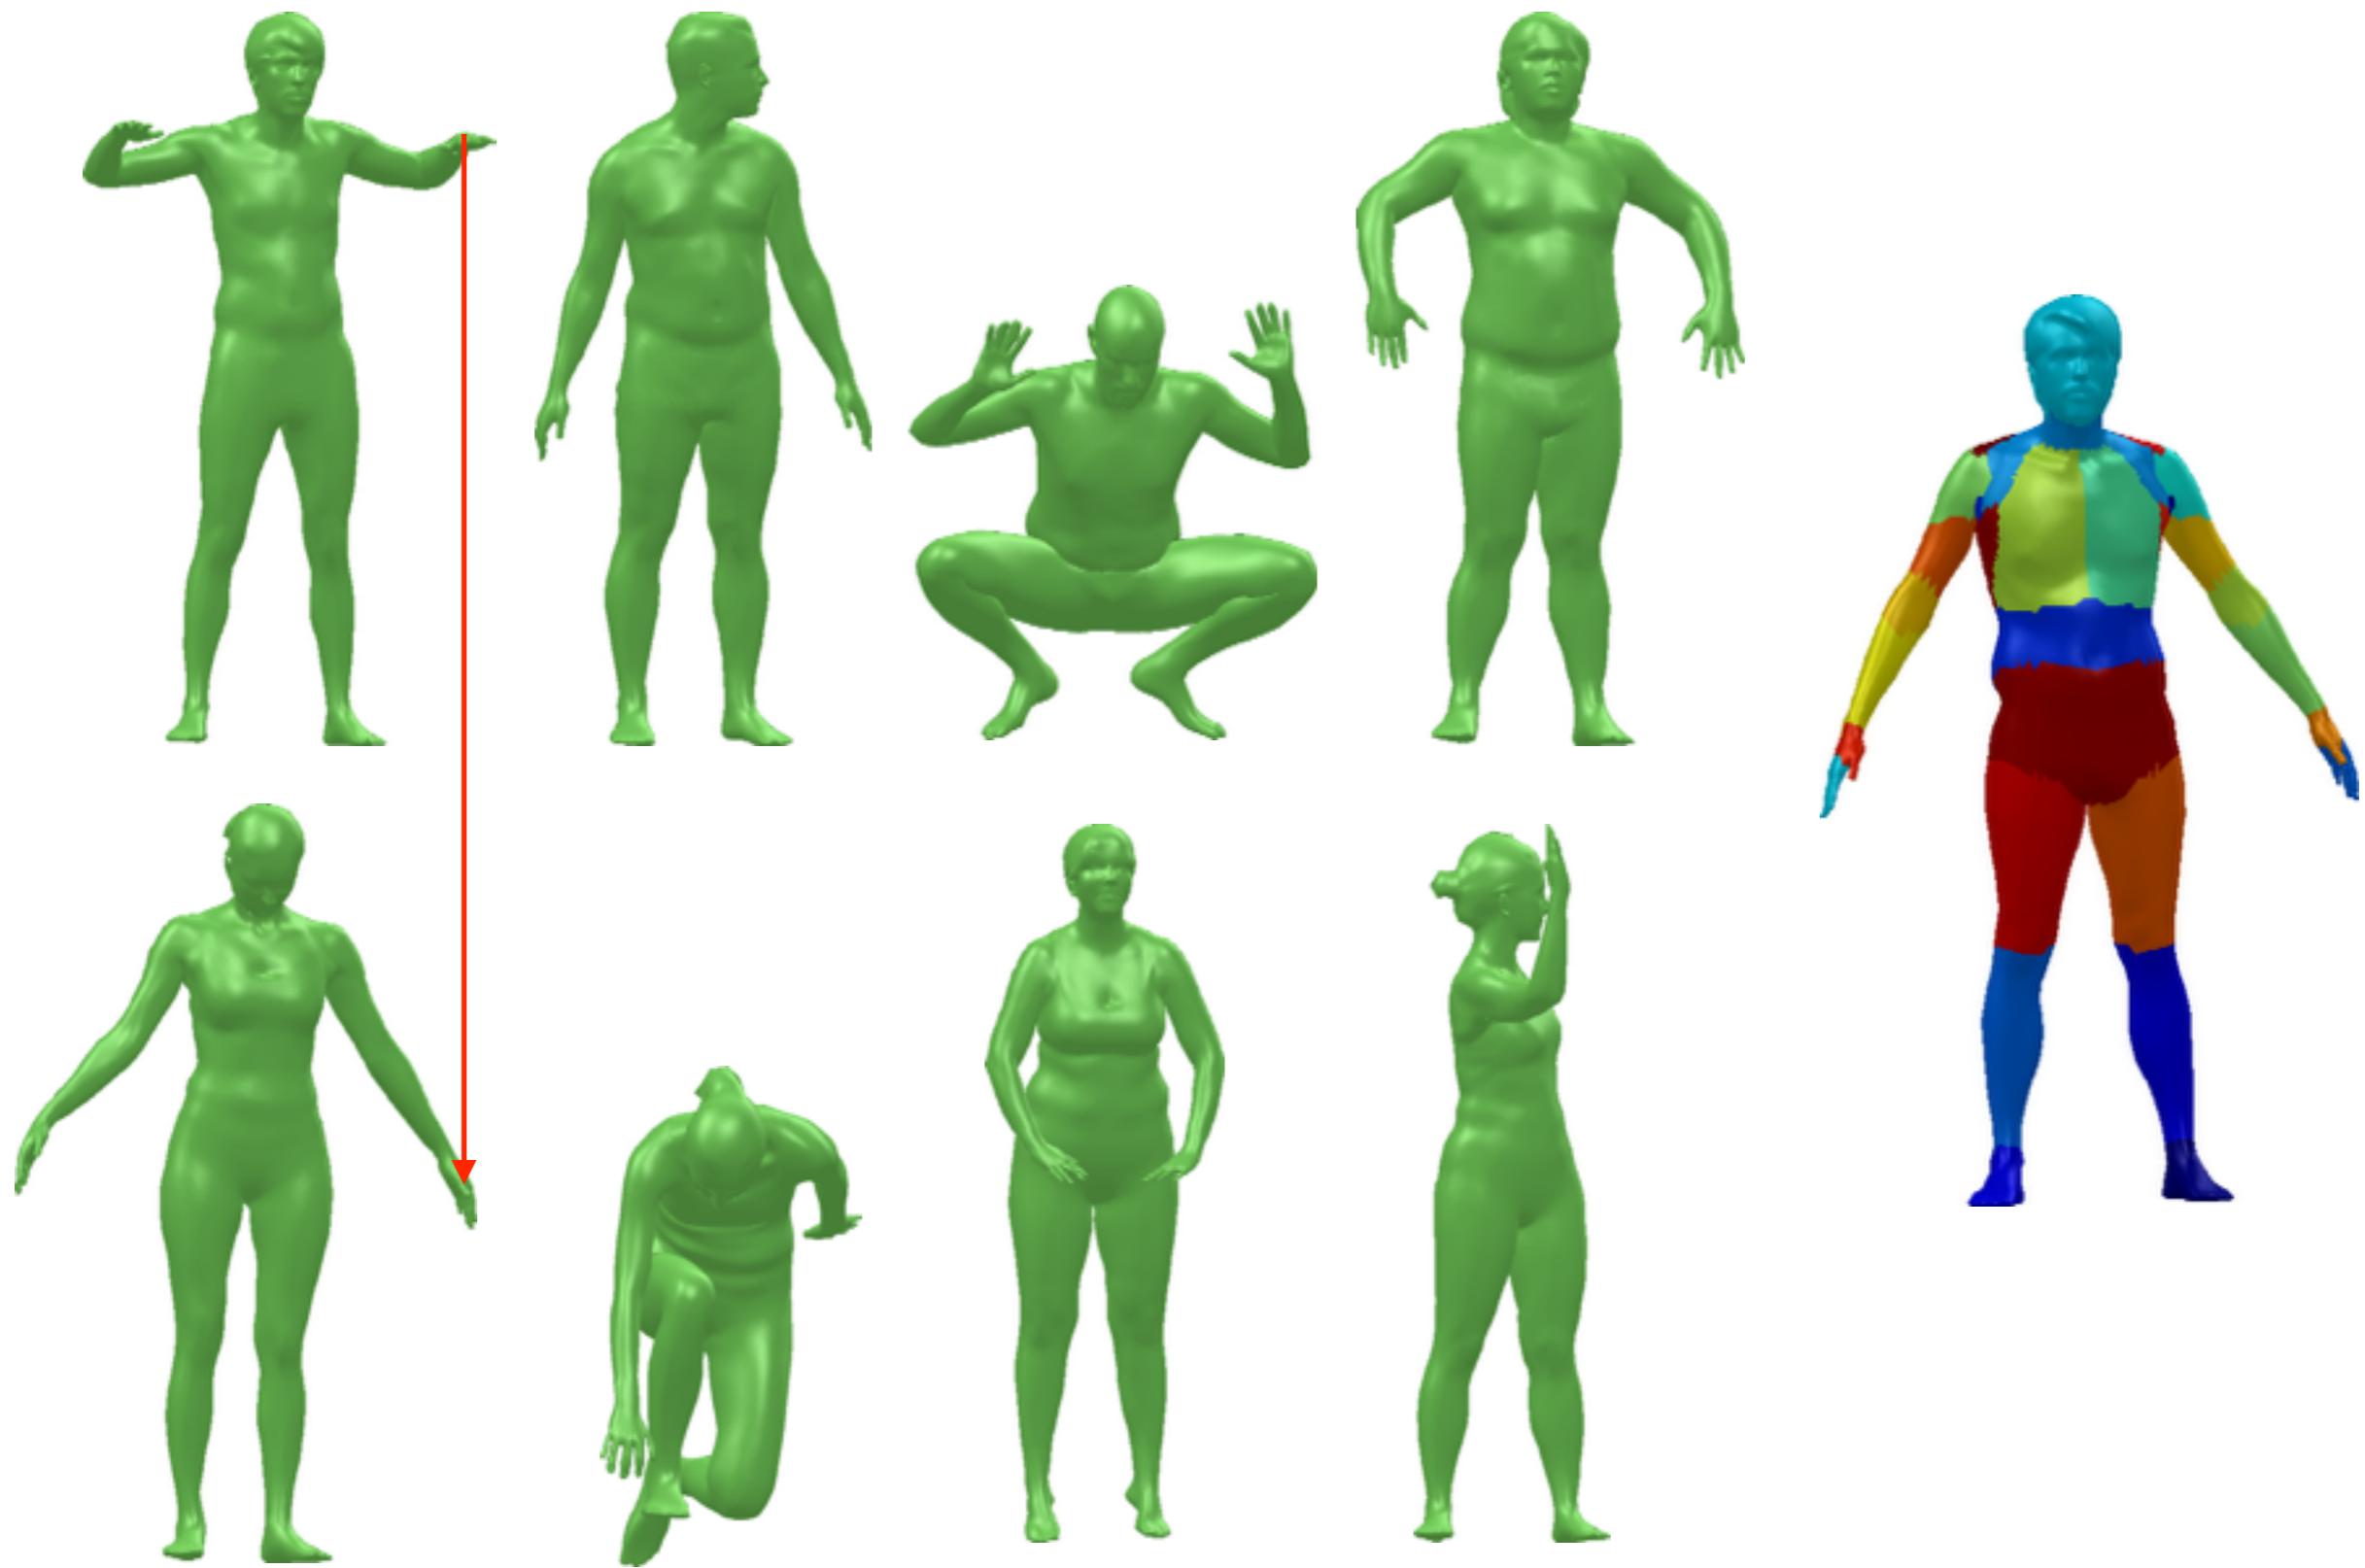
$$p(c_m = n \mid D, \alpha) \propto \begin{cases} f(d_{mn}) & \text{if } m \neq n, \\ \alpha & \text{if } m = n. \end{cases}$$

$$p(z(c) \mid D, \alpha) = \prod_n p(c_n \mid D, \alpha)$$

# Talk Outline

- Deformable 3D object segmentation
- Image and Video segmentation
- Layered decomposition of natural images
- Proposed Work

# Discovering Parts from Deformations

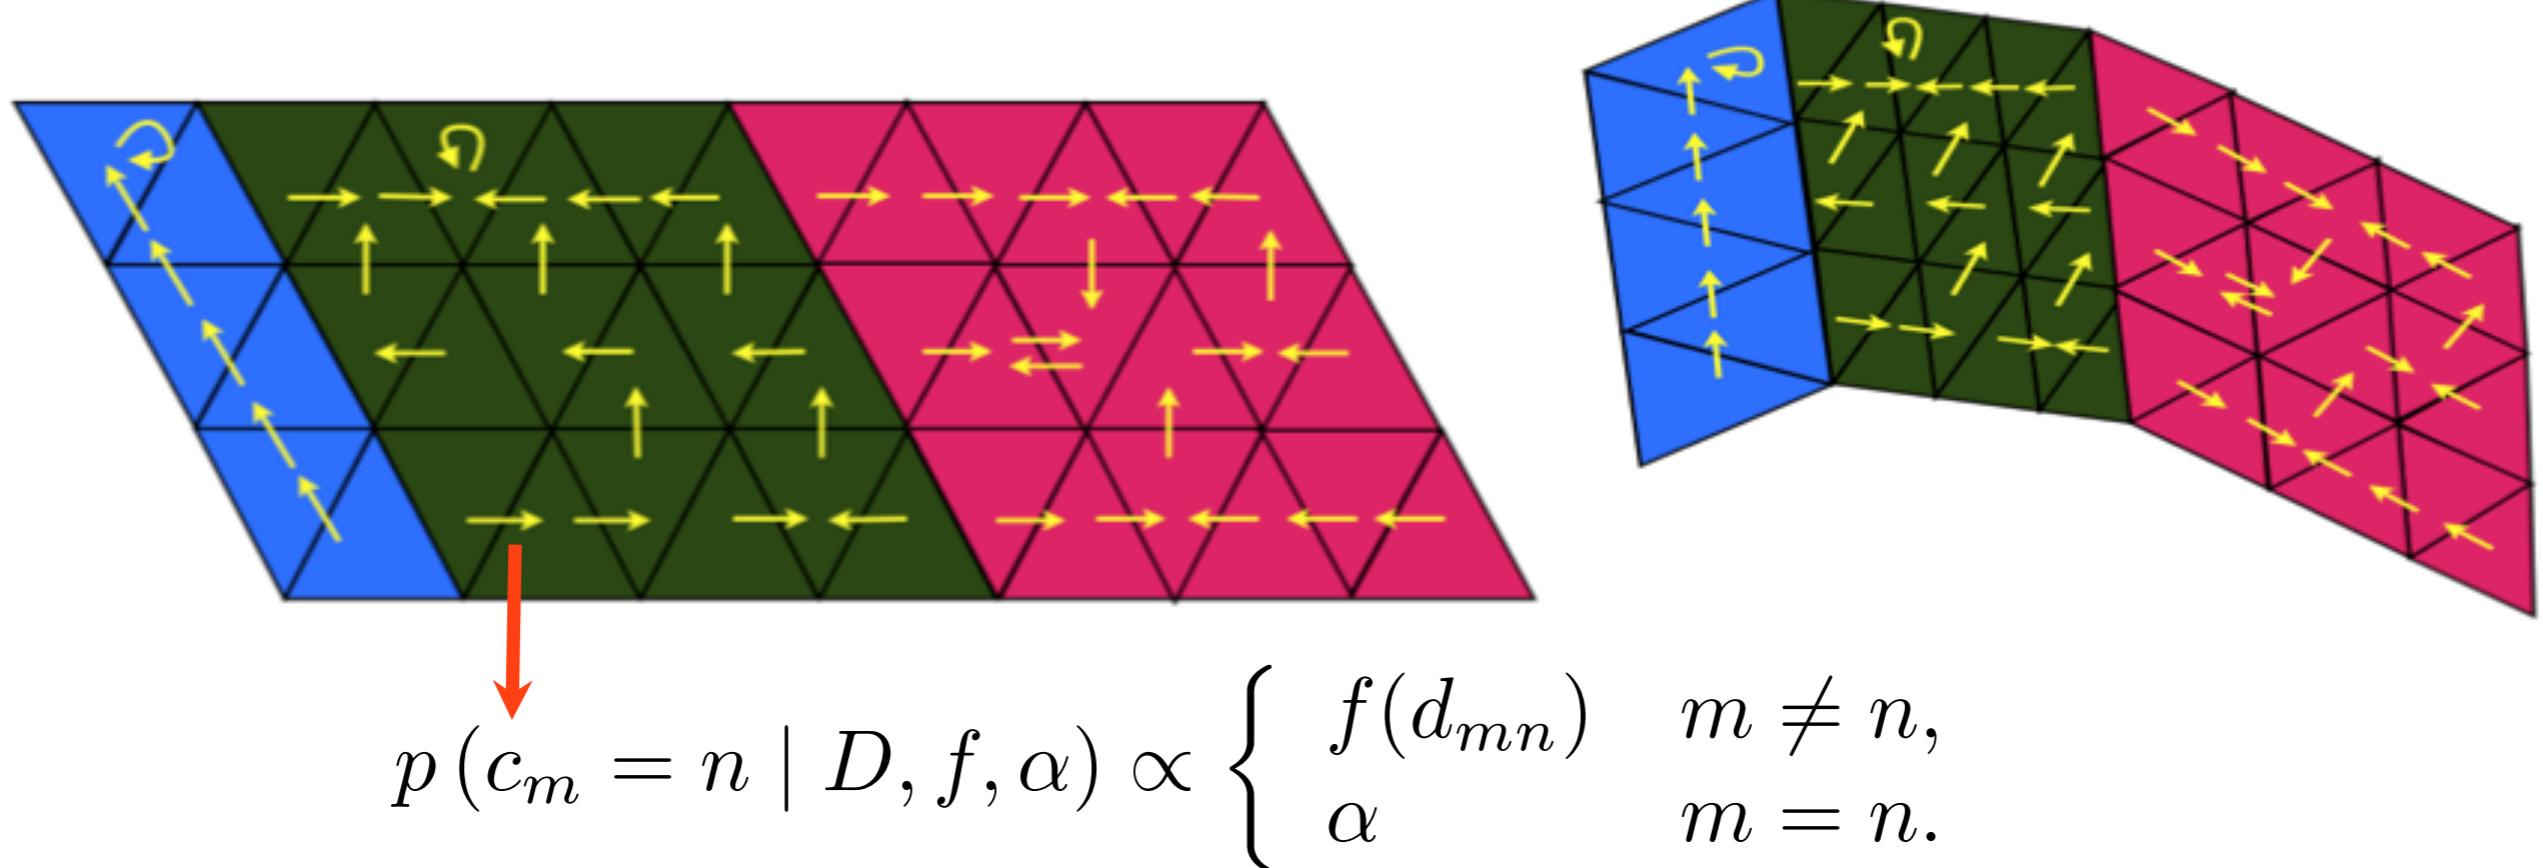


# Discovering Parts from Deformations: Big Picture

- **Cluster:** Mesh faces.
- **Prior:** over the space of plausible mesh partitions.
- **Likelihood:** Given segmentation into parts, model how multiple bodies deform across many poses.
- **Posterior:** Explored through MCMC.



# ddCRP Prior over Mesh Partitions



- Mesh faces are only allowed to link to neighboring faces:

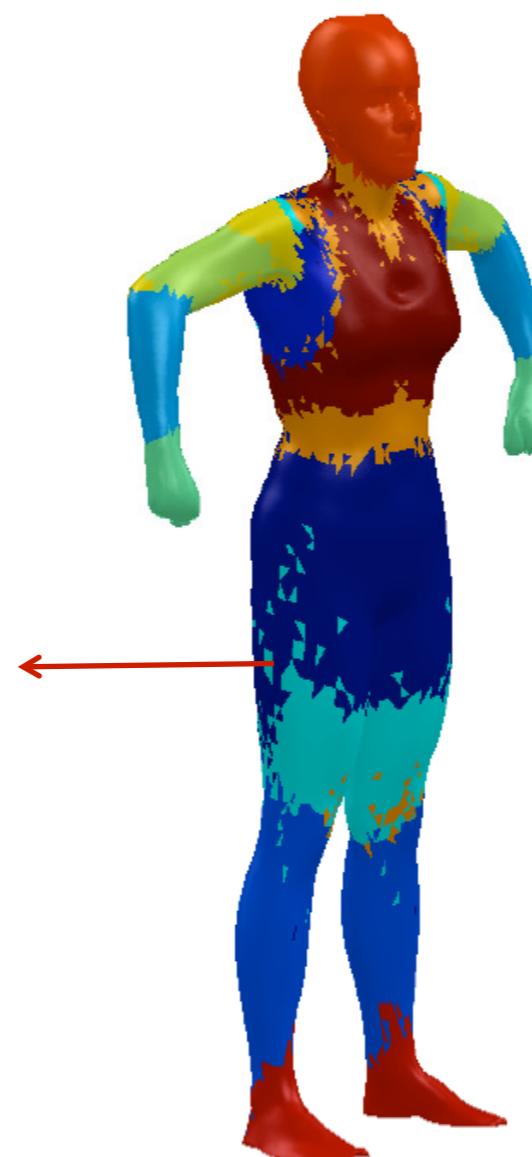
$$f(d_{mn}) = \mathbf{1}[d_{mn} \leq 1]$$

# Prior over plausible partitions



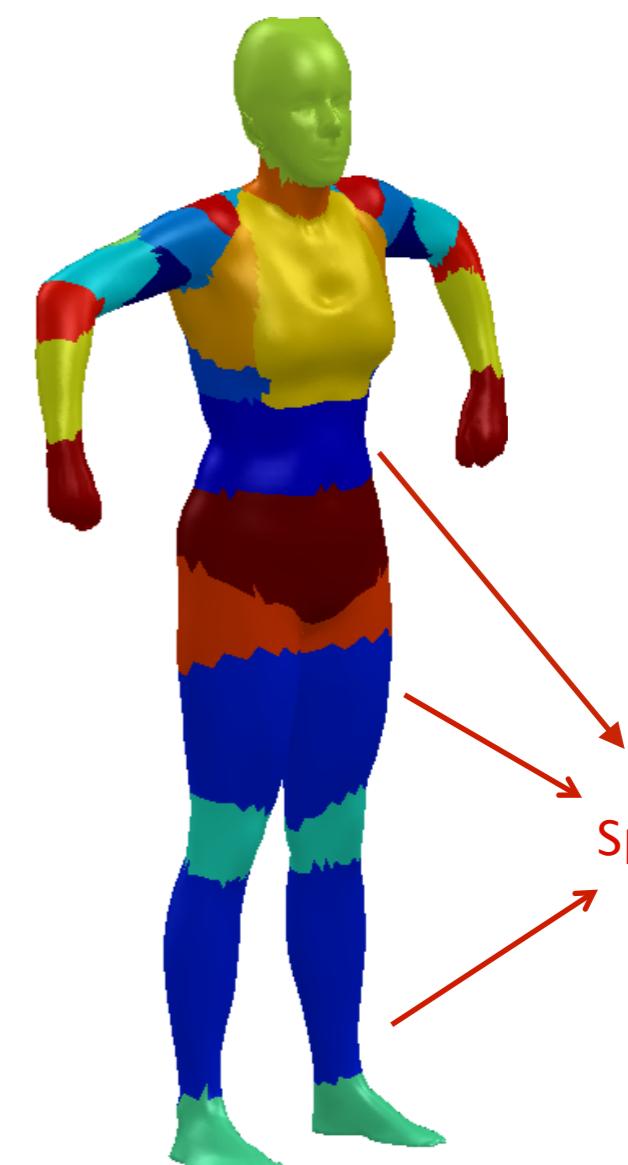
$$p(Z_1) > 0$$

*Desirable*



$$p(Z_2) = 0$$

*Avoid: Noisy Parts*



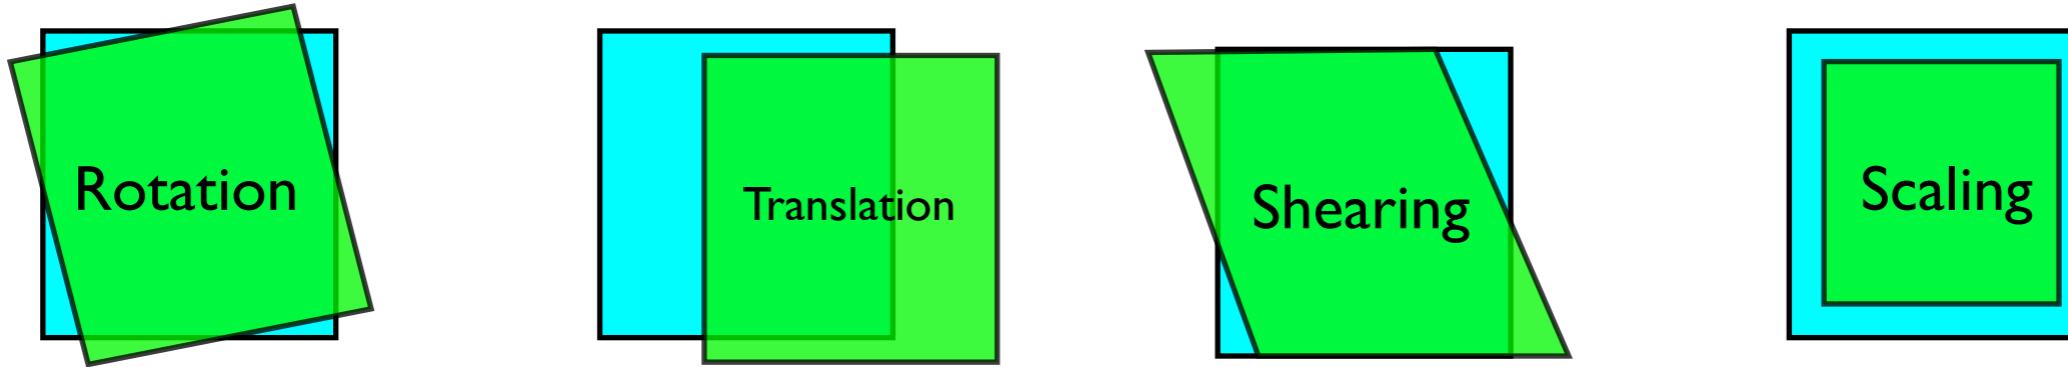
$$p(Z_3) = 0$$

*Avoid: Disconnected Parts*

Noise

Split limbs

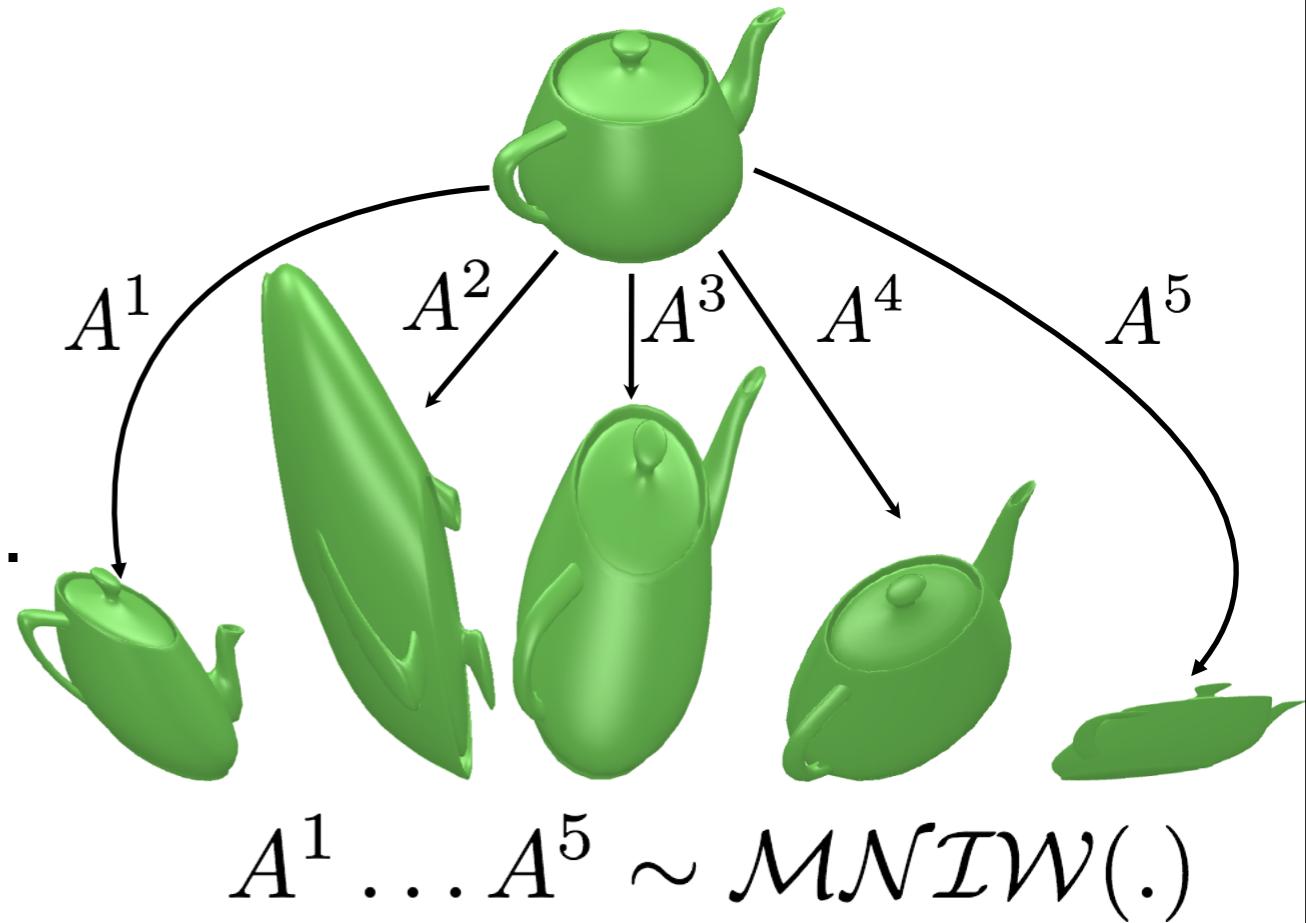
# Distributions over Affine Transformations



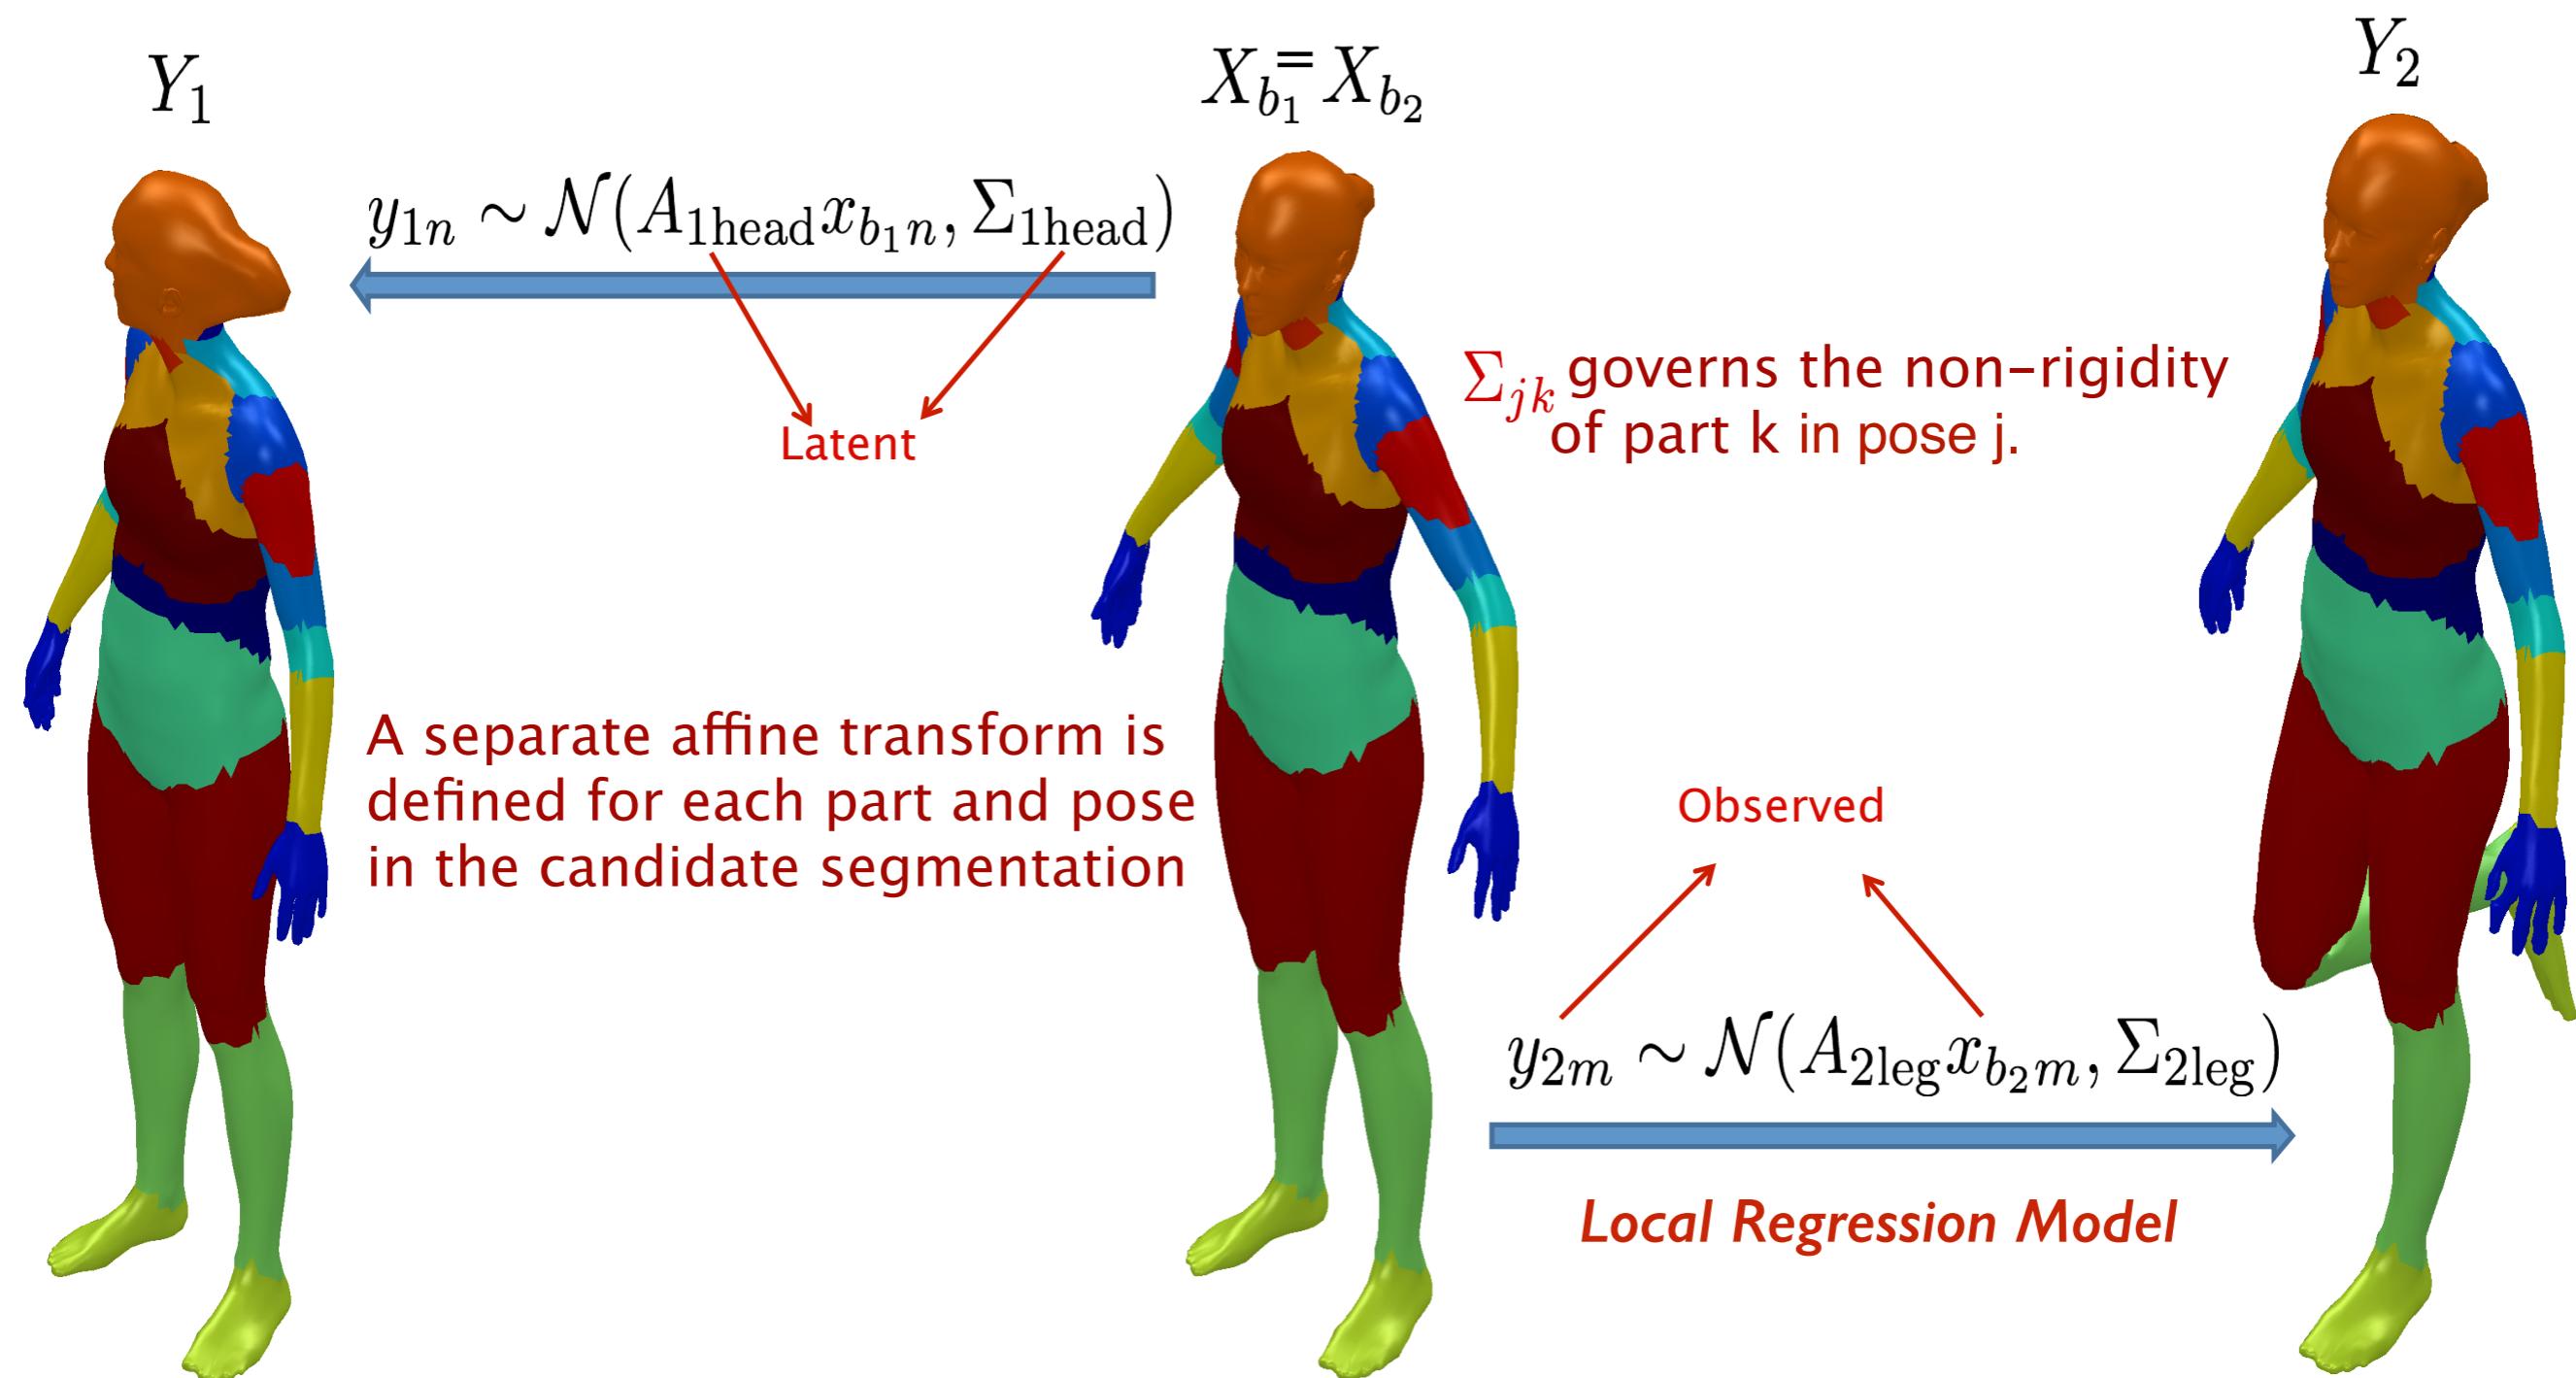
*Matrix Normal Inverse Wishart:*

$$\begin{aligned}\Sigma &\sim \mathcal{IW}(n_0, S_0) \\ A \mid \Sigma &\sim \mathcal{MN}(M, \Sigma, K)\end{aligned}$$

where  $A \in \mathbb{R}^{3 \times 4}$  is an **affine** transformation.



# Generative Affine Likelihoods



# Marginal Affine Likelihoods

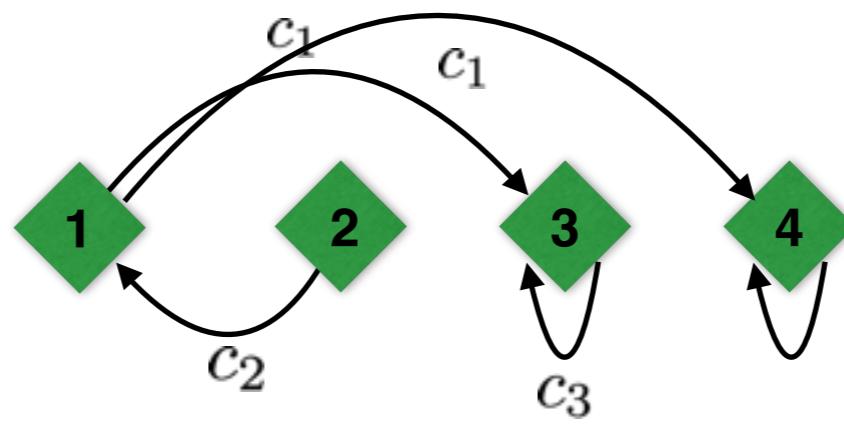
- We do not need to estimate the latent affine transformations. For each part and pose combination we can analytically marginalize them

$$p(Y_{jk} \mid X_{jk}) = \int p(Y_{jk}, A_{jk}, \Sigma_{jk} \mid X_{jk}) dA_{jk}, d\Sigma_{jk}$$

**Marginal Likelihood**

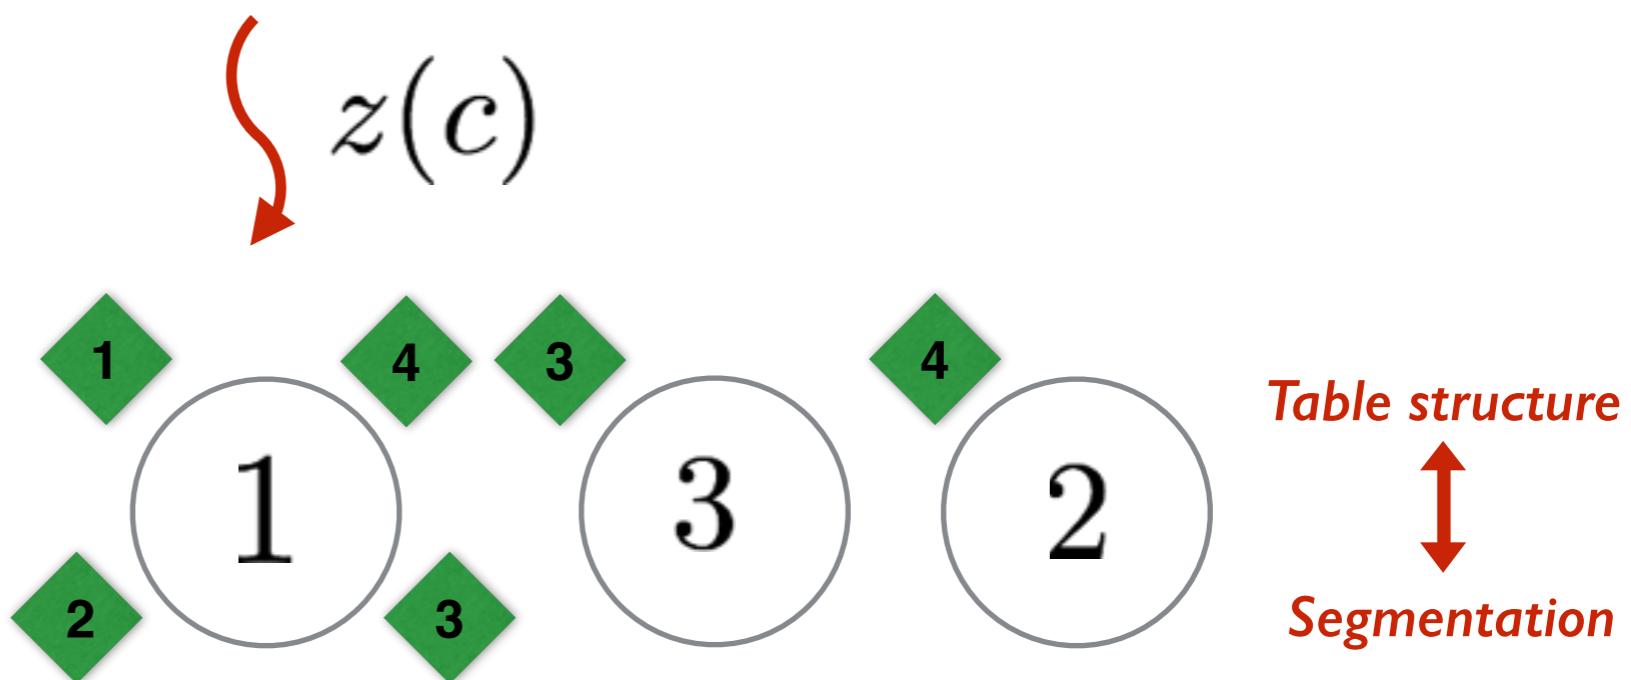
- Bayesian Model Selection:
  - Improper merges have low marginal likelihoods
  - Improper splits are “suspicious coincidences” and end up with lower marginal likelihoods

# Inference through Gibbs Sampling



*Customers = Mesh Faces  
Tables = Object Parts*

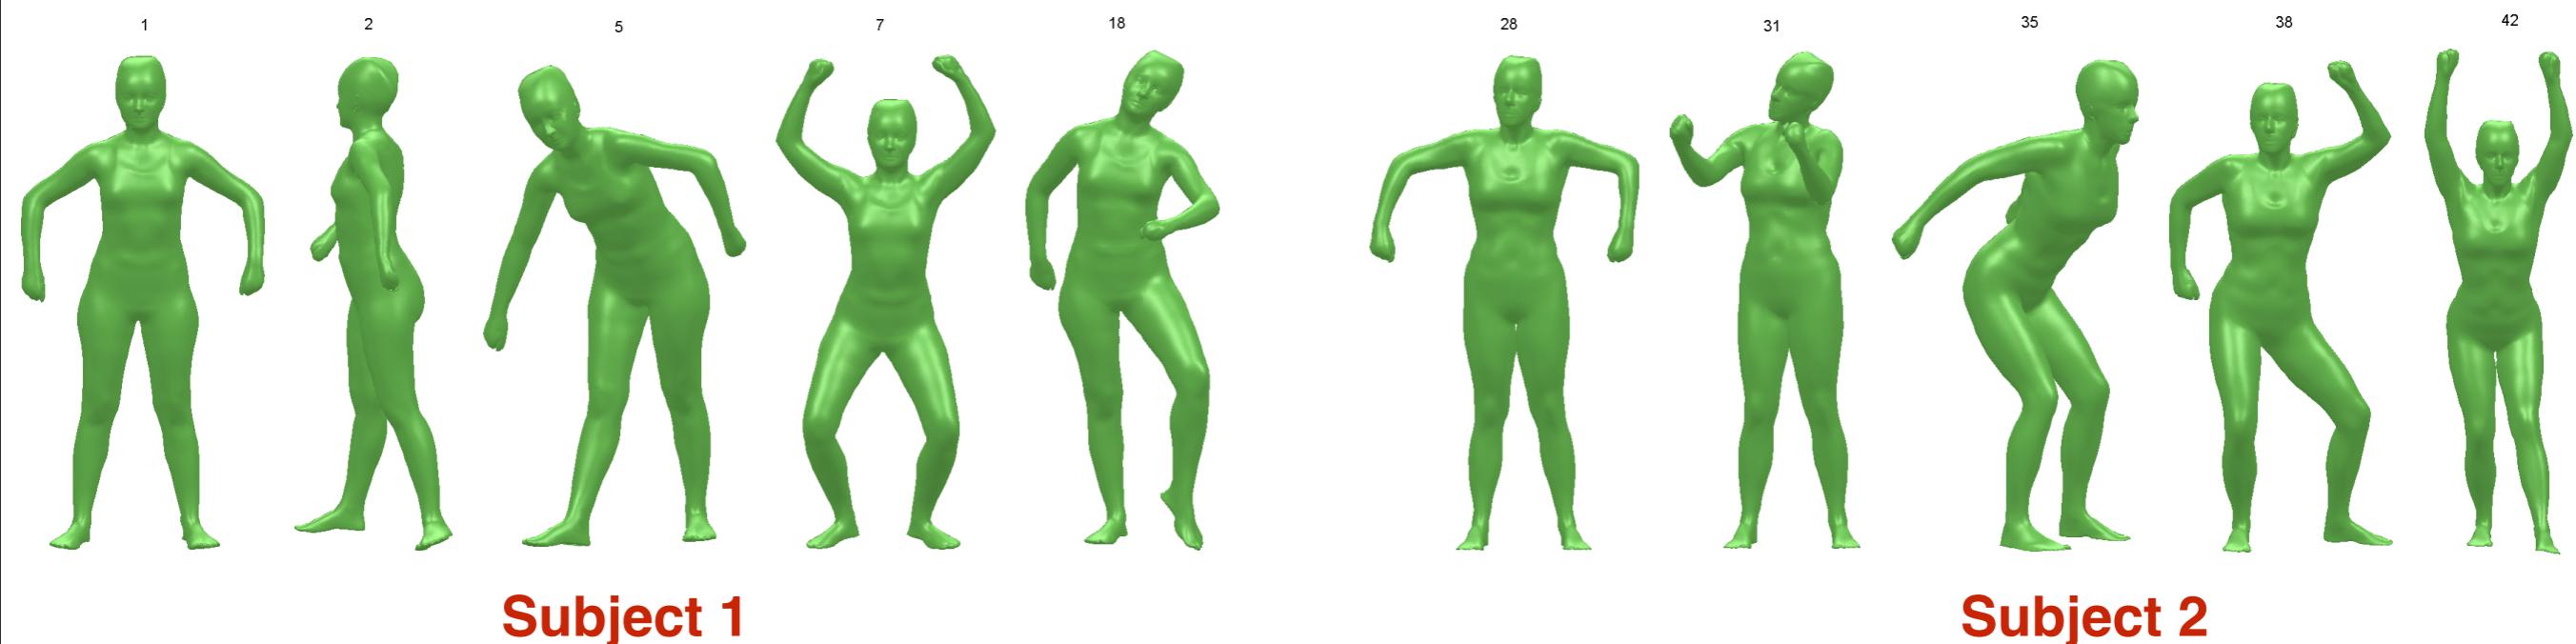
**Collapsed Sampler:**  
Only need to sample links,  
other random variables are  
analytically marginalized out.



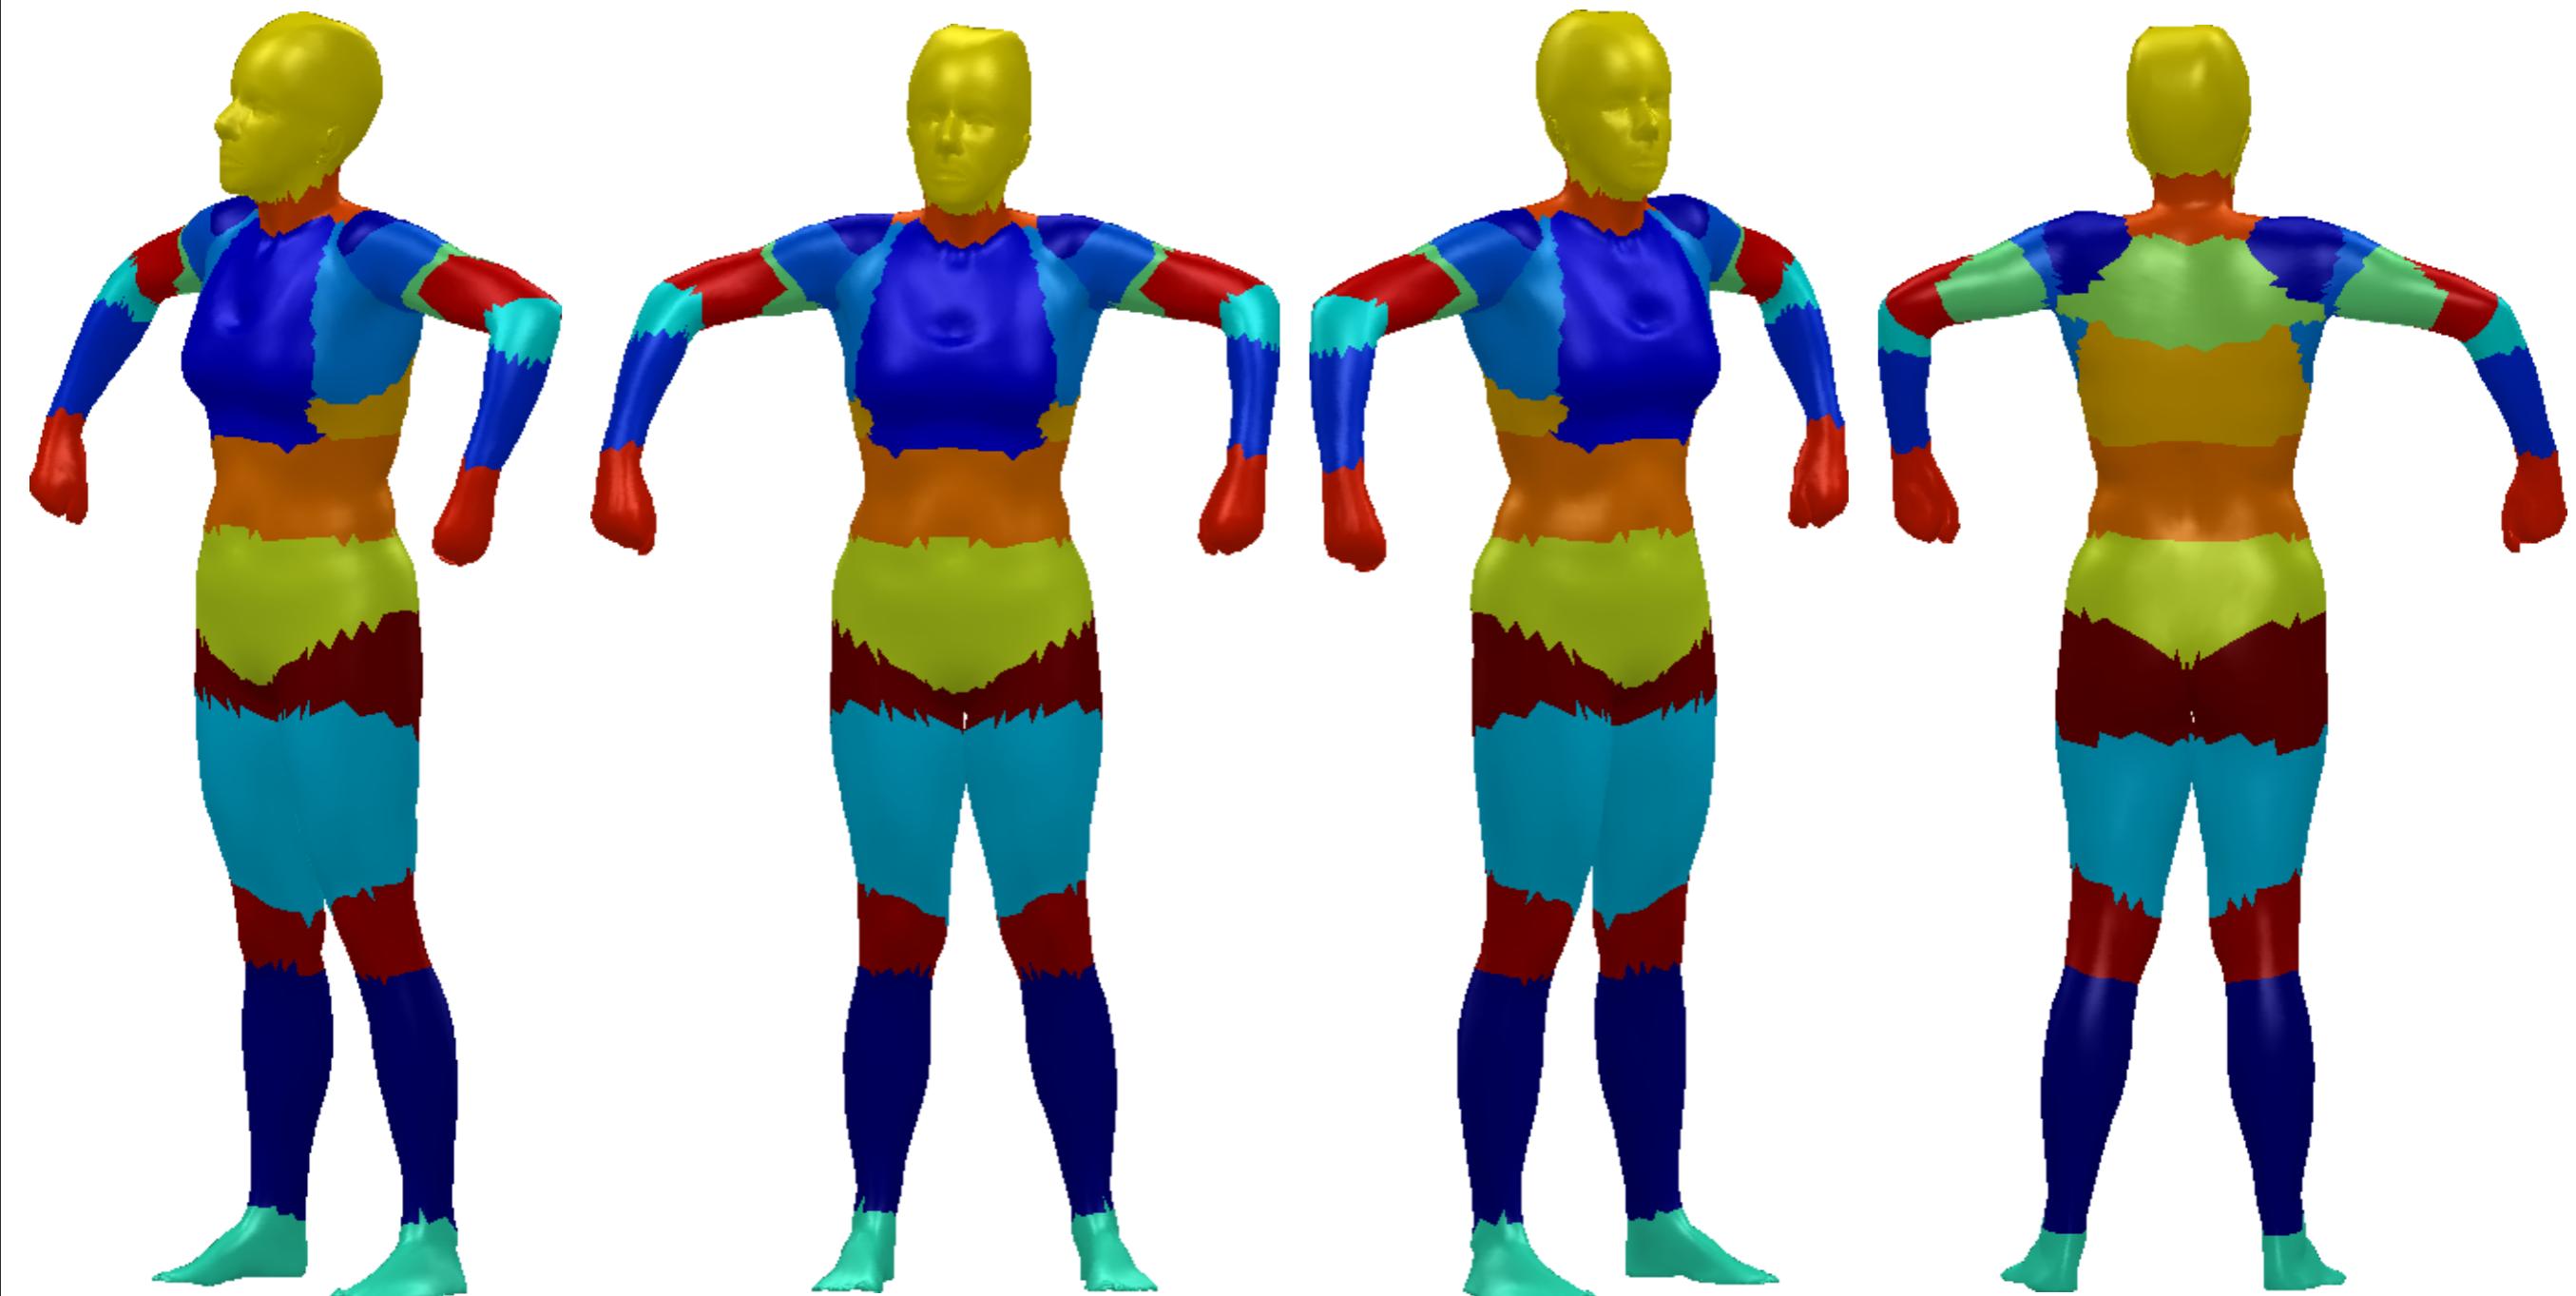
***Local changes in the link structure lead to large changes in the partition structure***

# Analysis of Human Bodies in Motion

- 56 Aligned scans from two human subjects
- Wide variability in poses, limited variability in body shapes



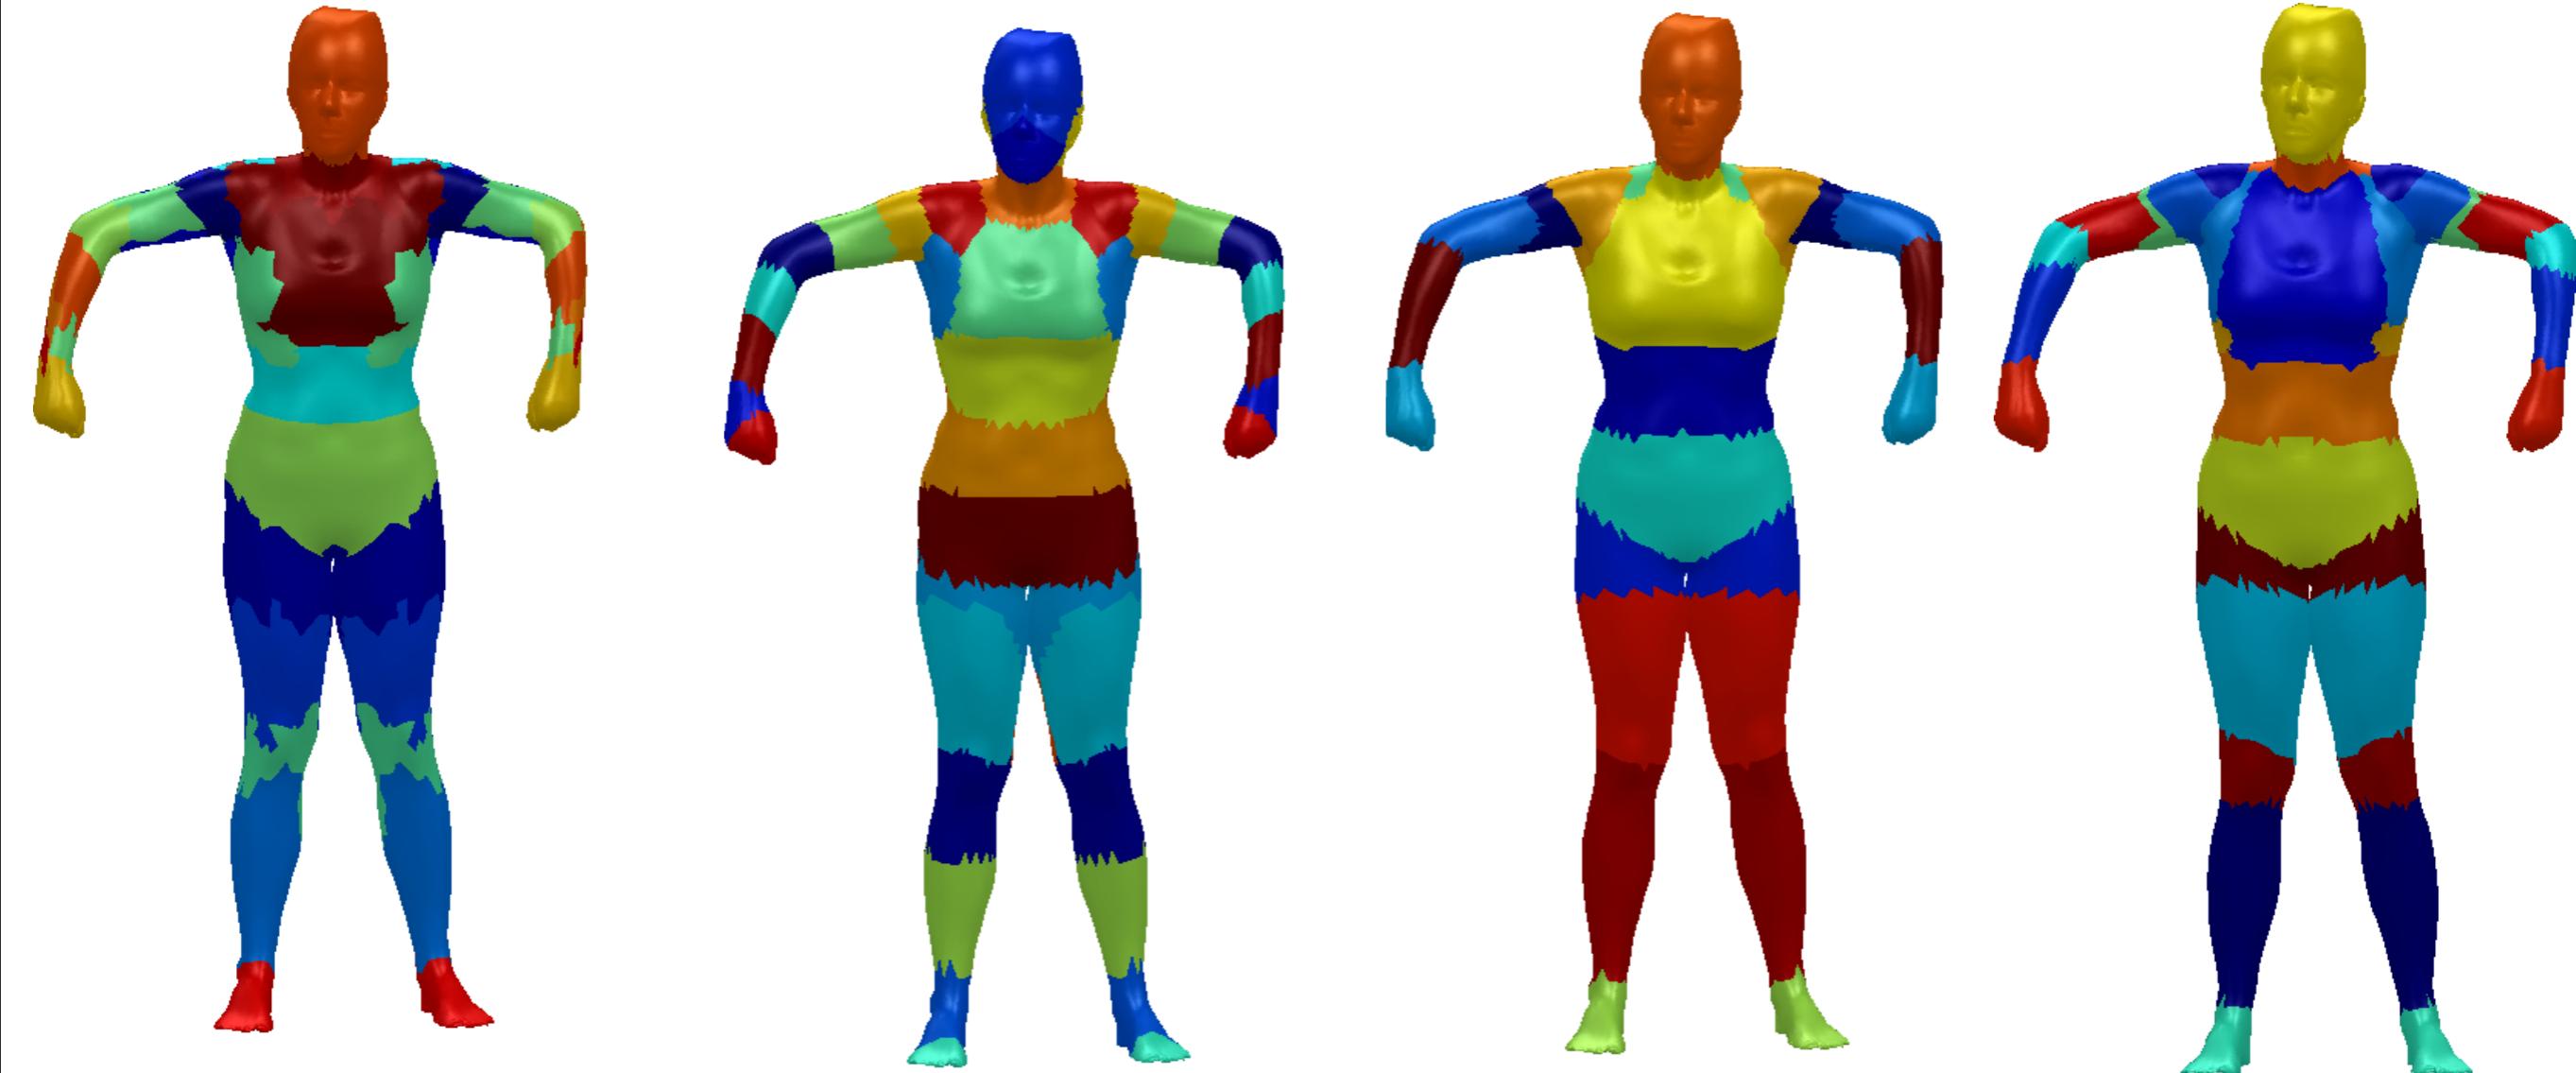
# Inferred Segmentation



**Discovered a segmentation with 20 Parts**

Ghosh et al., NIPS 2012

# Visual Comparisons



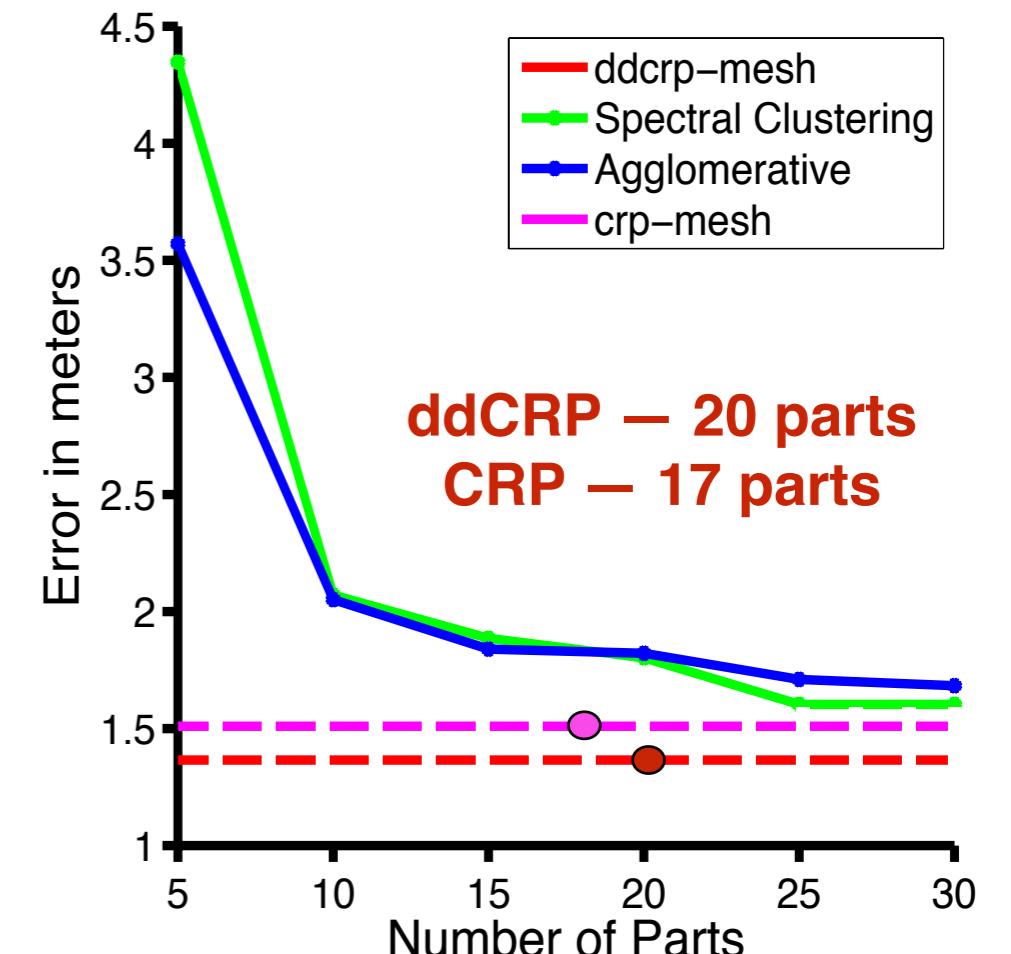
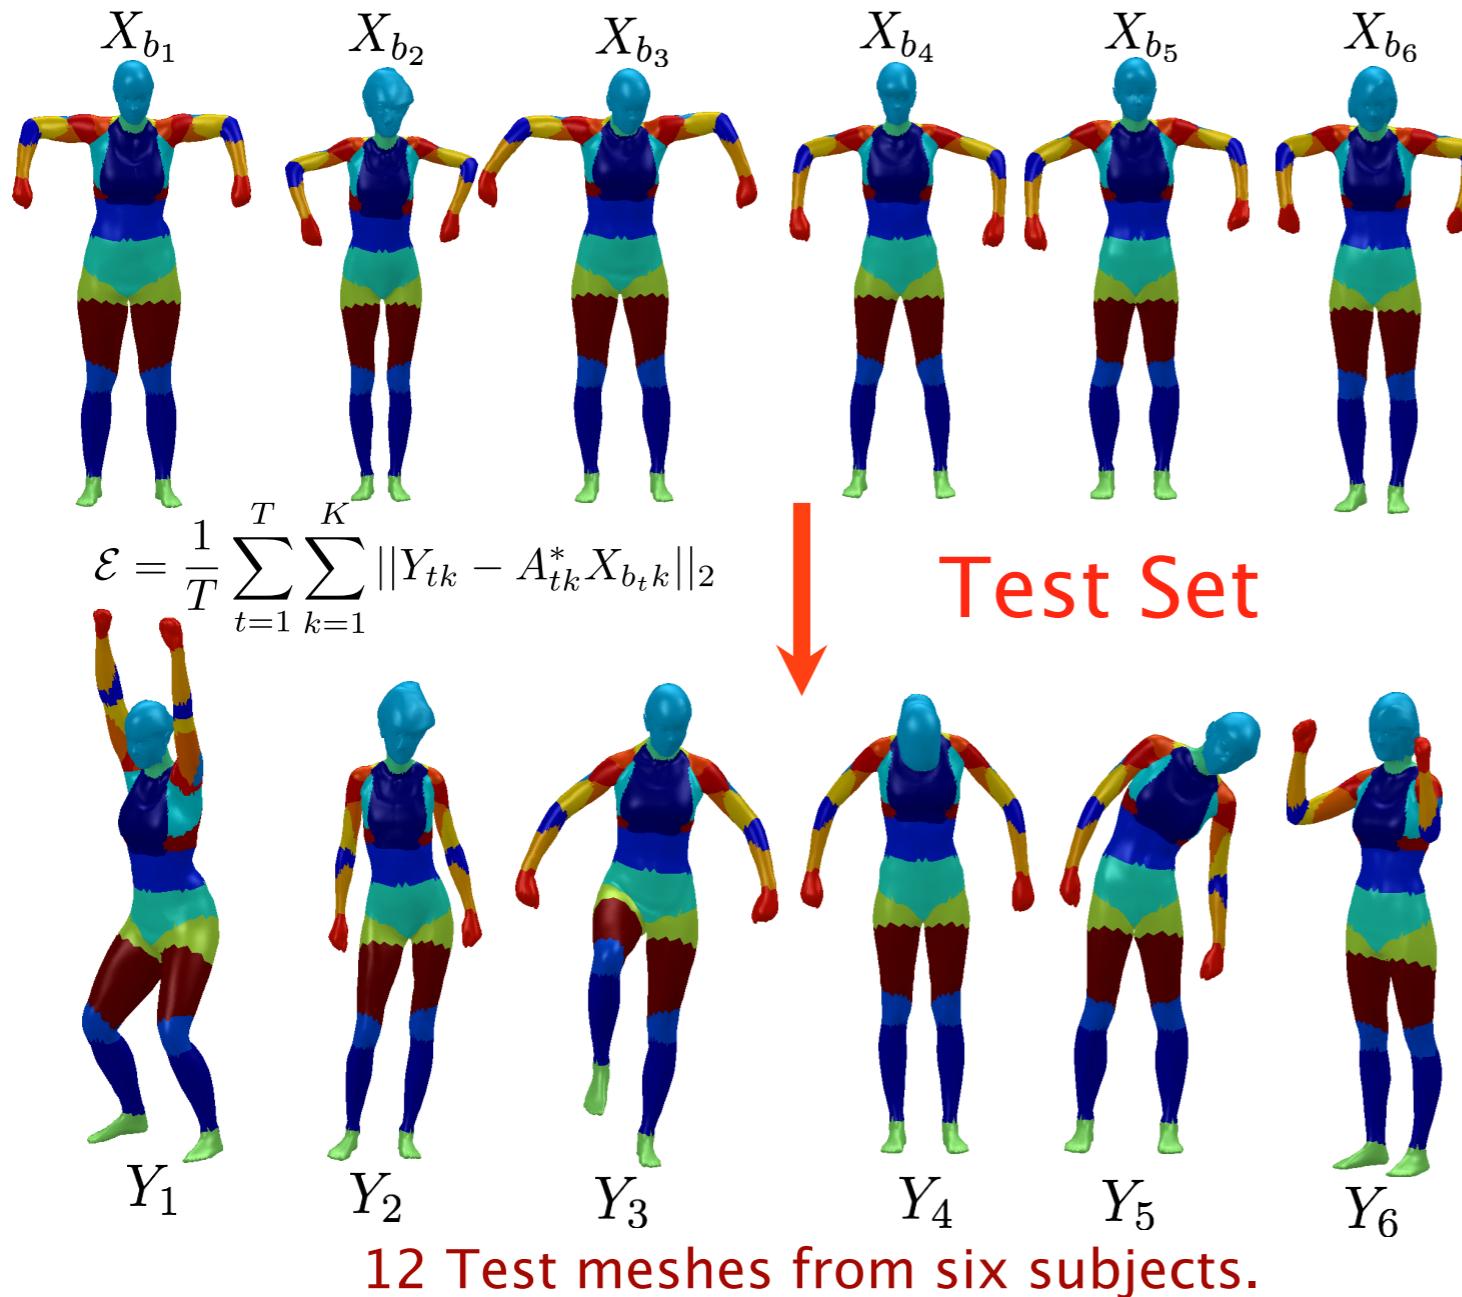
Agglomerative

Spectral Clustering  
*Liu & Zhang, 2004*

CRP

ddCRP

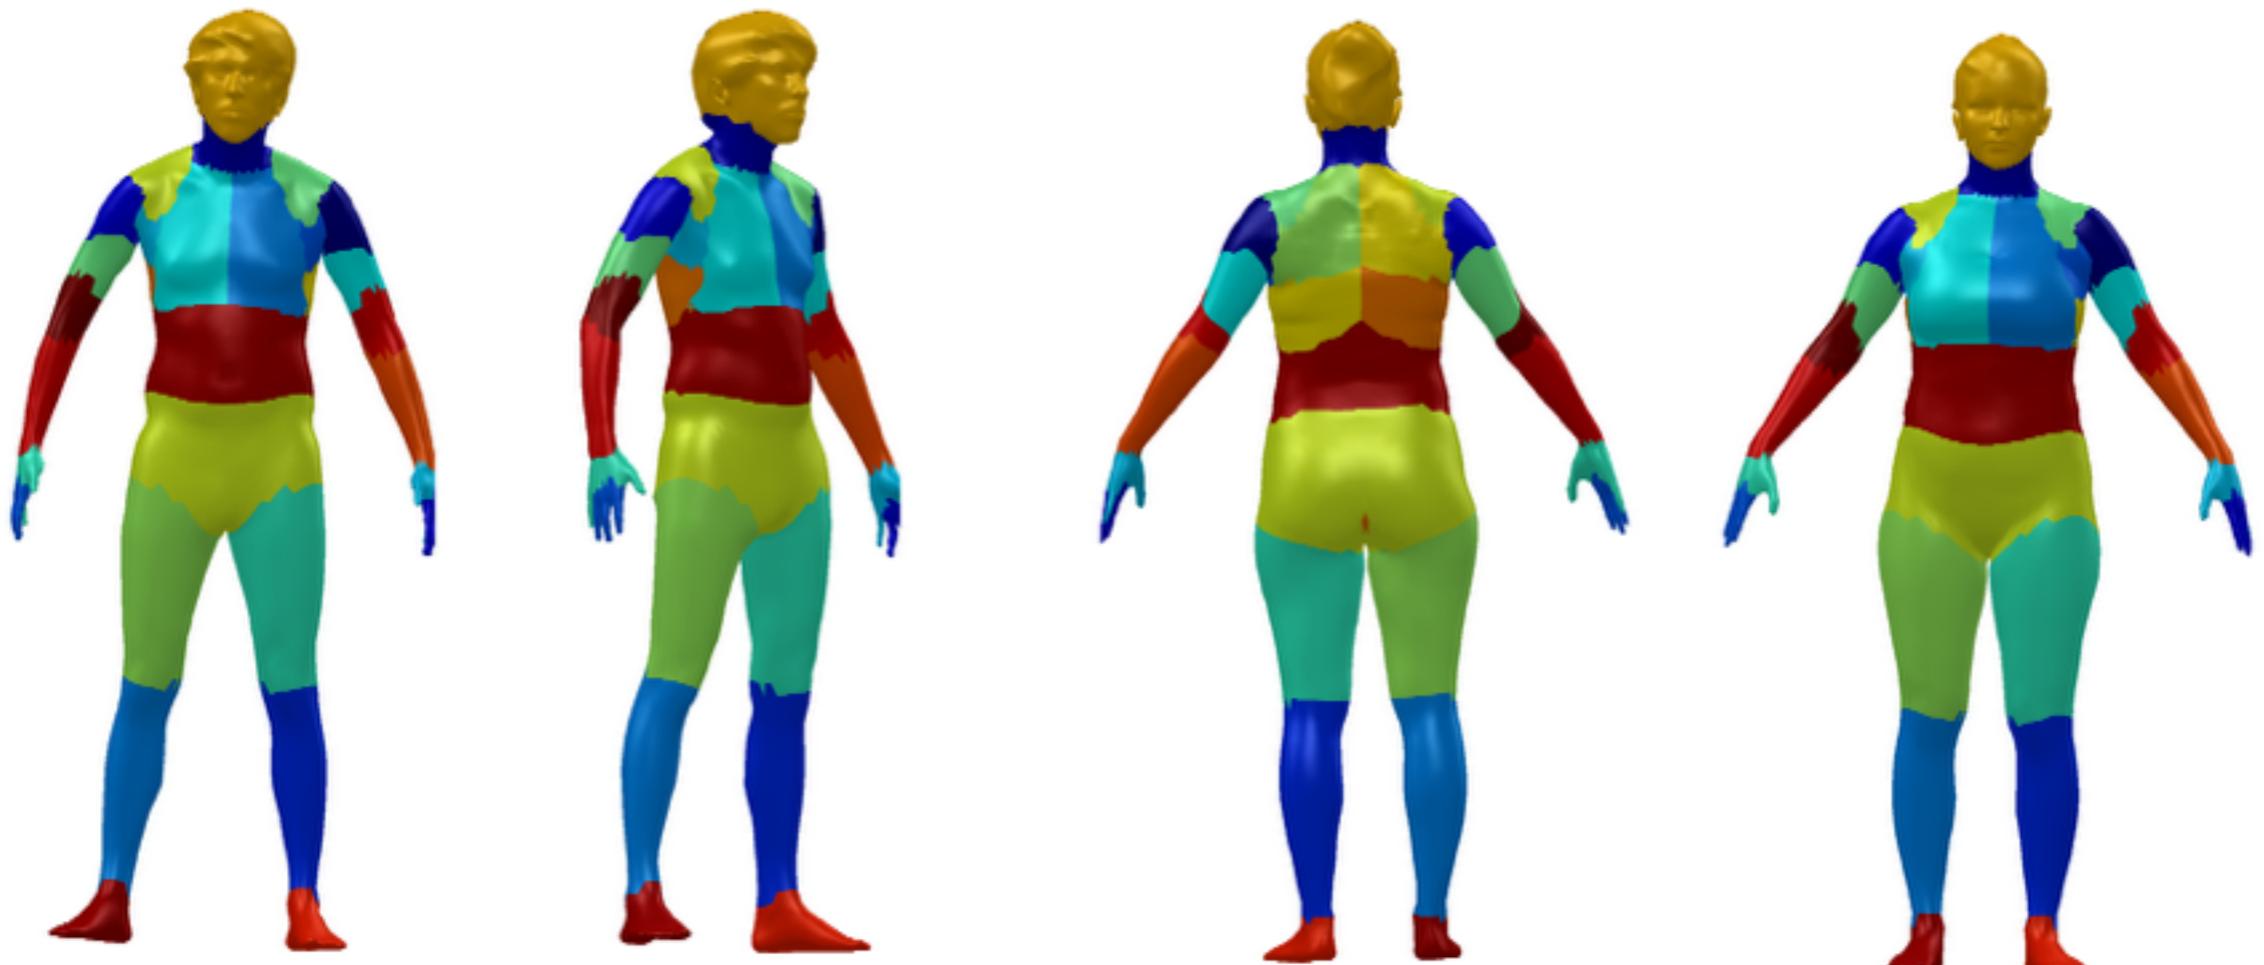
# Quantitative Evaluation



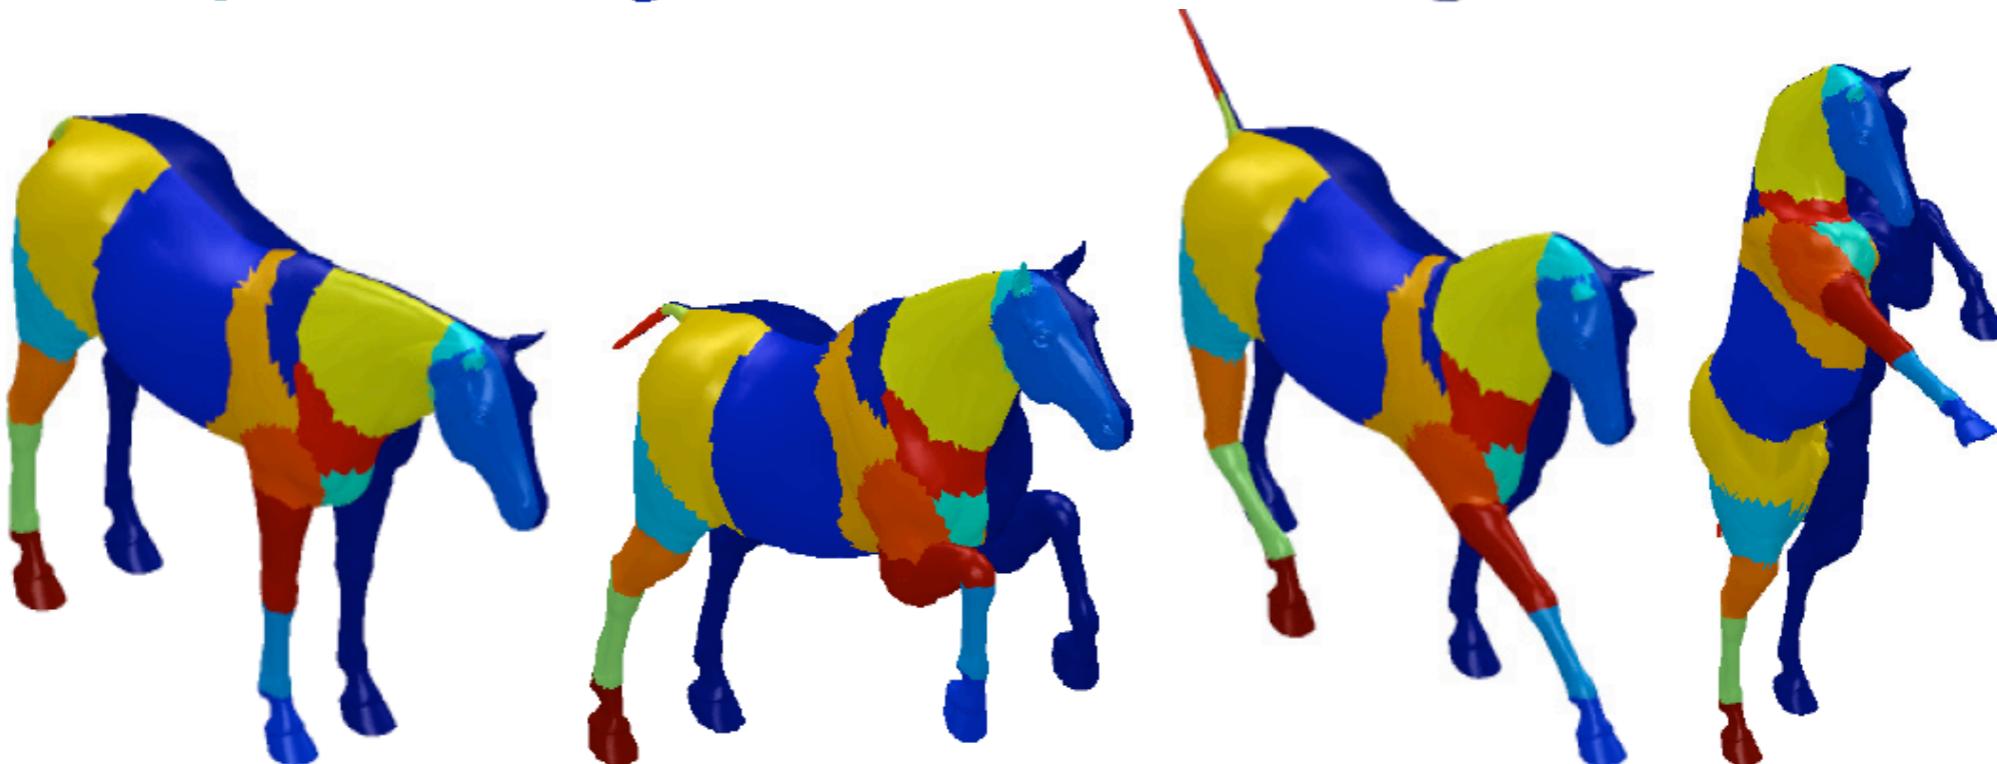
- Measure error in predicted motion for the candidate segmentations

# Larger Studies

- 1732 meshes, 78 subjects.
- Wide variability in **both** body **shapes** and **poses**.
- Similar results — part boundaries are cleaner.



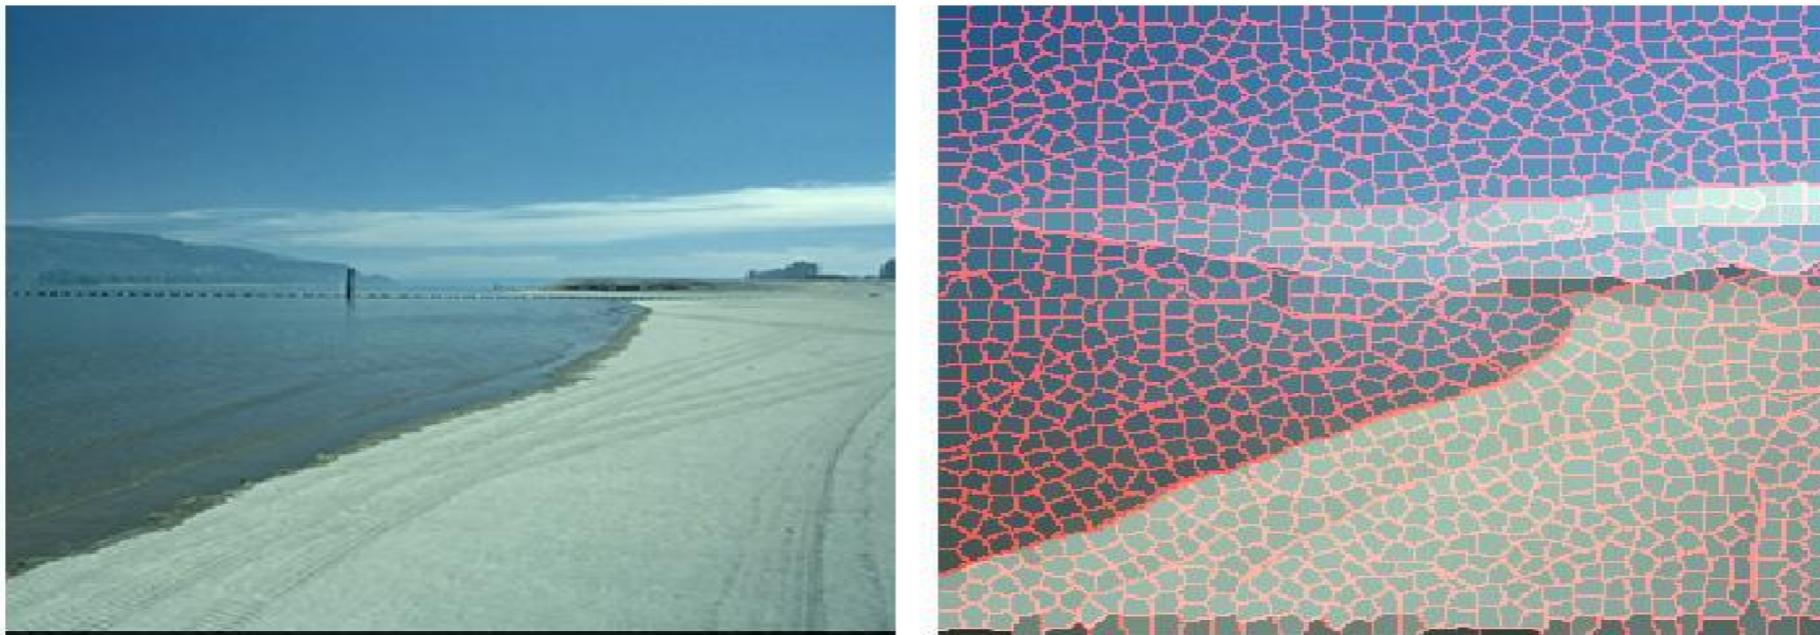
# Computer generated meshes



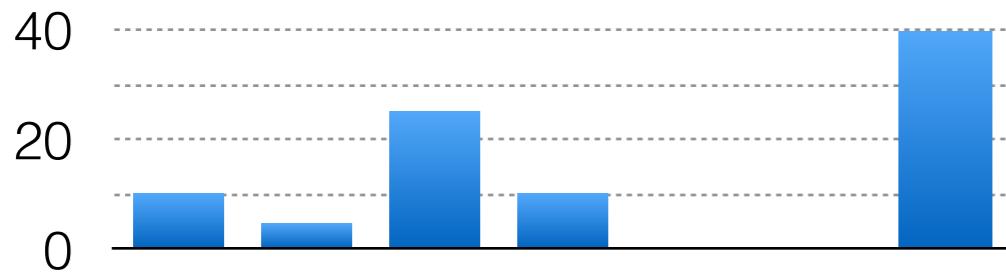
# Talk Outline

- Deformable 3D object segmentation
- Image and Video segmentation
- Layered decomposition of natural images
- Proposed Work

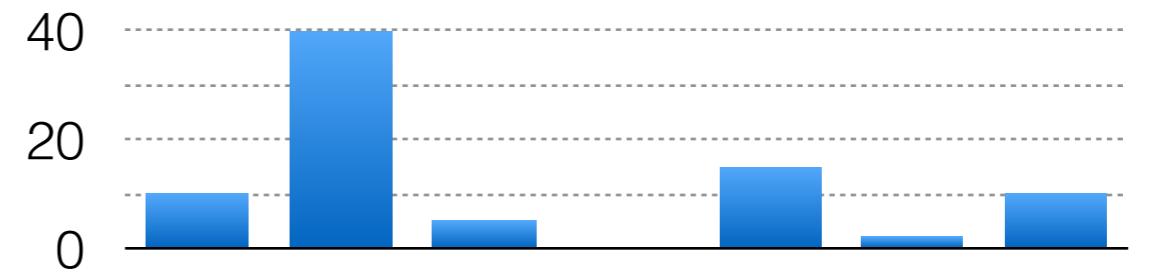
# Image Representation



- ~1000 super-pixels
- Each super-pixel is described through a histogram (~120 bin) of color and texture

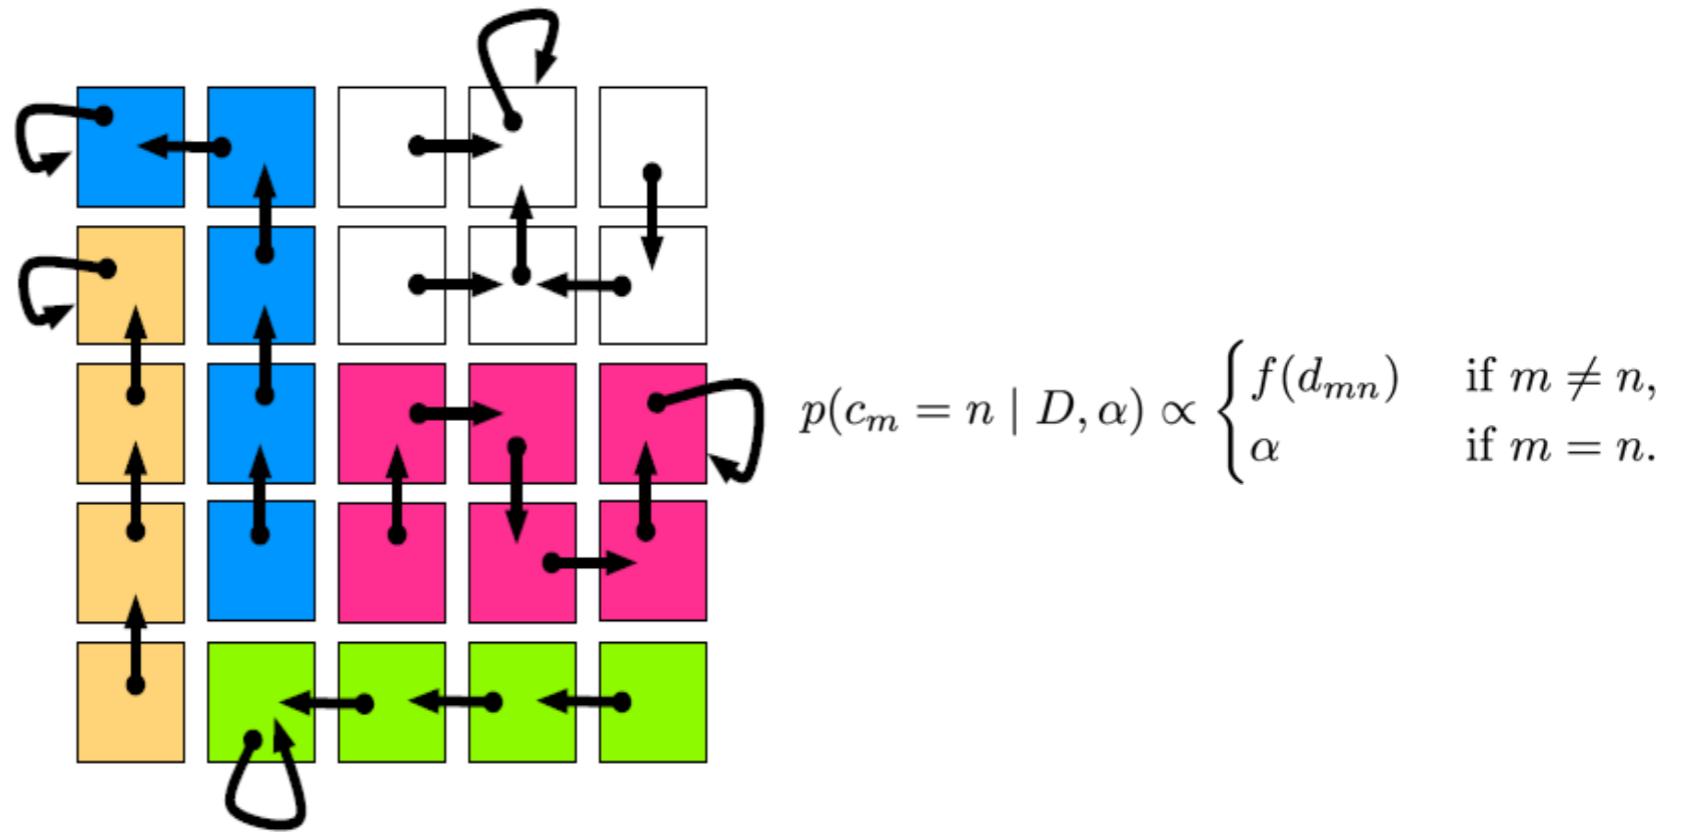


$$x_i^{color} \sim \text{Mult}(\theta^{color})$$



$$x_i^{texture} \sim \text{Mult}(\theta^{texture})$$

# Spatial ddCRP prior on Image Segmentations



*ddCRP1*

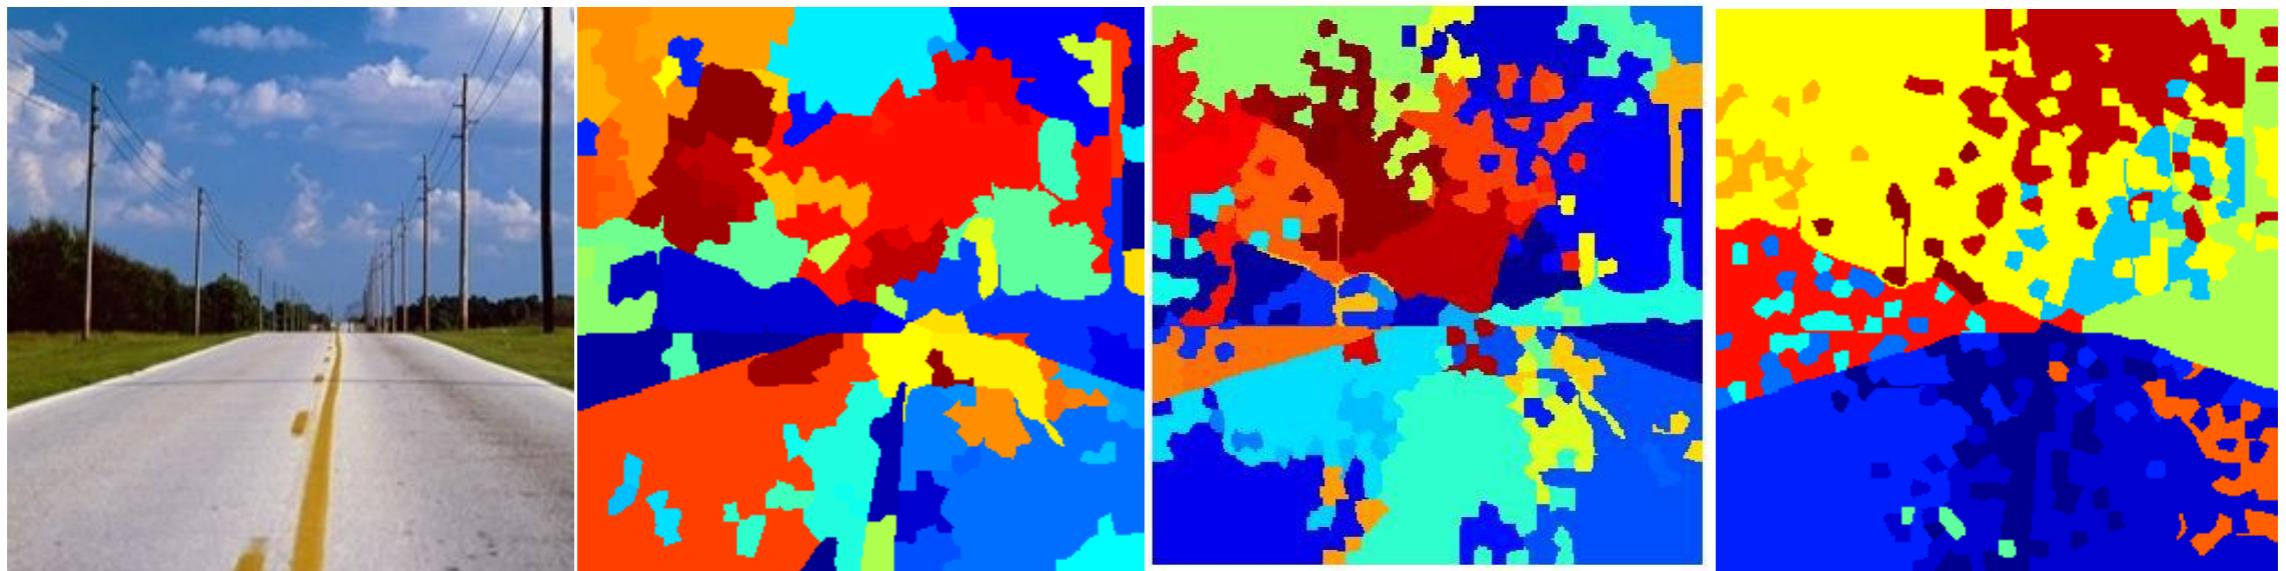
$$f(d_{mn}) = \mathbf{1}[d_{mn} \leq 1]$$

*ddCRP2*

$$f(d_{mn}) = \mathbf{1}[d_{mn} \leq 2]$$

*ddCRP5*

$$f(d_{mn}) = \mathbf{1}[d_{mn} \leq 5]$$

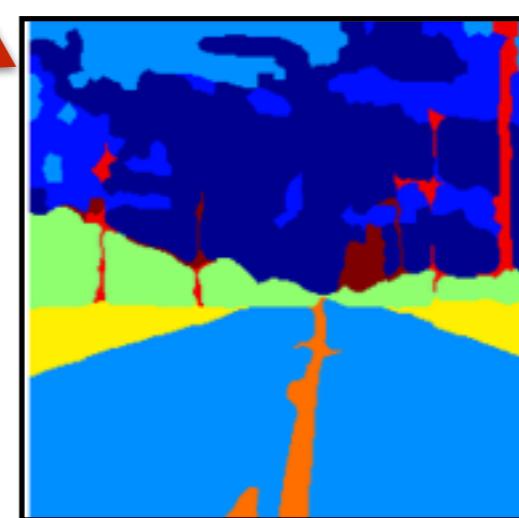
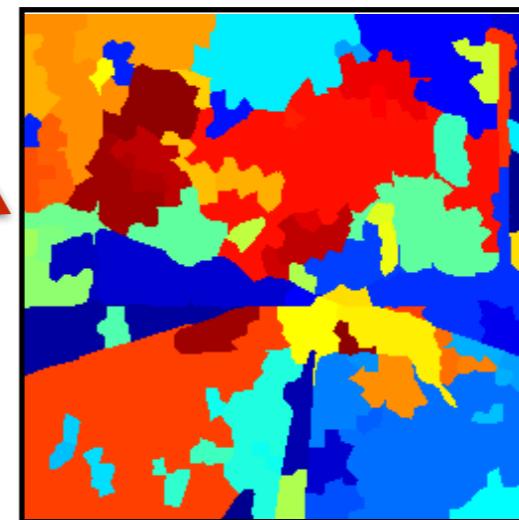
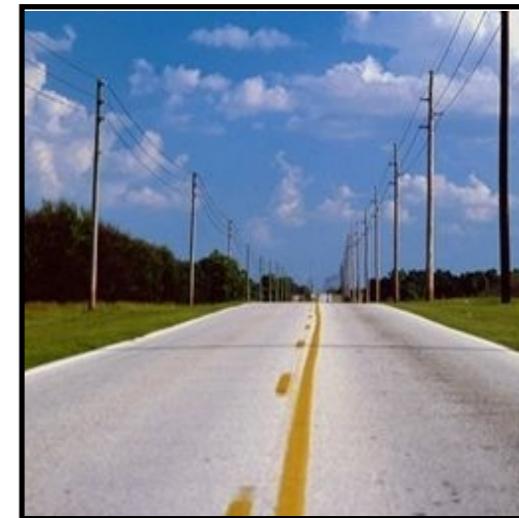


# Hierarchical Variant - “region” ddCRP

- Each super-pixel connects to a spatial neighbor. This produces an oversegmentation of the image.

$$c_i \sim \text{ddCRP}(\alpha, f, D)$$

- Group the segments into regions according to a Chinese restaurant process.
- Each region  $k$  maintains an appearance distribution parametrized by  $\phi_k$
- Color and texture histograms at super-pixels are then generated by sampling the appropriate region appearance distributions:  $x_i \sim p(x_i | \phi_{z_i})$



Ghosh et al.,  
NIPS 2011

# Inference - Gibbs Sampling

- Repeat till exhausted:
  - Sample super-pixel links (identical to ddCRP):
$$c_i \sim p(c_i \mid c_{-i}, k, x)$$
  - These define an over segmentation of the image.
  - *Extra wrinkle*: sample region assignments for each segment.

$$k_t \sim p(k_t \mid c, k_{-t}, x)$$

# Eight Natural Scene Category Dataset (LabelMe)



*800 images in all*

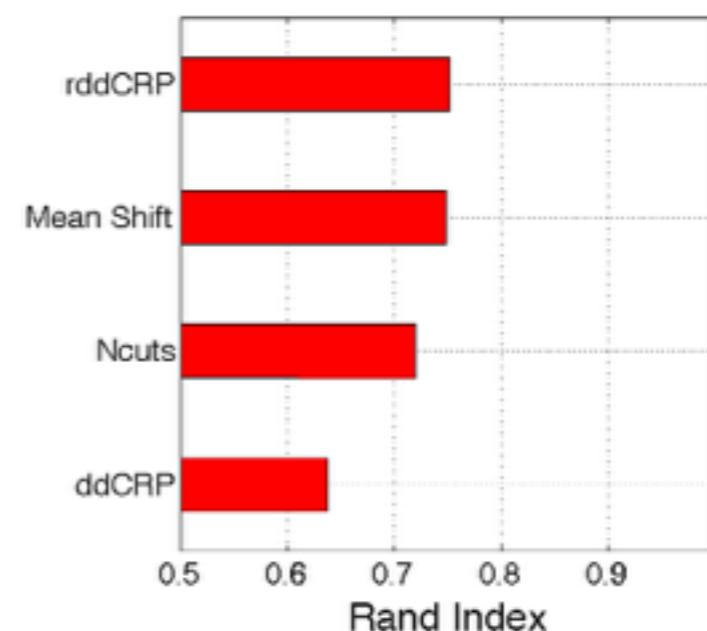
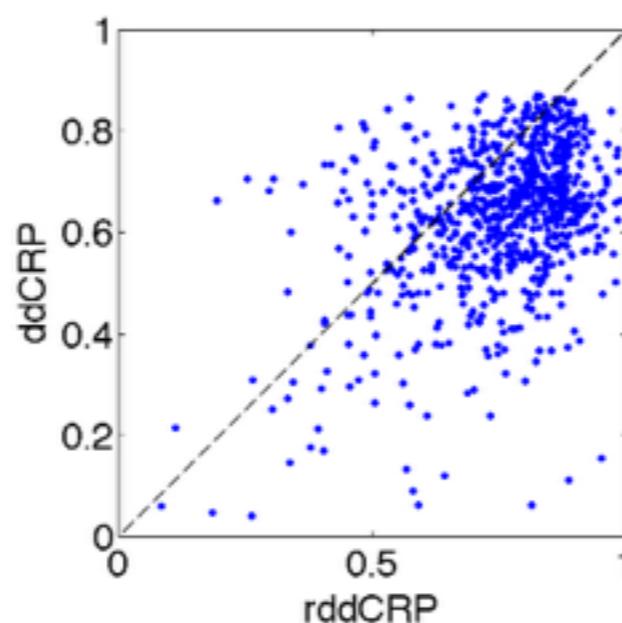
*Oliva and Torralba 2001*

# Segmentation Results



## Rand Index

Measures agreement  
with human segmentations



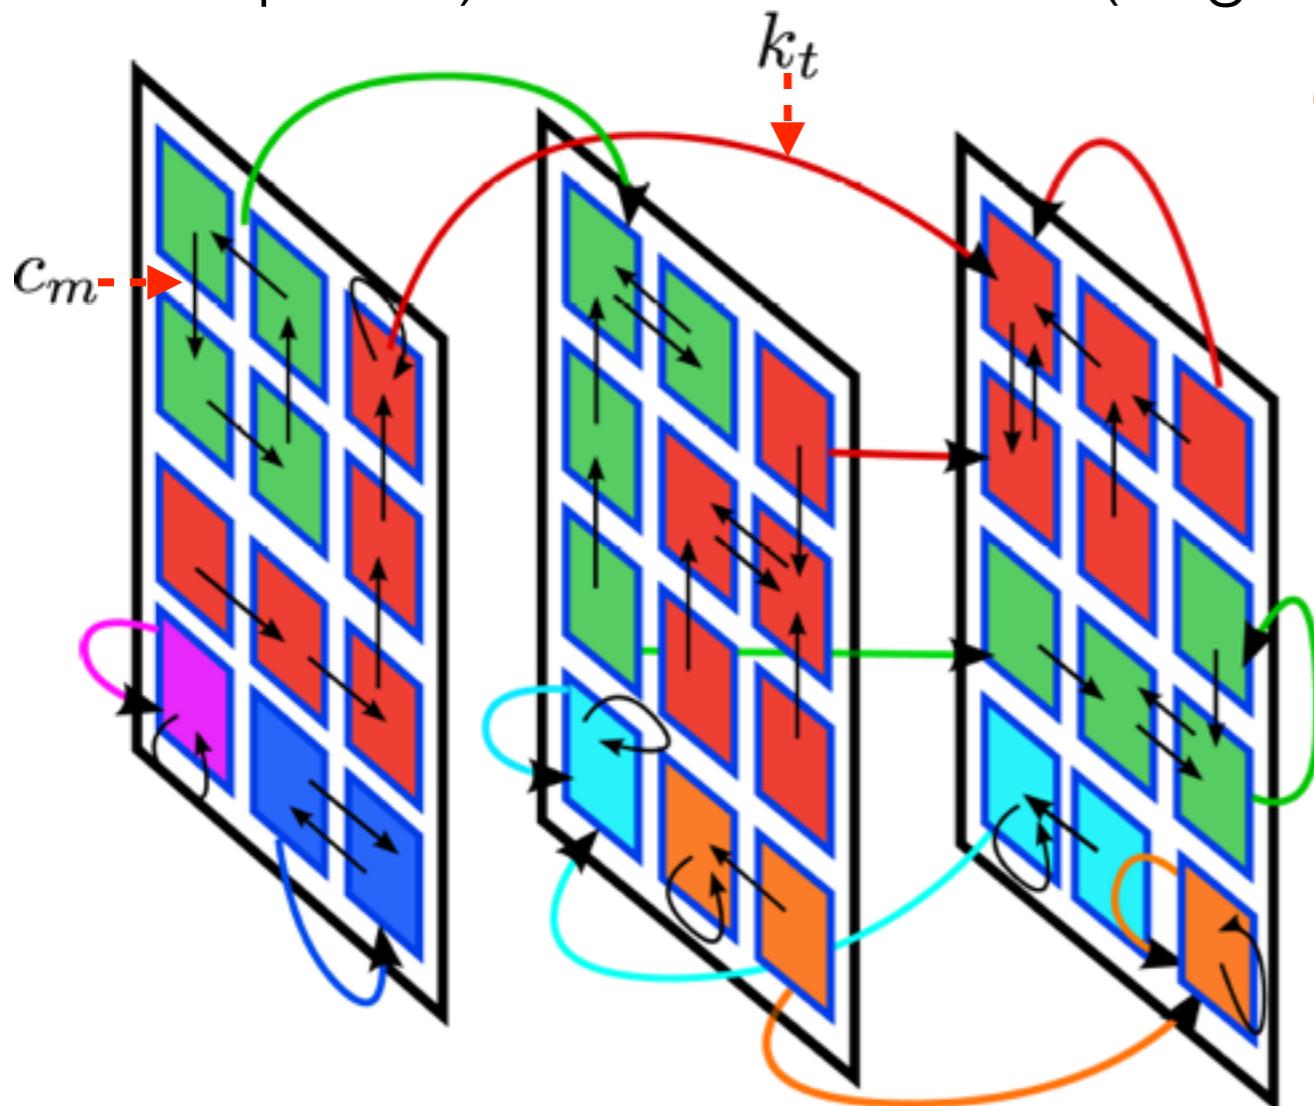
# Video Segmentation

- “region” ddCRP models spatial correlations between pixels
- But tables are grouped with a CRP— Inter table correlations are not modeled



# Generalized Hierarchical ddCRPs

- Model for grouped data.
- Key Idea: Allow distances between both customers (super-pixels) and *latent* tables (segments).



**Group indexed ( $j$ ) ddCRPs :**

$$p(c_{jm} | D_j, f_j, \alpha_j) \propto \begin{cases} f_j(d_j(mn)) & \text{if } jm \neq jn \\ \alpha_j & \text{if } jm = jn \end{cases}$$

**Tables are coupled using a inter-table ddCRP :**

$$p(k_t = t' | D_0(c), f_0, \alpha_0) \propto \begin{cases} f_0(d_0(t, t', c)) & \text{if } t \neq t' \\ \alpha_0 & \text{if } t = t' \end{cases}$$

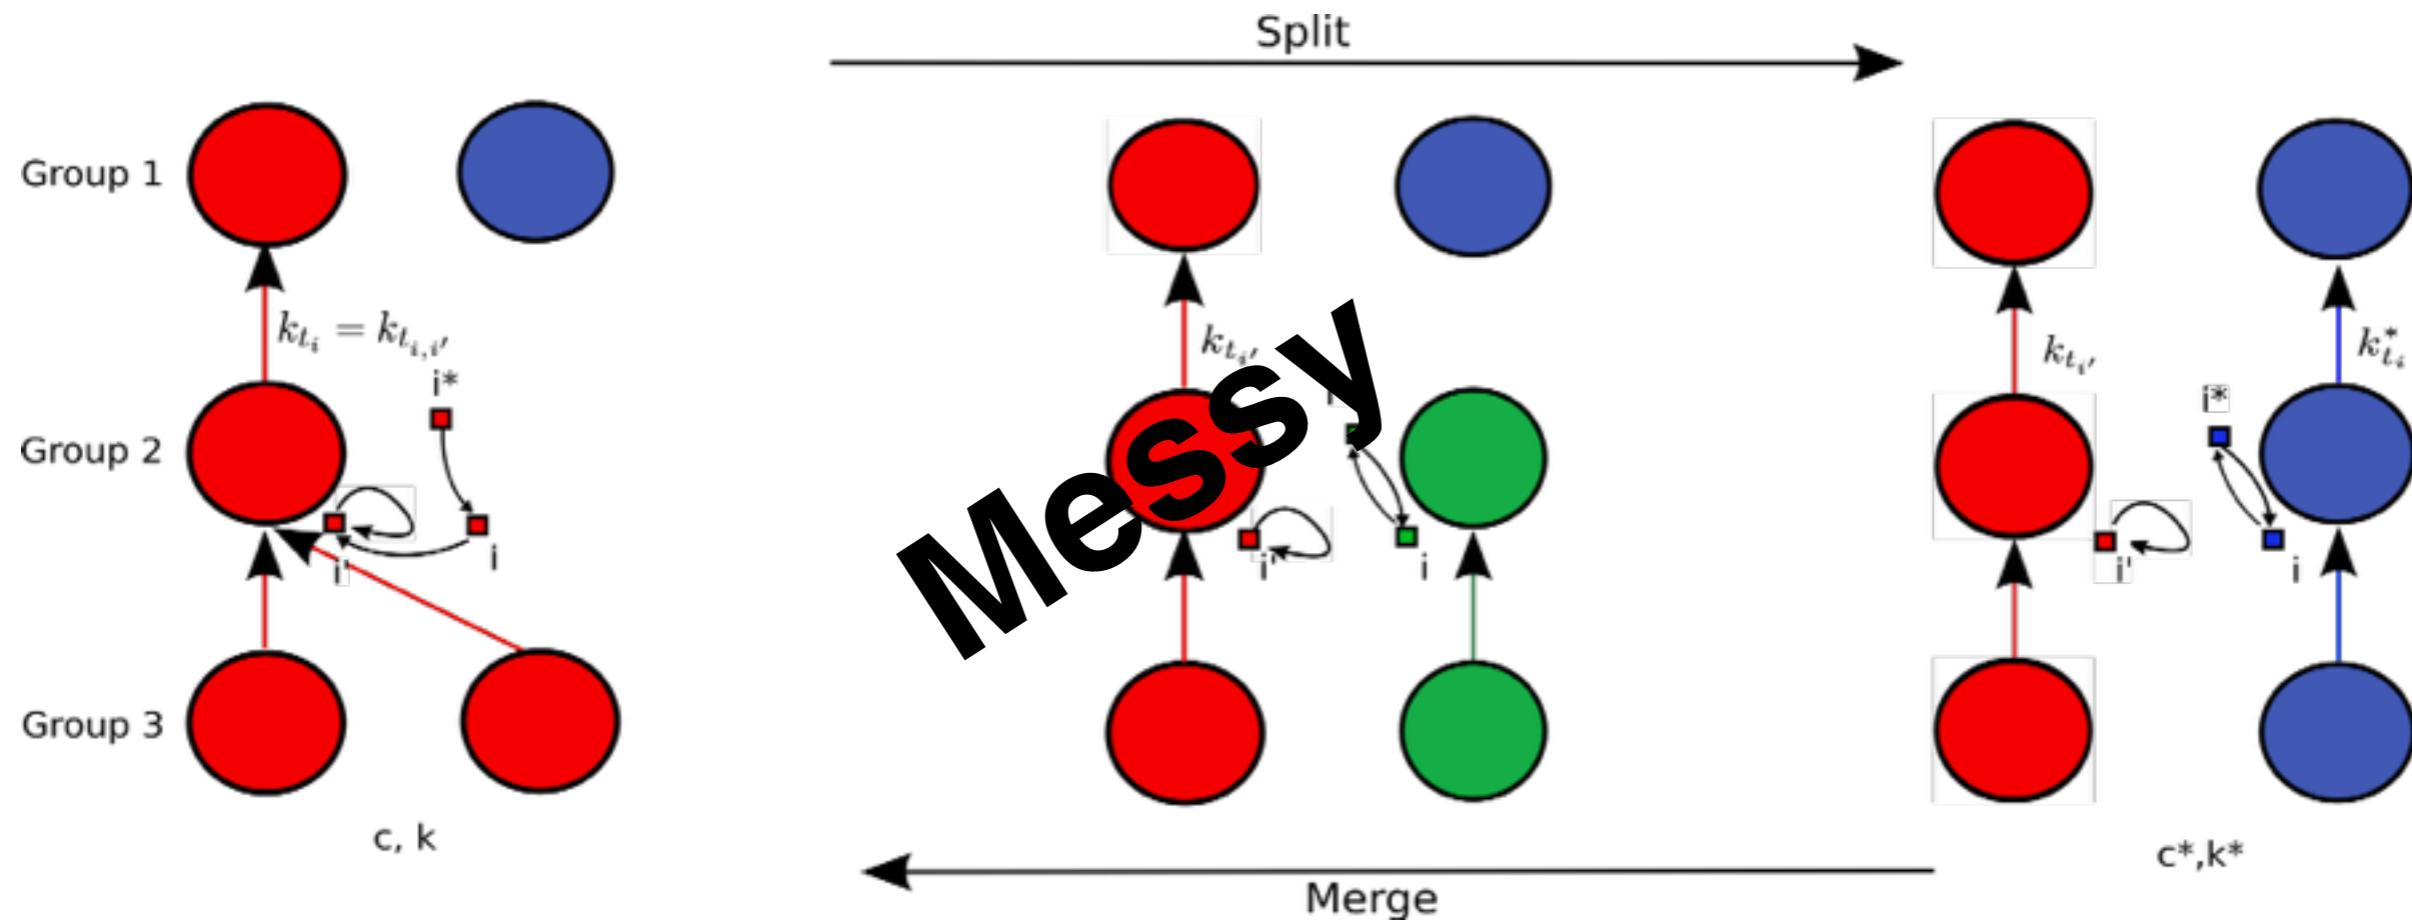
- Table distances could be functions of:
- i) table sizes
  - ii) table shapes
  - iii) table locations
  - iv) table neighborhoods

# Inference - Metropolis Hastings

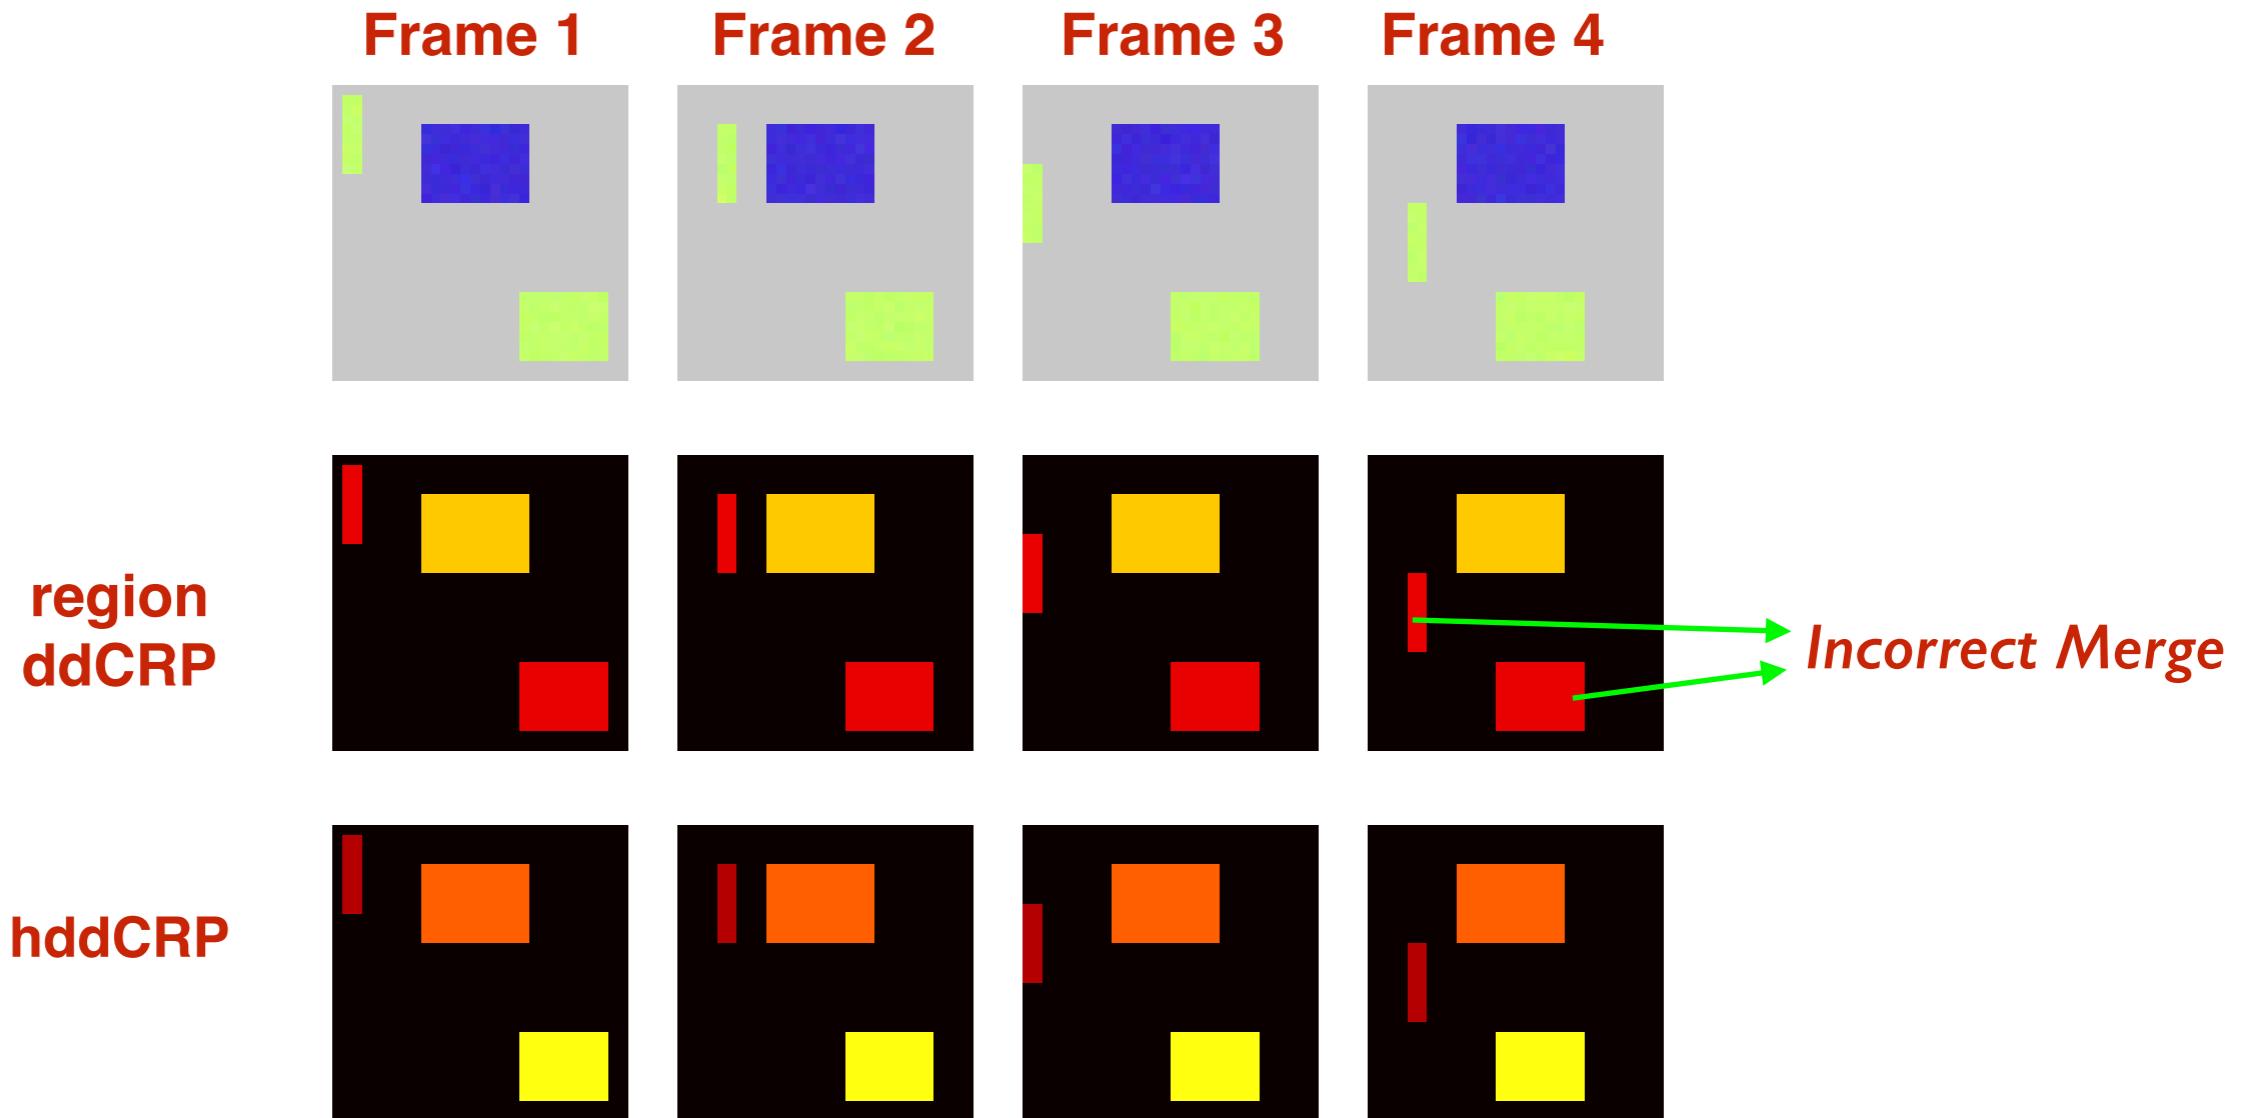
- Composite moves — make coordinated changes to customer and table links.

$$q(c_i^*) \propto p(c_i^* | \alpha, D) \Gamma(\mathbf{X}, \mathbf{z}, \lambda),$$

$$\Gamma(\mathbf{X}, \mathbf{z}, \lambda) = \begin{cases} \frac{p(\mathbf{X}_{\mathbf{z}(\Delta)=m_a} \cup \mathbf{X}_{\mathbf{z}(\Delta)=m_b} | \lambda)}{p(\mathbf{X}_{\mathbf{z}(\Delta)=m_a} | \lambda)p(\mathbf{X}_{\mathbf{z}(\Delta)=m_b} | \lambda)} & \text{if } c_i^* \text{ merges dishes } m_a \text{ and } m_b \\ 1 & \text{otherwise.} \end{cases}$$



# Video Segmentation-Toy Data



Size based inter-table distances.

*Groups those tables (segments) together which in addition to having *similar appearance* have *similar sizes*.*

# Video Segmentation Results

Frame 1



Frame 5



Frame 10



HGVS - Grundman'10

*Rand Index - 0.49*

# Video Segmentation Results

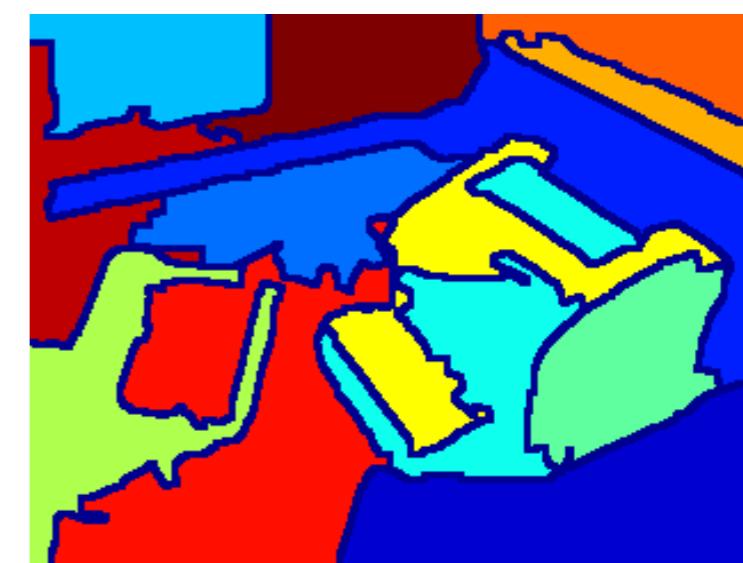
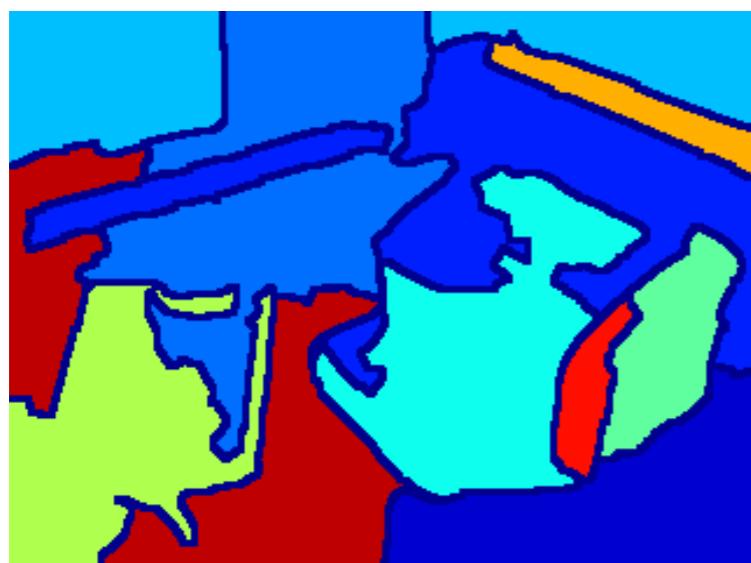
Frame 1



Frame 5



Frame 10



region ddCRP

*Rand Index - 0.53*

# Video Segmentation Results

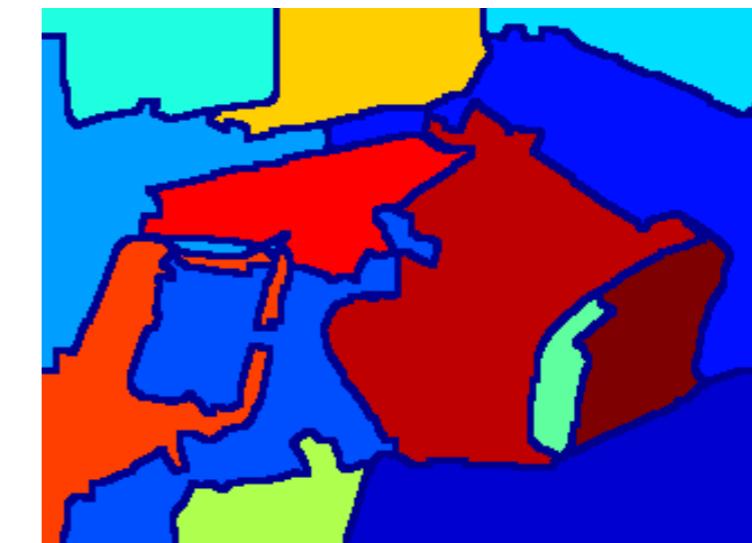
Frame 1



Frame 5



Frame 10

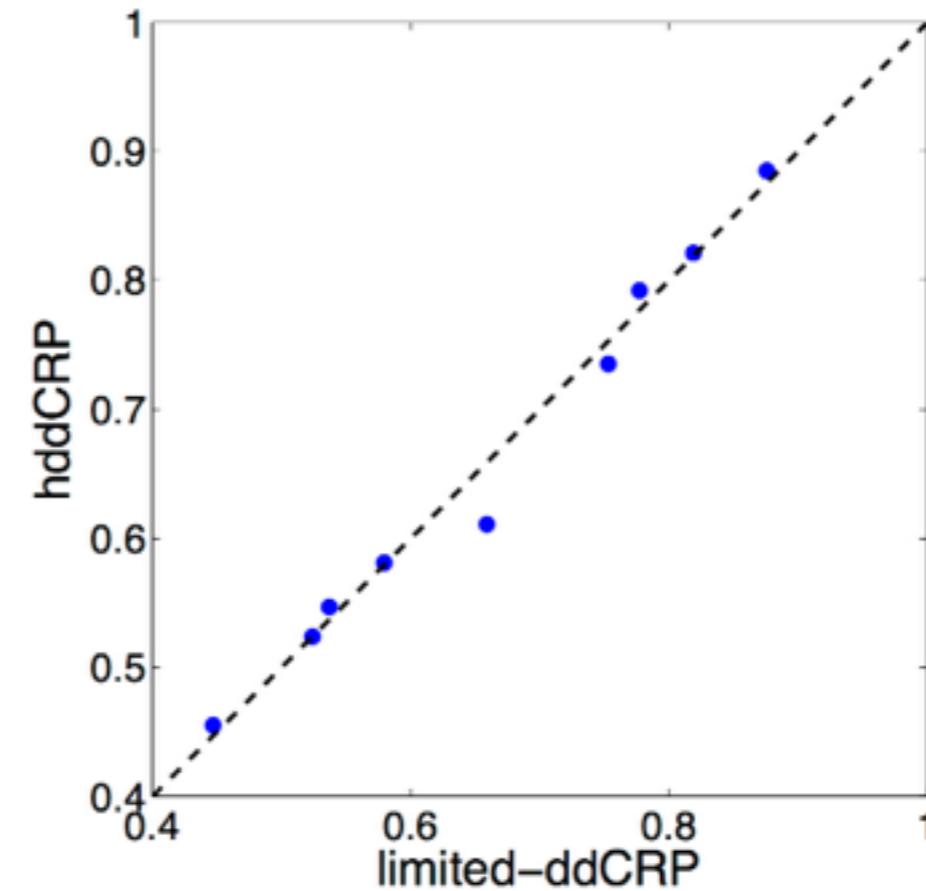
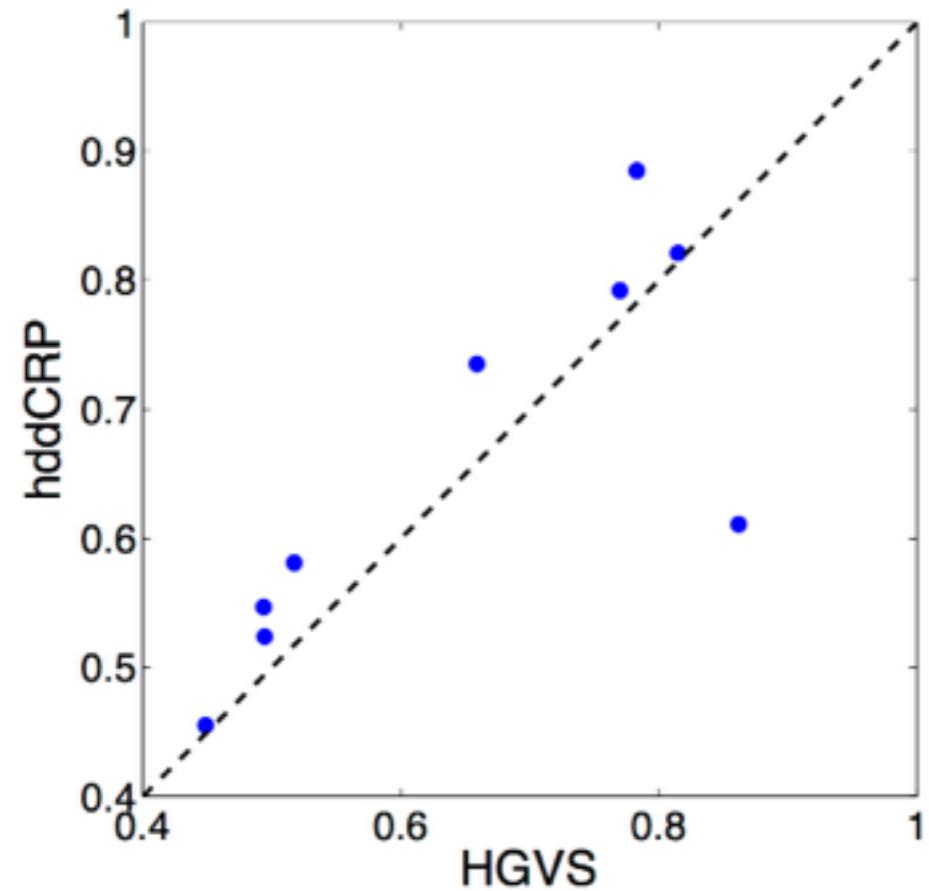


**hddCRP**

*Rand Index - 0.55*

# Quantitative Performance

- Nine human segmented video sequences. (Liu et al., 08)



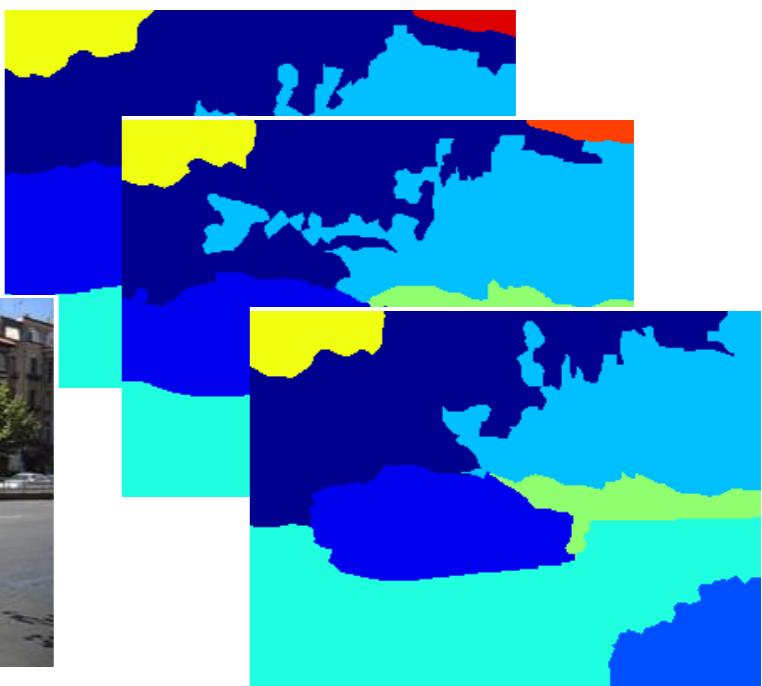
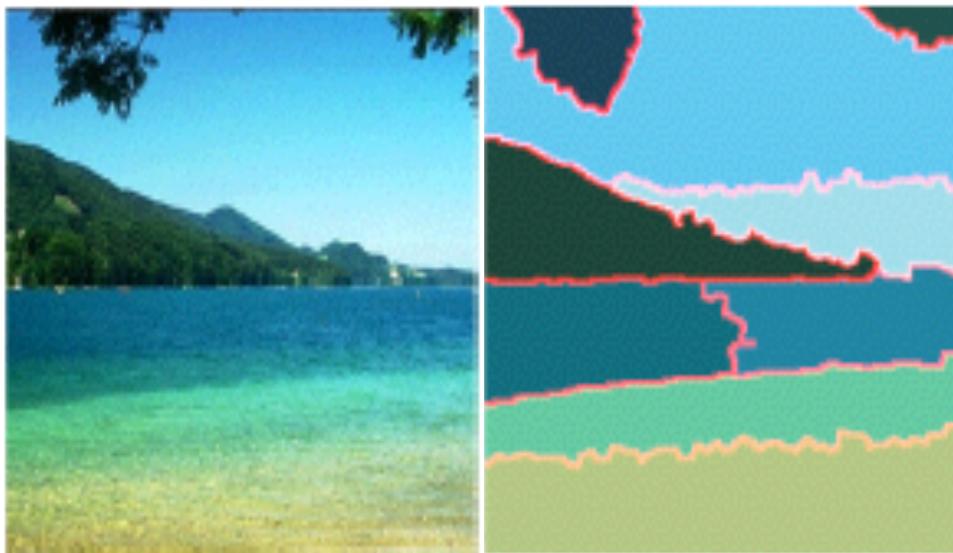
**Rand Index – Higher = Better**

- **Proposed Work:**

- Explore whether the modest improvement (over rddCRP) is an inference problem or a modeling issue.
- Use fancier inter-table distances.

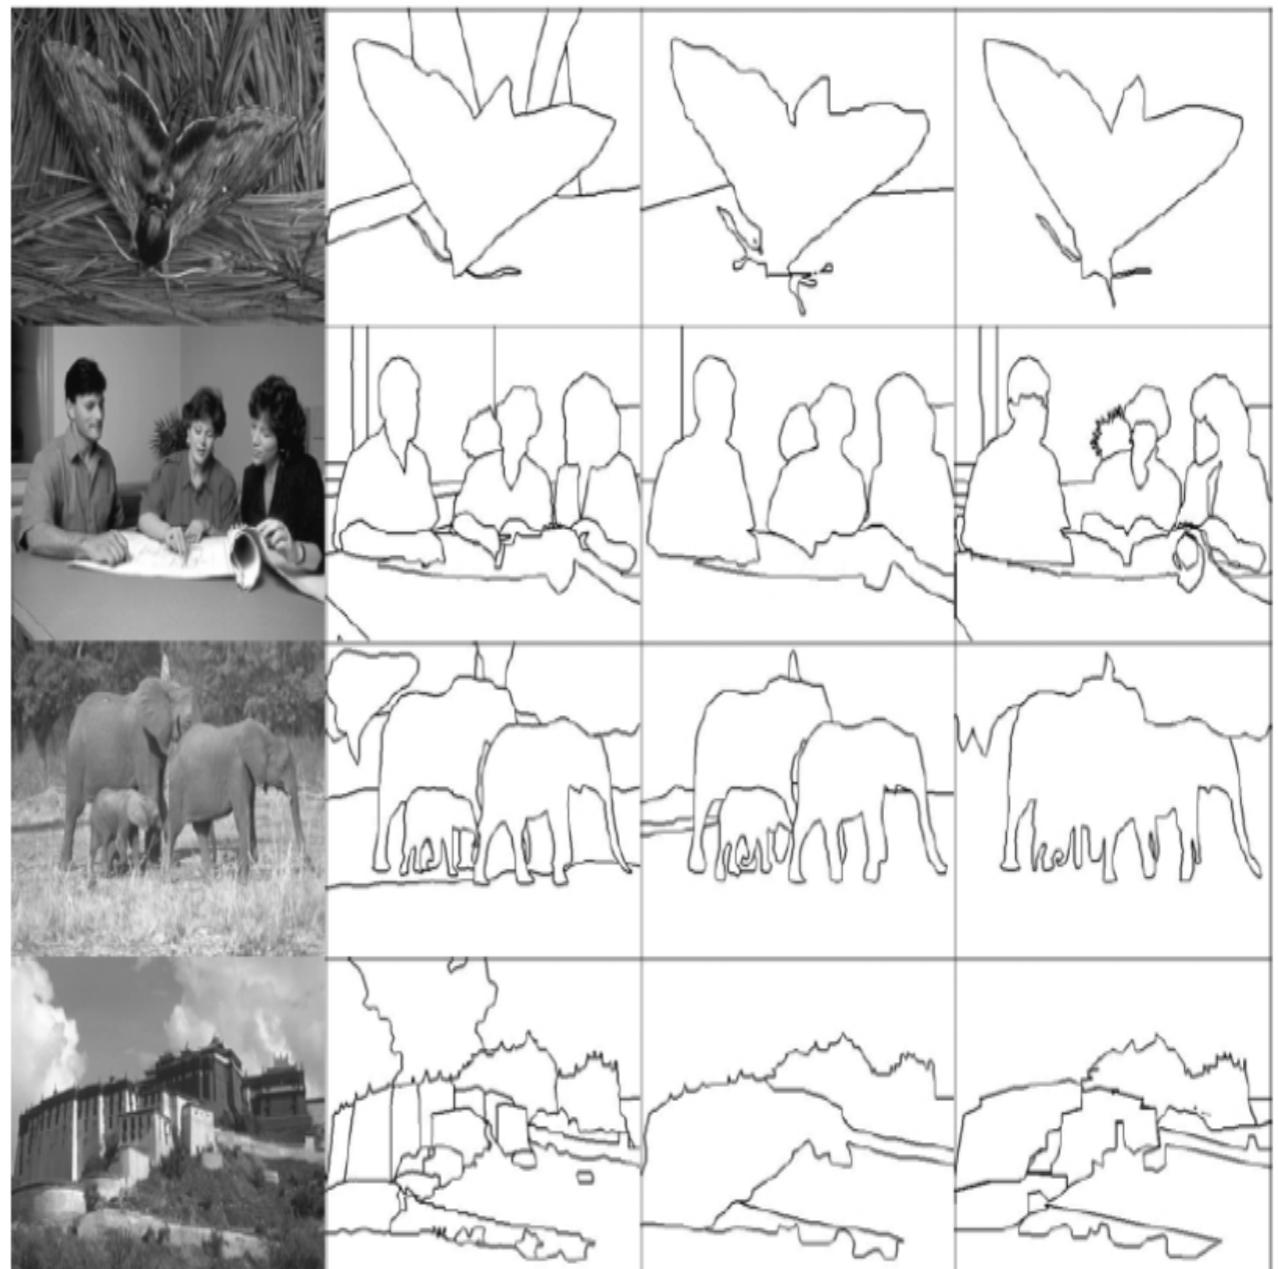
# Summary

- ddCRP mixture of local regression models for articulated object segmentation
- “region” ddCRP and the generalized hierarchical ddCRP for image and video segmentation

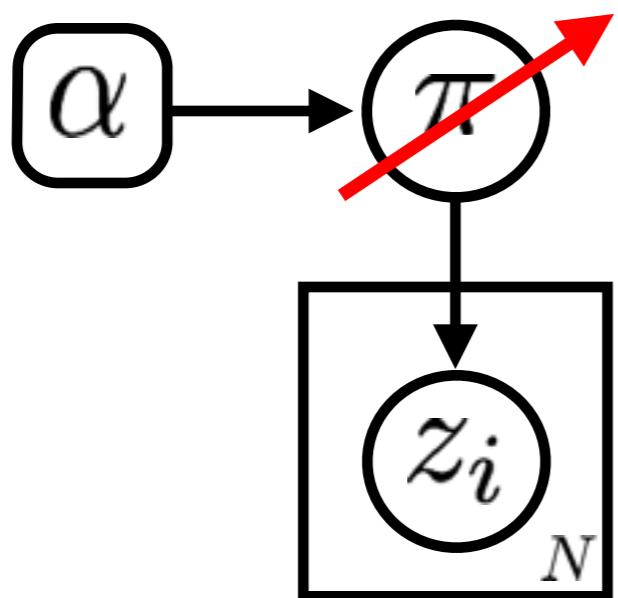


# Statistics of Human Segments

- Human segment sizes follow power law behavior.



# Influencing Segment Sizes



$$\pi = \{\pi_1, \pi_2, \dots, \pi_\infty\}$$

$$\sum_{k=1}^{\infty} \pi_k = 1 \quad 0 \leq \pi_k \leq 1$$

$$z_i \sim \pi$$

**Distribution over partitions - Generalized CRP**

$$p(Z \mid \alpha) = \int p(\pi \mid \alpha) \prod_i p(z_i \mid \pi) d\pi$$

# Pitman-Yor Process

- The Pitman-Yor process defines a distribution on infinite discrete measures, or partitions

$$G(\theta) = \sum_{k=1}^{\infty} \pi_k \delta_{\theta_k}(\theta) \sim \text{PY}(\alpha, H)$$

$$\theta_k \sim H(\lambda)$$

## Power Law Behavior:

*Expected size of sorted component k:*

$$E[\pi_k] \propto O(k^{-\frac{1}{\alpha_a}})$$

- Use as a distribution over segment sizes.

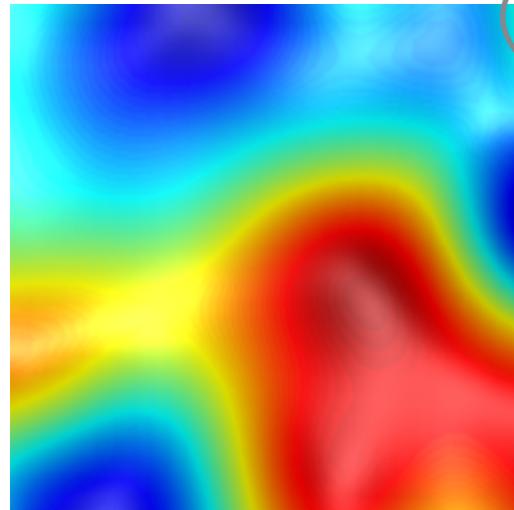
# Talk Outline

- Deformable 3D object segmentation
- Image and Video segmentation
- Layered decomposition of natural images
- Proposed Work

# Spatial Coupling through Layers

## Smooth Layers

$$u_1 \sim \text{GP}(0, K)$$

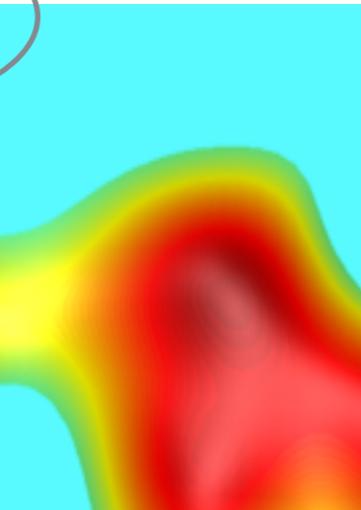


$$\begin{aligned} w_1 &\sim \text{Beta}(1 - \alpha_a, \alpha_b + \alpha_a) \\ \delta_1 &= \Phi^{-1}(w_1) \end{aligned}$$

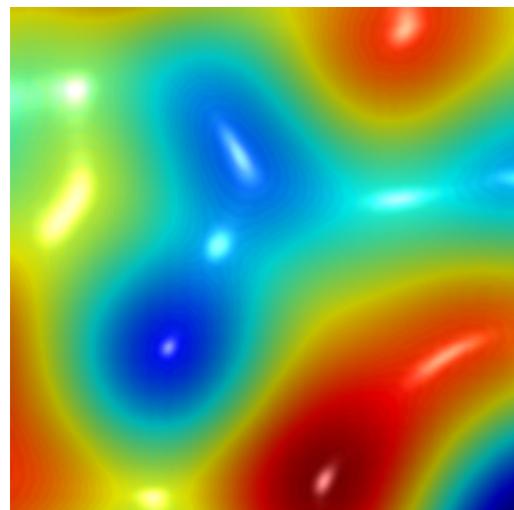
$$u_1 < \delta_1$$



## Thresholded layer support

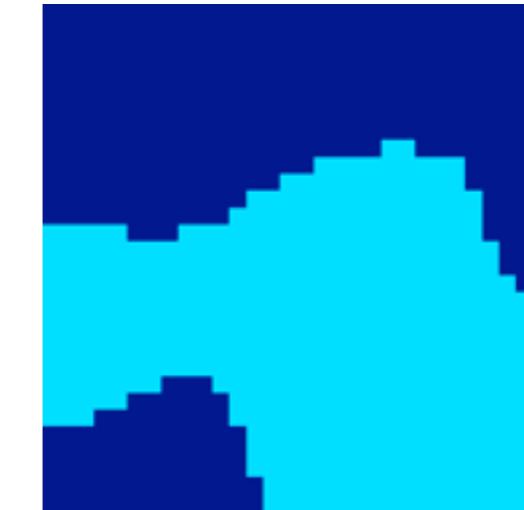


$$u_2 \sim \text{GP}(0, K)$$



$$\begin{aligned} w_2 &\sim \text{Beta}(1 - \alpha_a, \alpha_b + 2\alpha_a) \\ \delta_2 &= \Phi^{-1}(w_2) \end{aligned}$$

$$u_2 < \delta_2$$



Occlude



## Image Partition

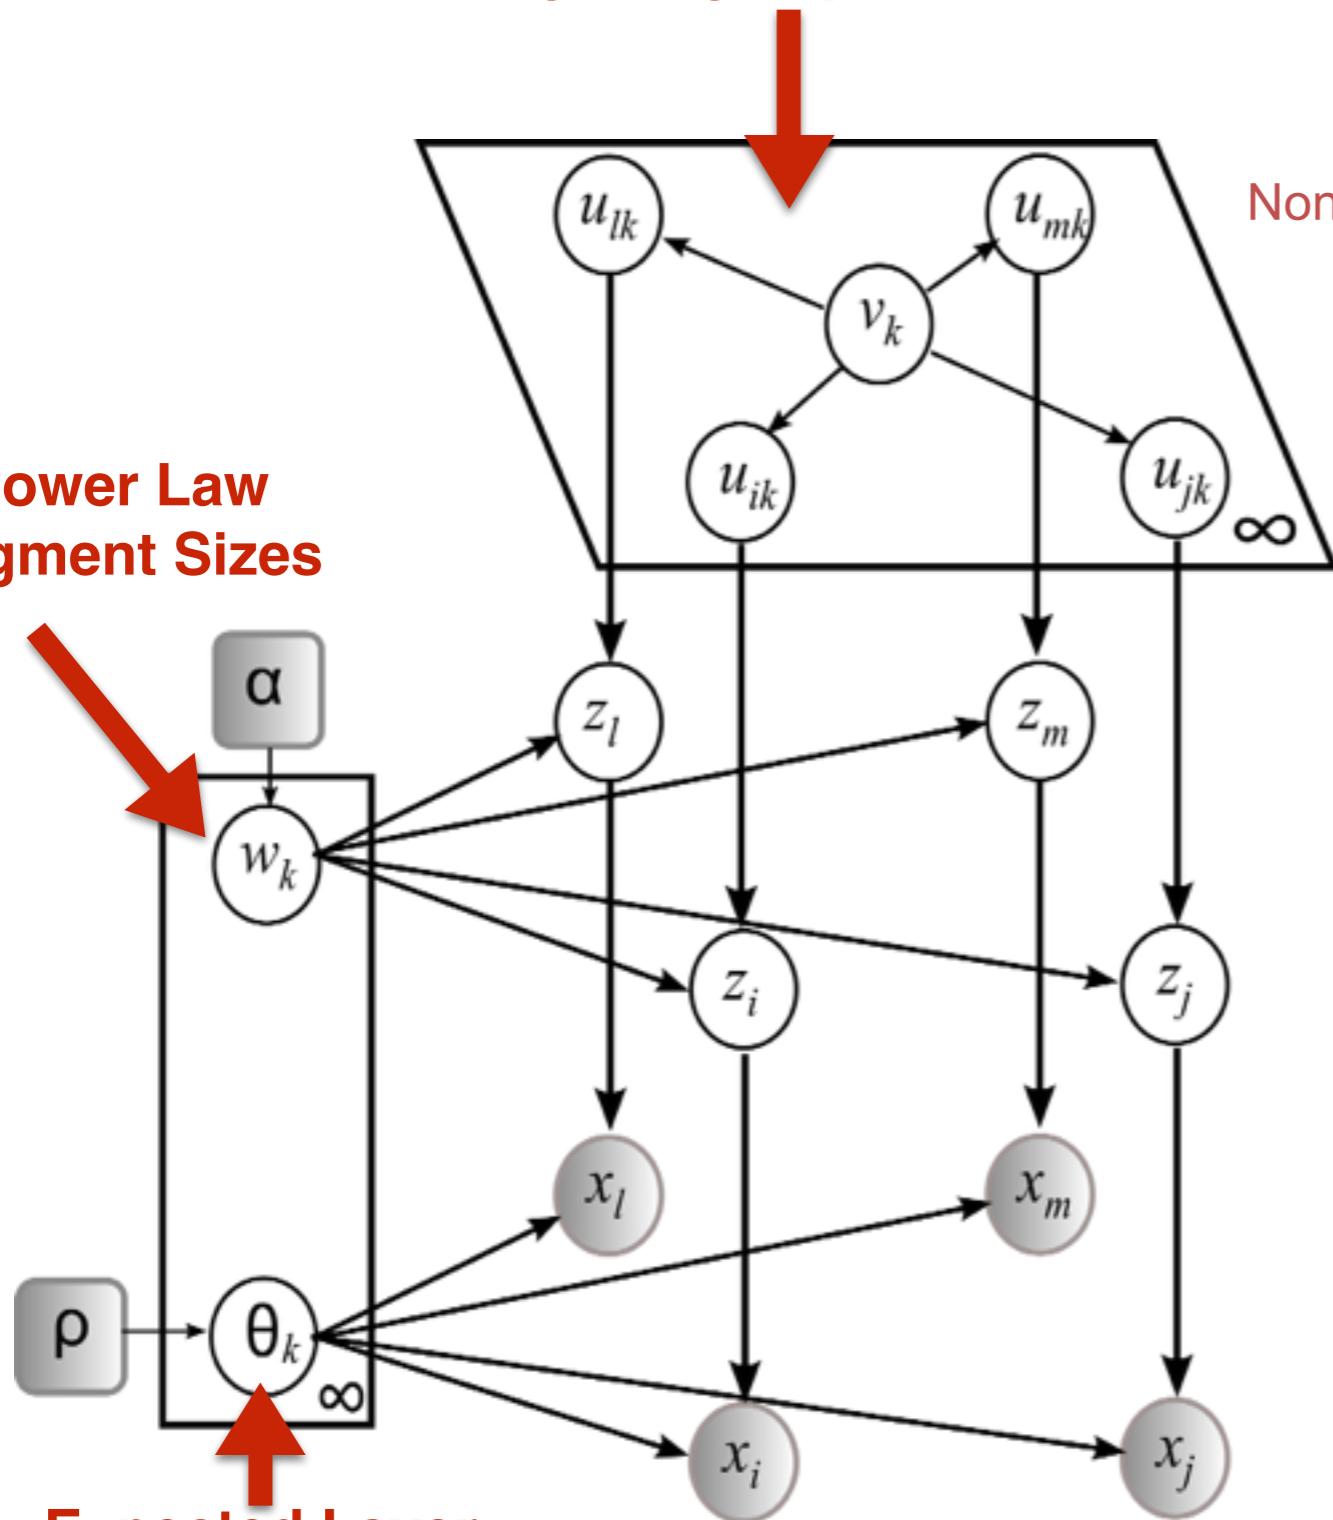


$$z_n = \min\{k \mid u_{kn} < \delta_k\}$$

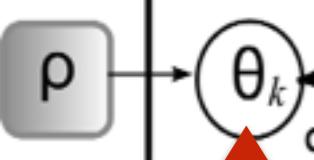
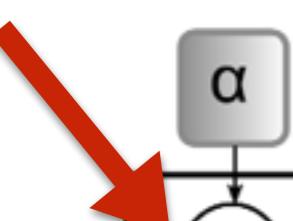
*Sudderth & Jordan, 2008  
Ghosh & Sudderth, 2012*

# Spatially Coupled PY Processes

**Model Long Range Spatial Correlations**



**Power Law Segment Sizes**



**Expected Layer Appearance**

Non-Markov Gaussian Processes: Spatial Dependence

$$\mathbf{u}_k = A\mathbf{v}_k + \boldsymbol{\epsilon}_k \quad \mathbf{v}_k \sim \mathcal{N}(0, \mathbf{I}_D)$$

$$u_{kn} \sim \mathcal{N}(0, 1) \quad \boldsymbol{\epsilon}_k \sim \mathcal{N}(0, \boldsymbol{\Psi})$$

Pitman-Yor prior: Segment sizes follow a *power law* distribution

$$\Pr[z_n = k] = w_k \prod_{\ell=1}^{k-1} (1 - w_\ell)$$

$$w_k \sim \text{Beta}(1 - \alpha_a, \alpha_b + k\alpha_a)$$

Segment Assignments: Layers

$$z_n = \min\{k \mid u_{kn} < \Phi^{-1}(w_k)\}$$

Features: Color & Texture

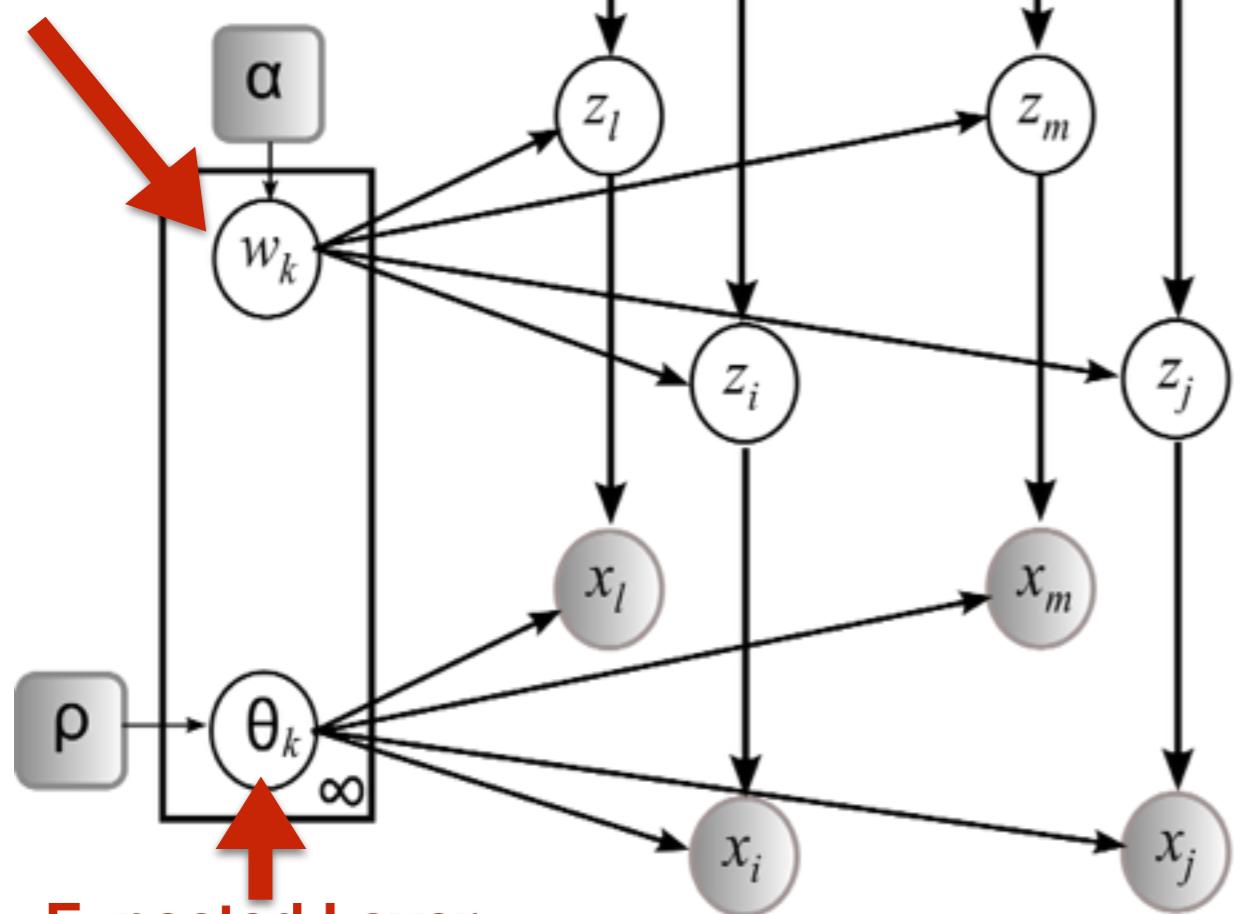
$$x_n^c \sim \text{Mult}(\theta_{z_n}^c)$$

$$x_n^t \sim \text{Mult}(\theta_{z_n}^t)$$

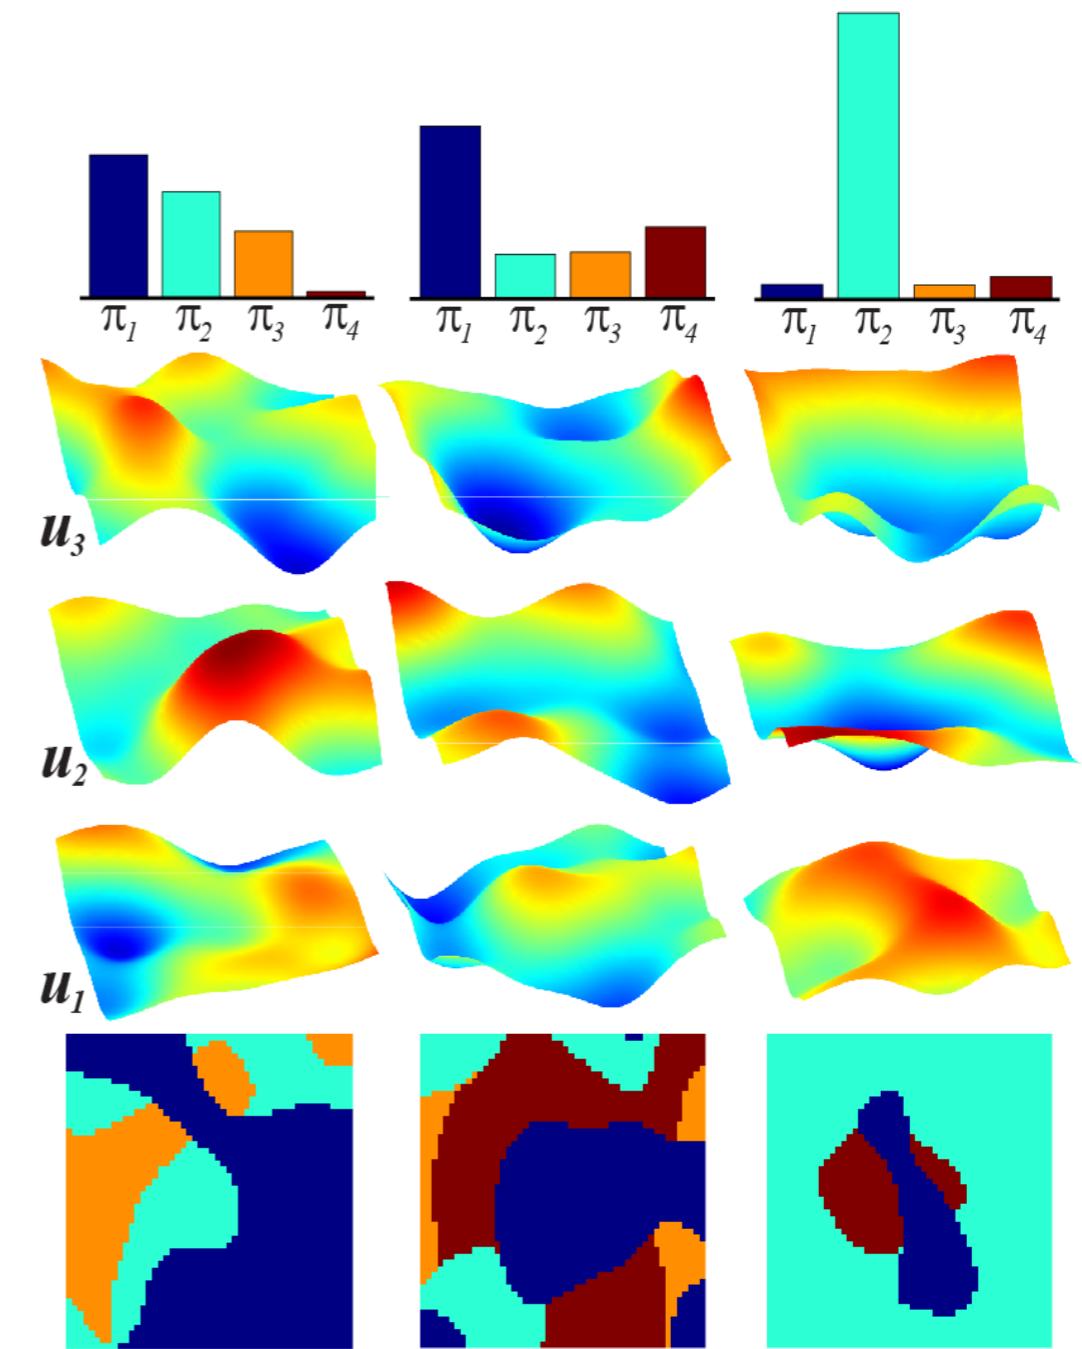
# Spatially Coupled PY Processes

Model Long Range Spatial Correlations

Power Law Segment Sizes

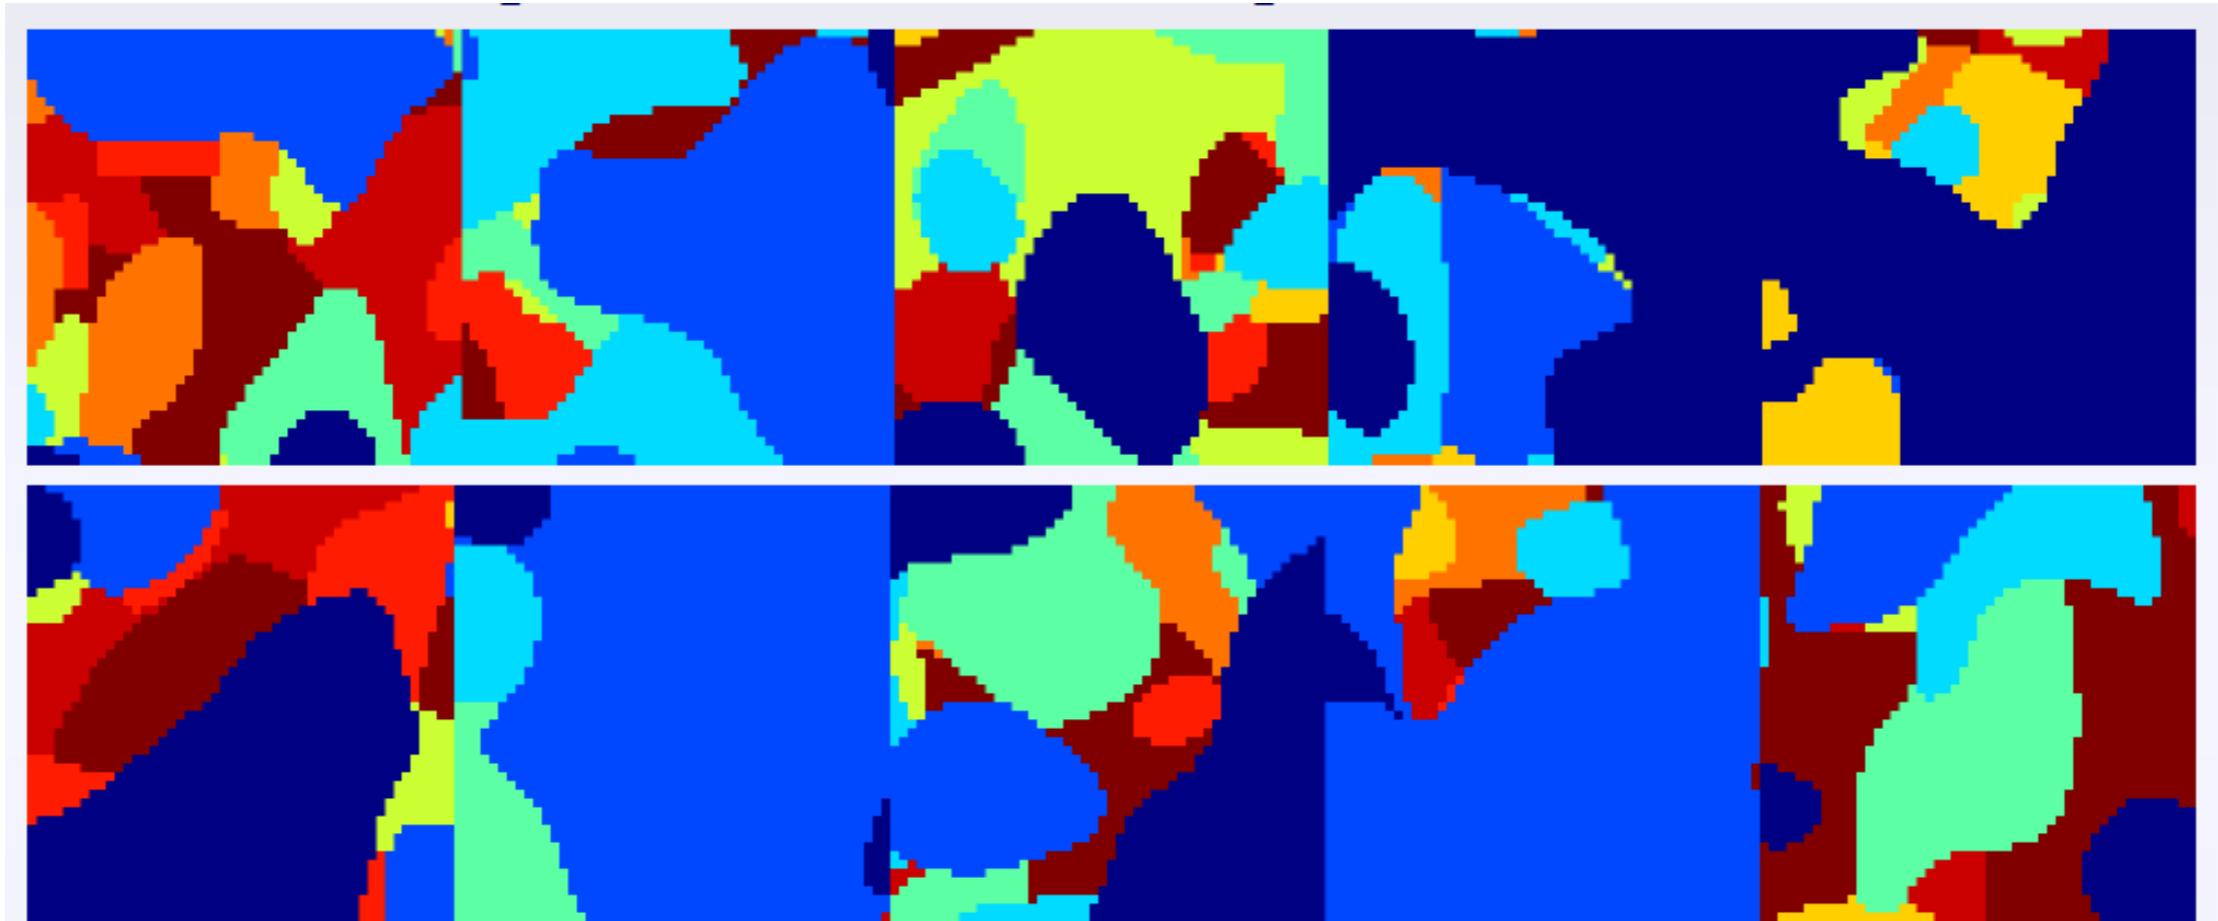


Expected Layer Appearance

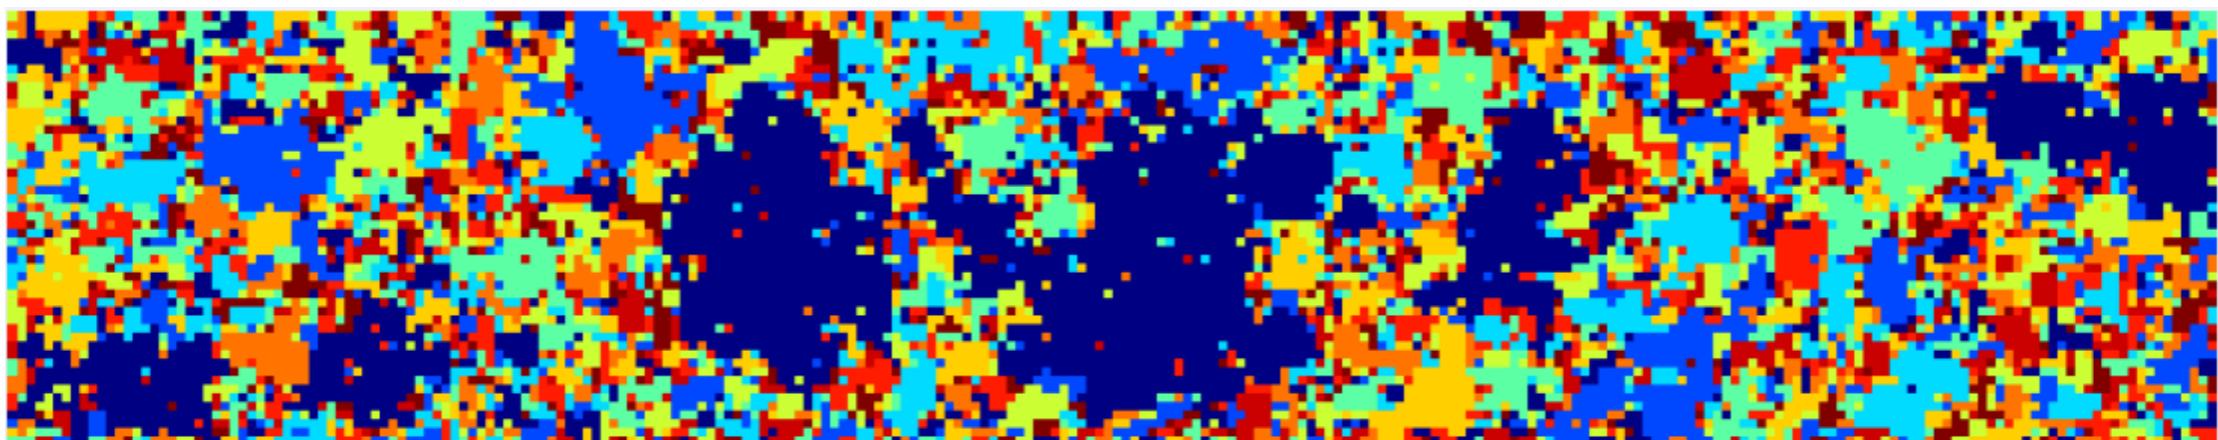


Sudderth & Jordan, 2008

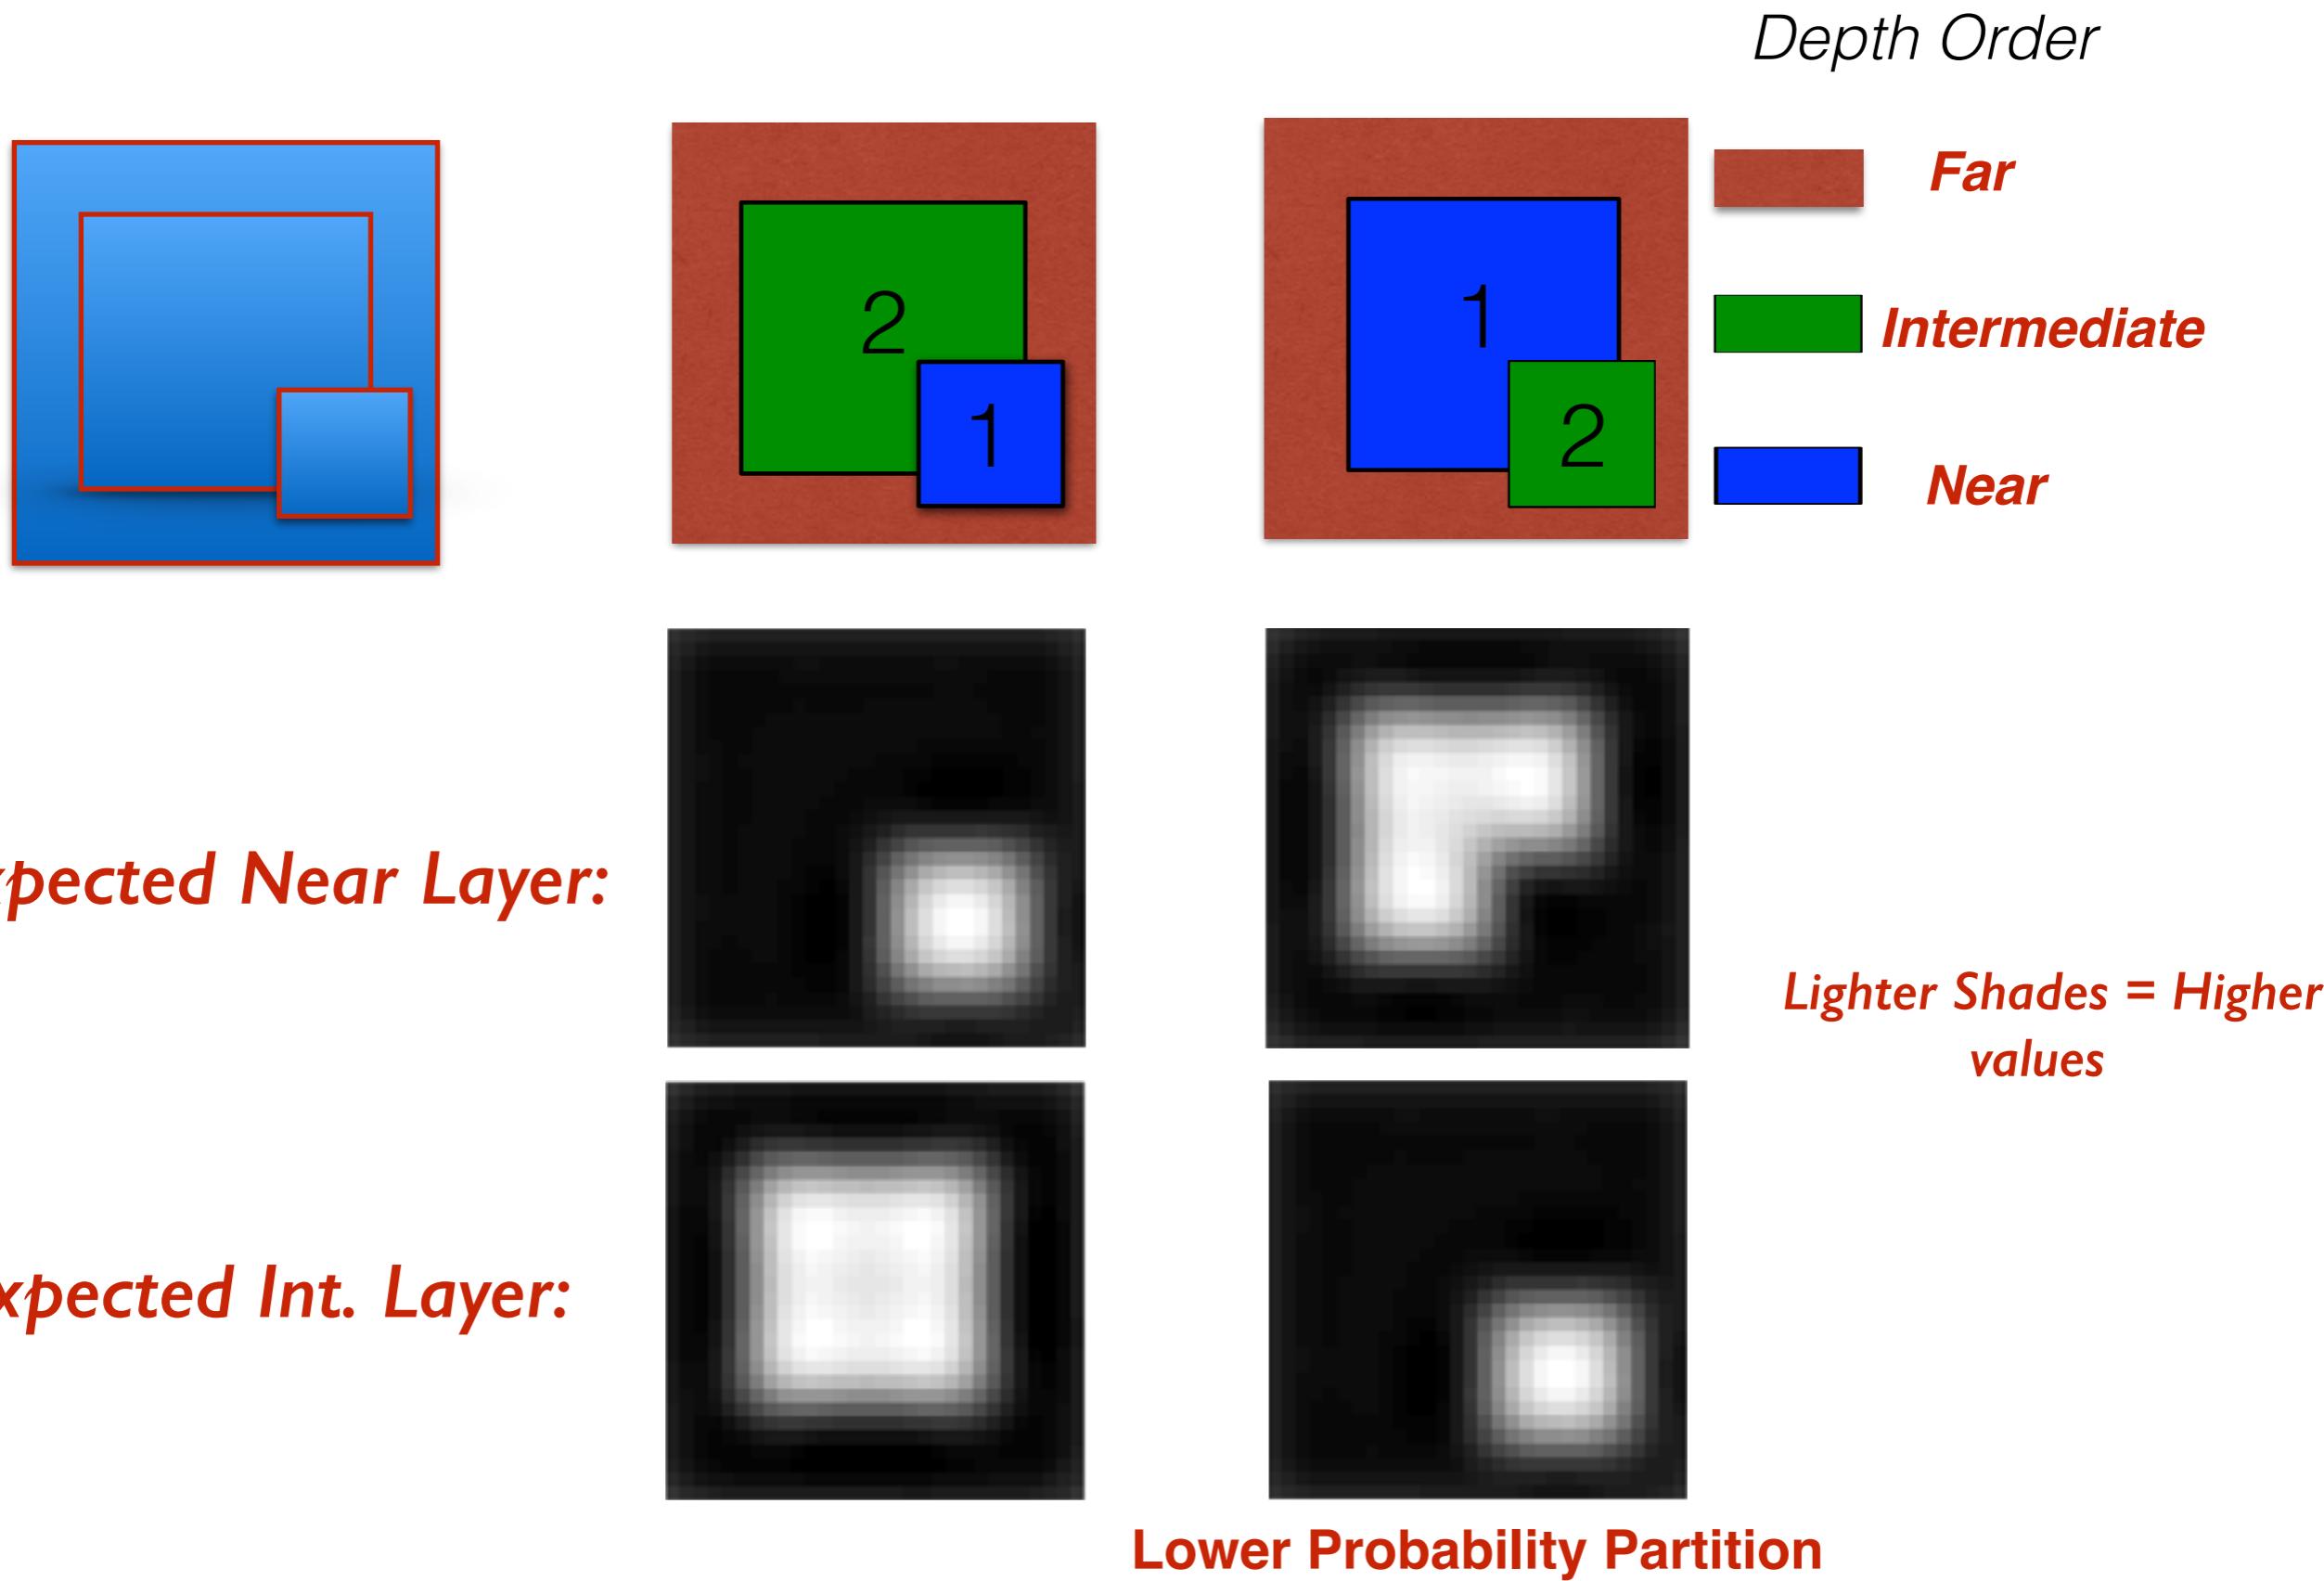
# Generative Samples from the Model



*Samples from a Potts model:*

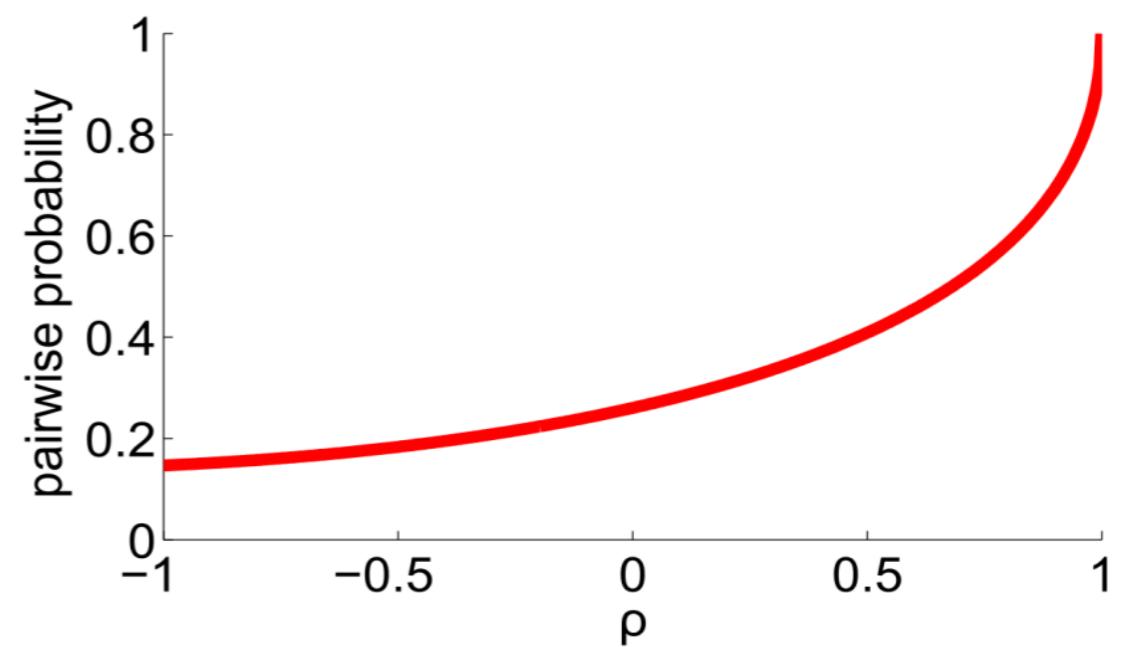


# Depth-ordered Layers



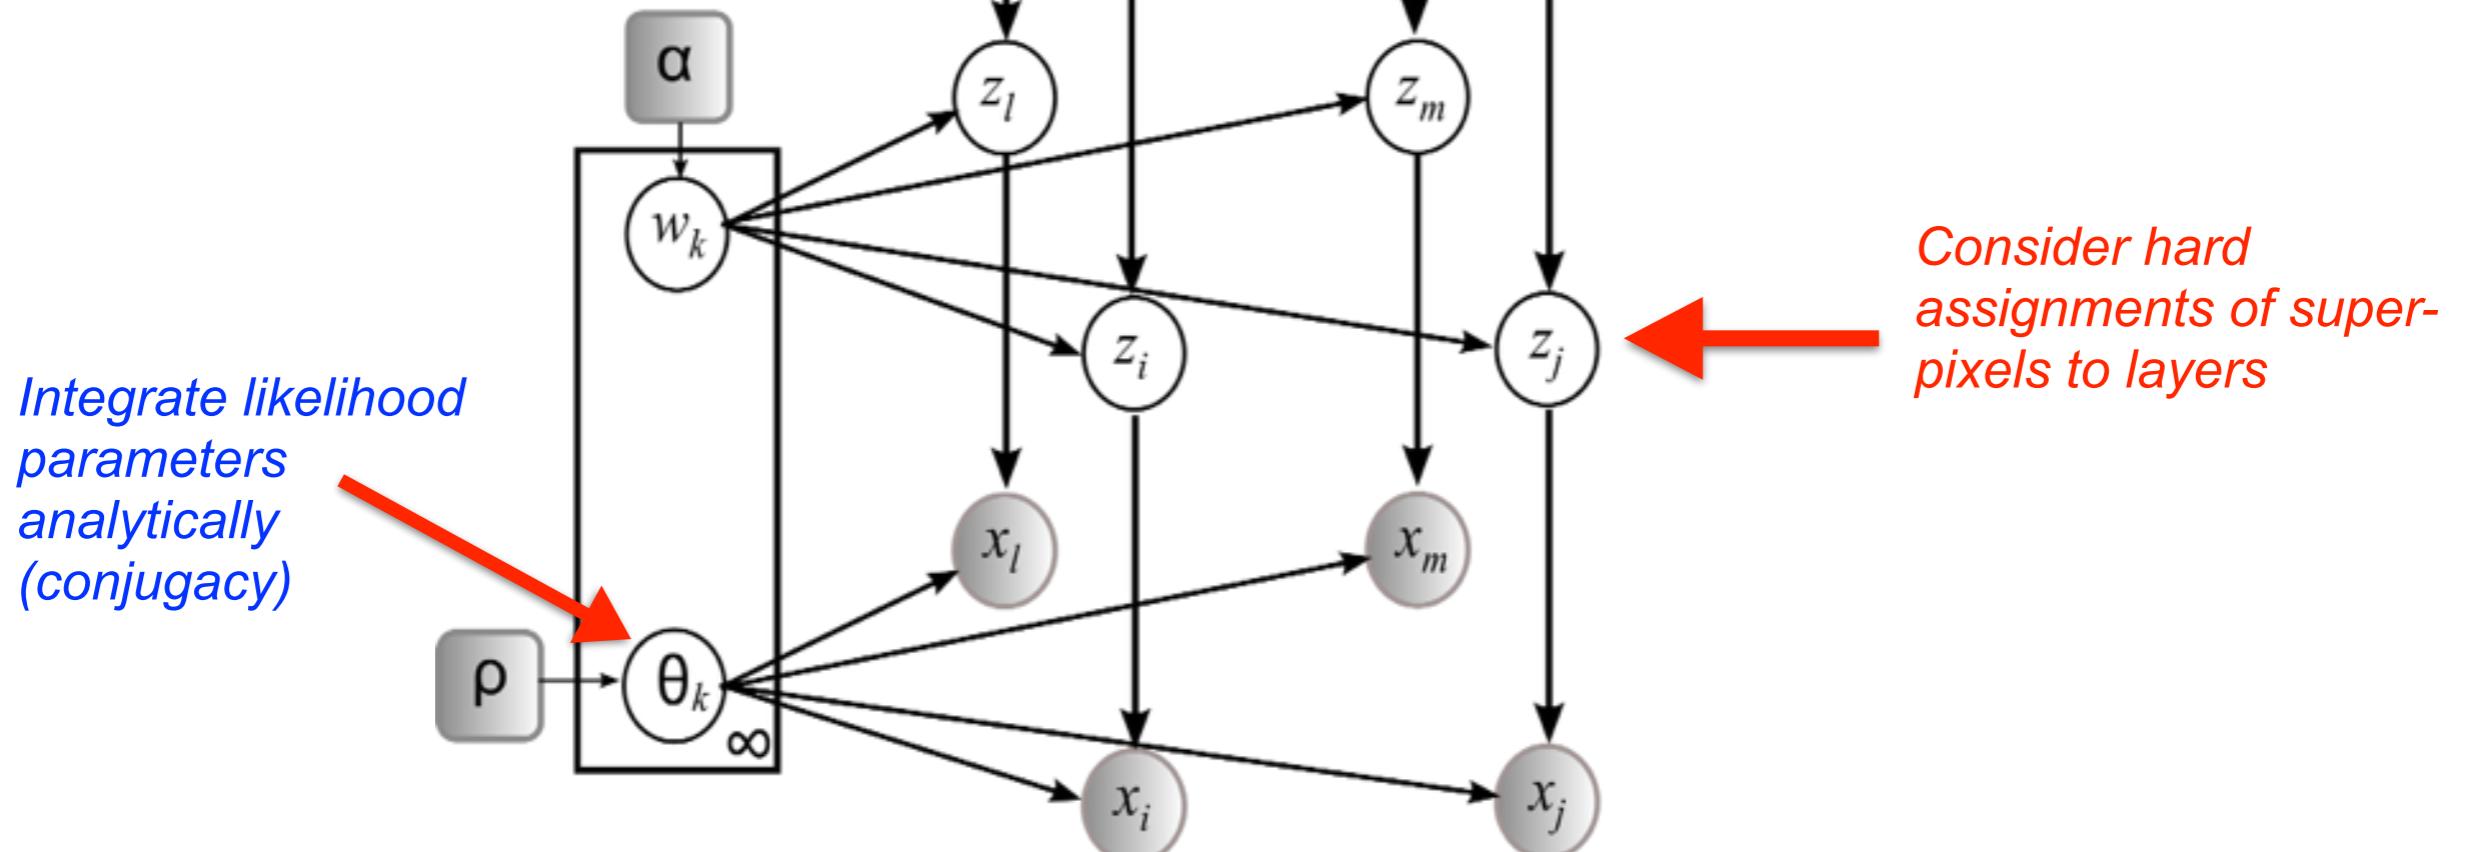
# Calibrating Gaussian Processes

- *Binary Classification*: Probability of two super-pixels belonging to the same segment, learned via regularized logistic regression.
- *Co-occurrence to covariance*: Injective mapping between probabilities and GP correlation
- *Valid Model*: Project independently learned covariance matrix to the closest PSD matrix.



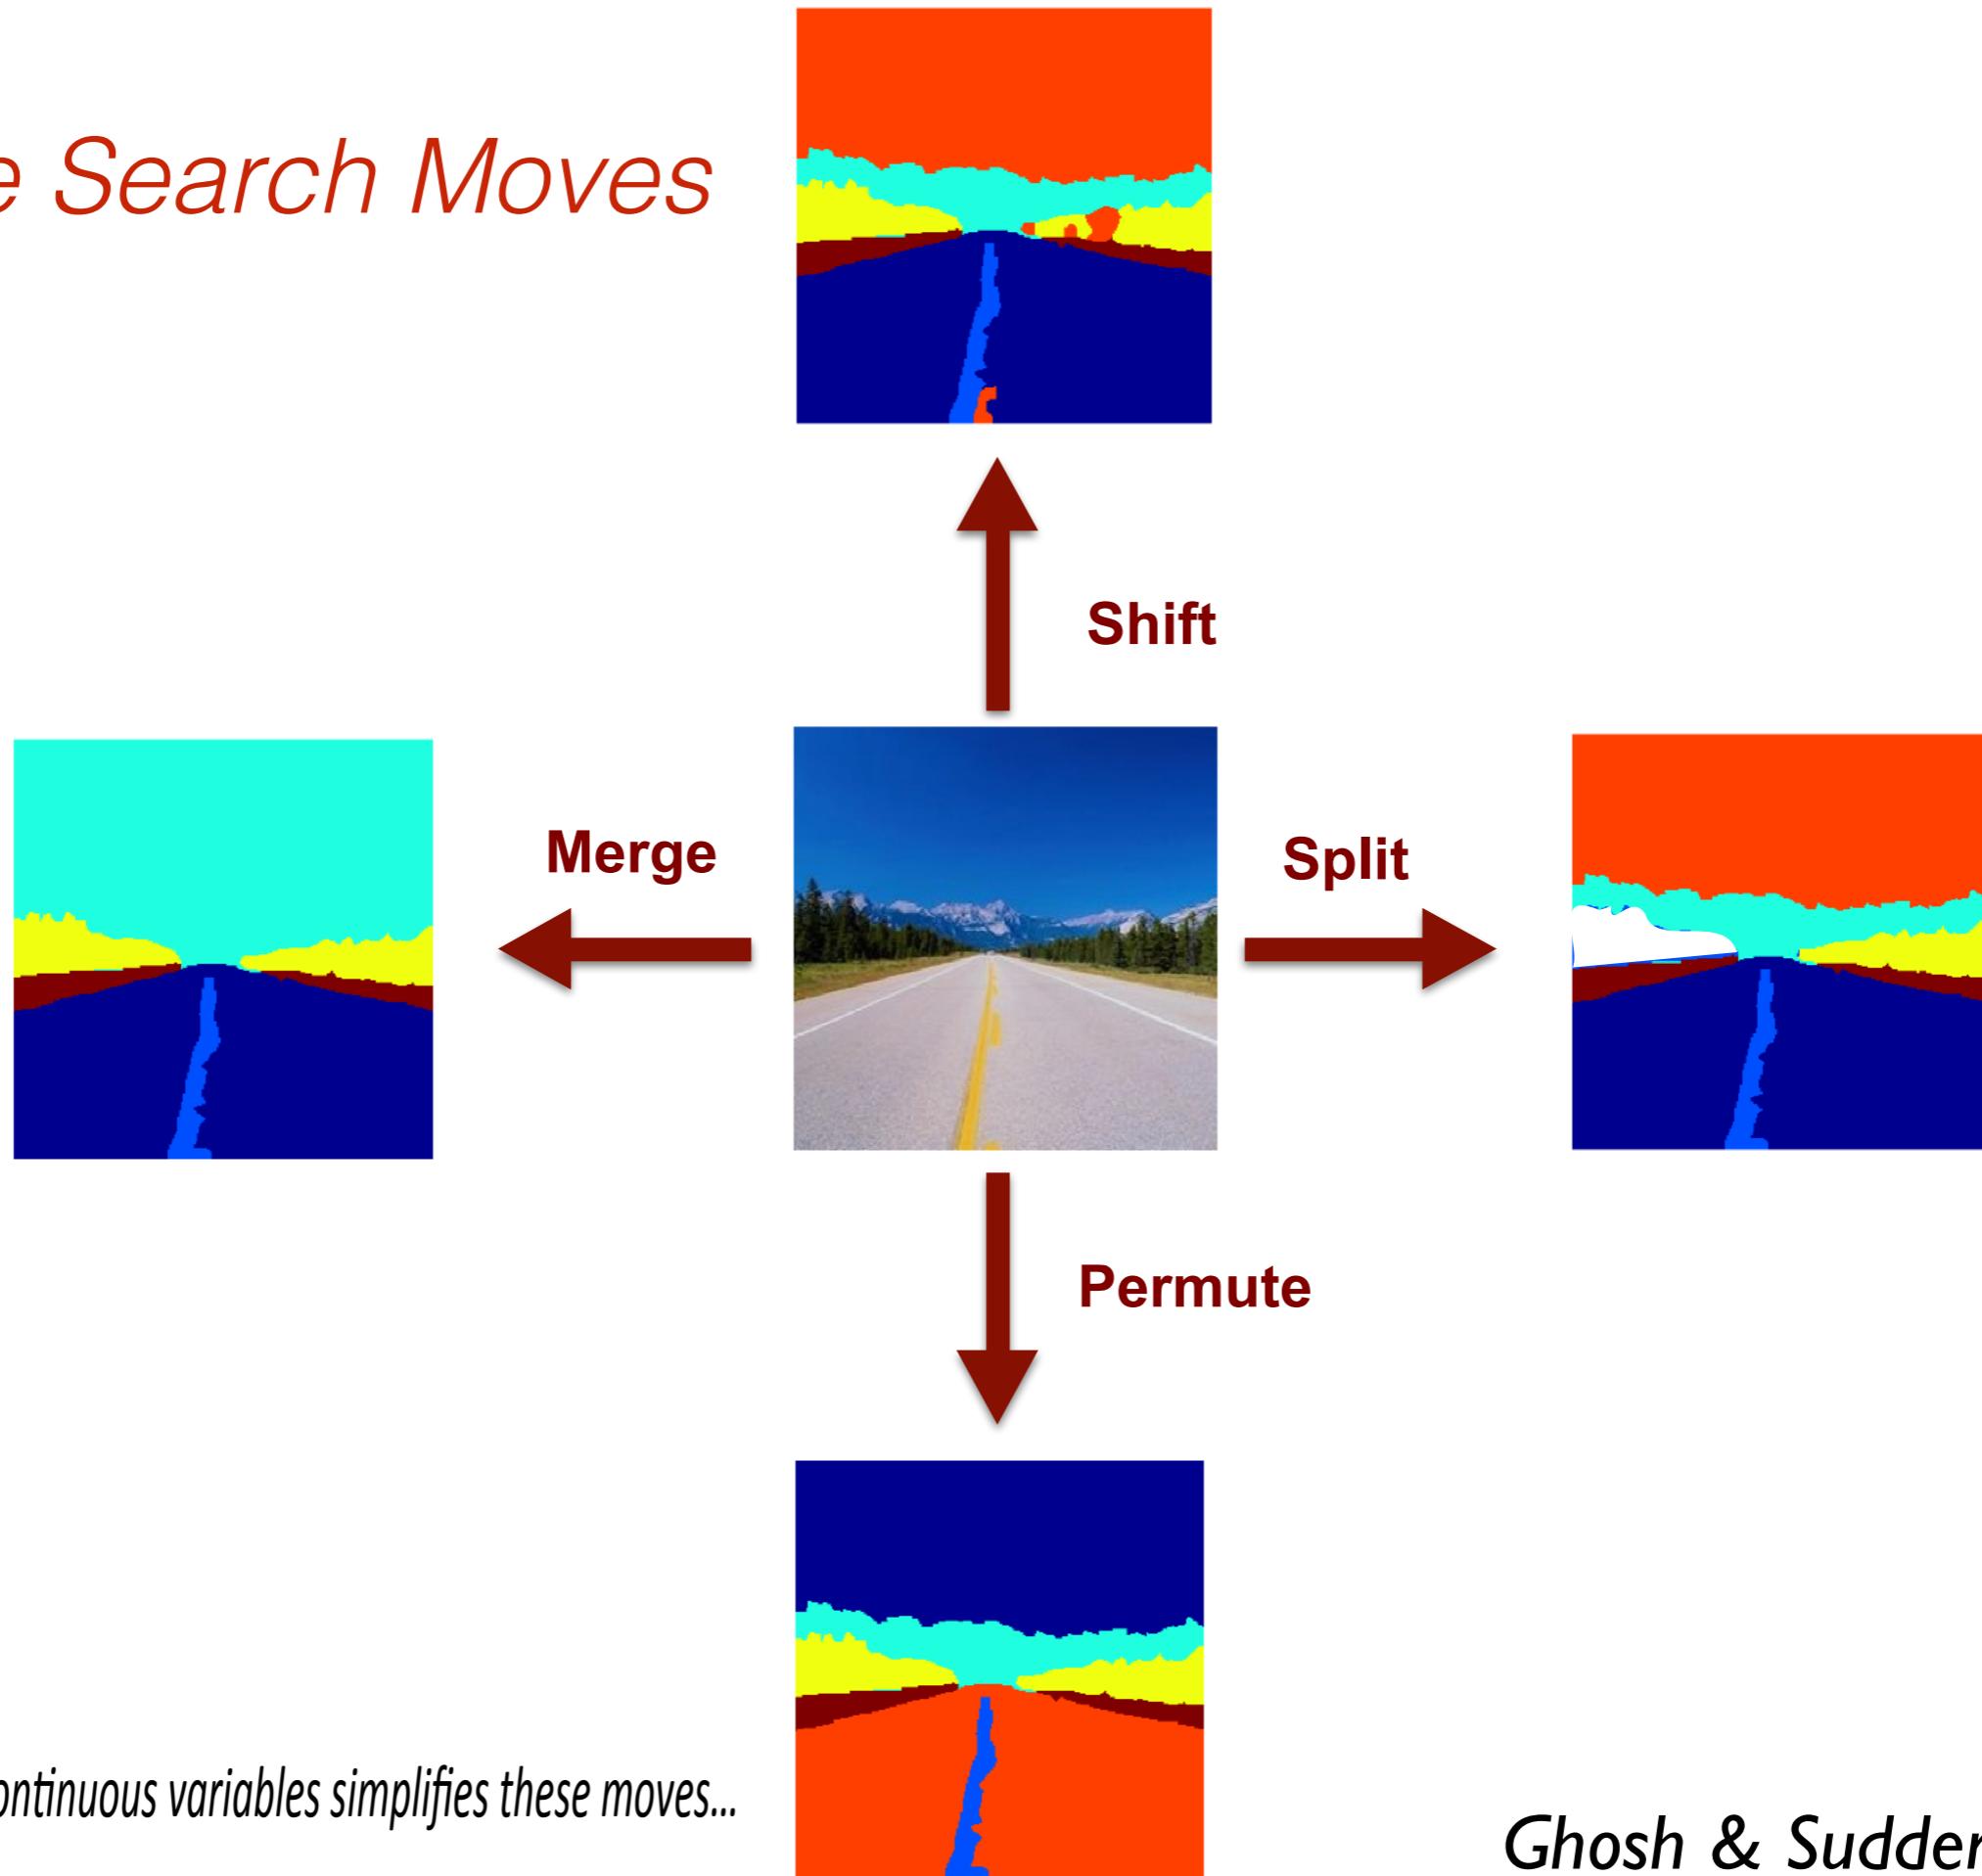
# Inference via Stochastic Search

Marginalize a **finite number** of layers via expectation propagation (EP):  
approximate but accurate

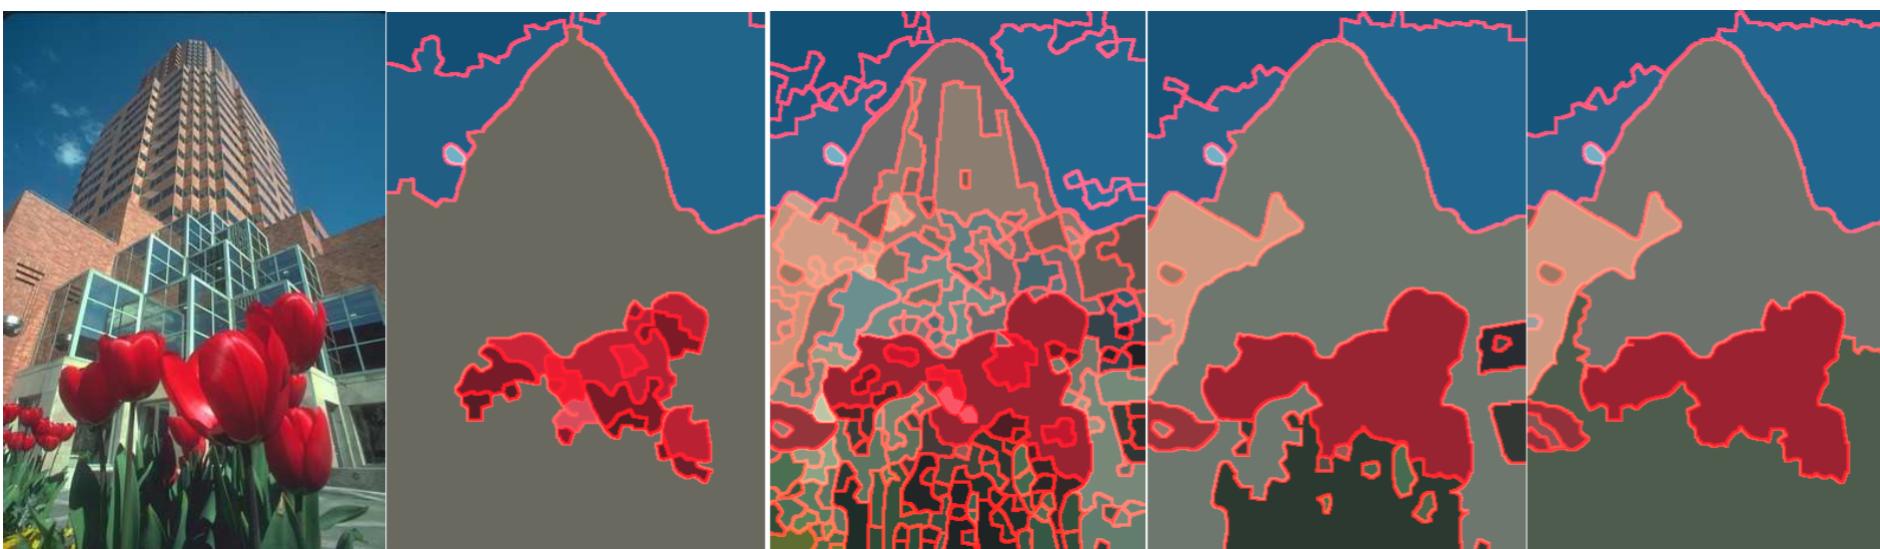
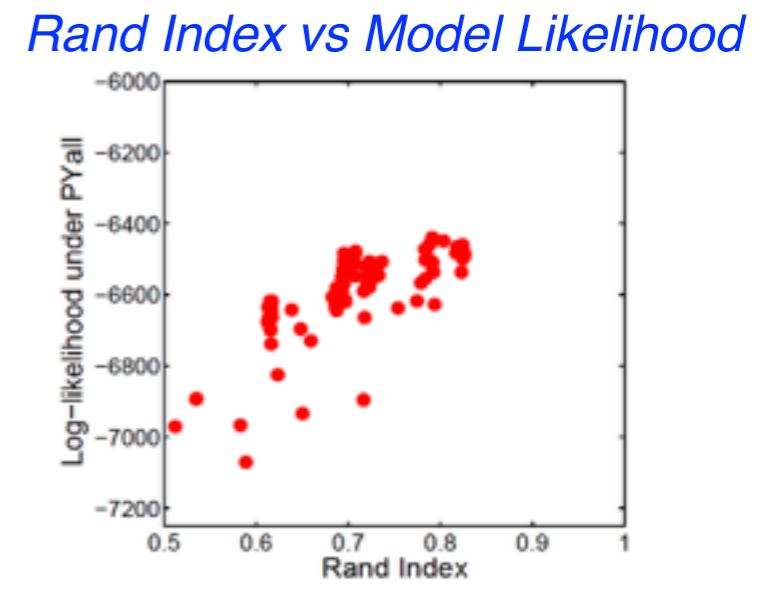
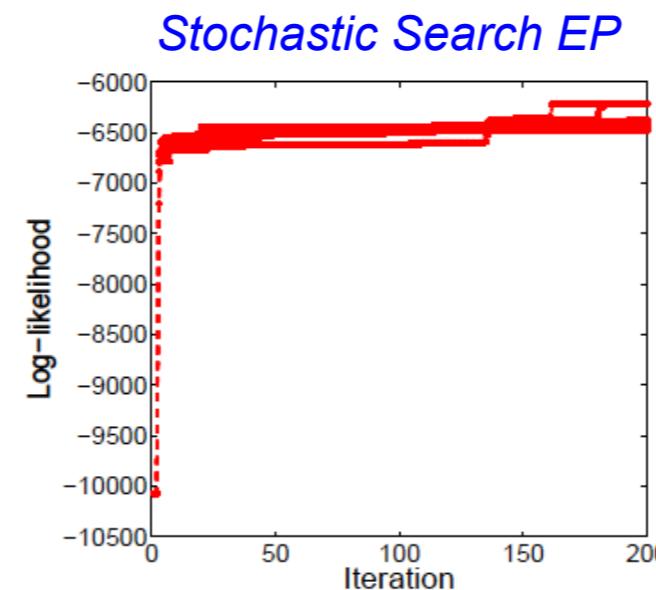
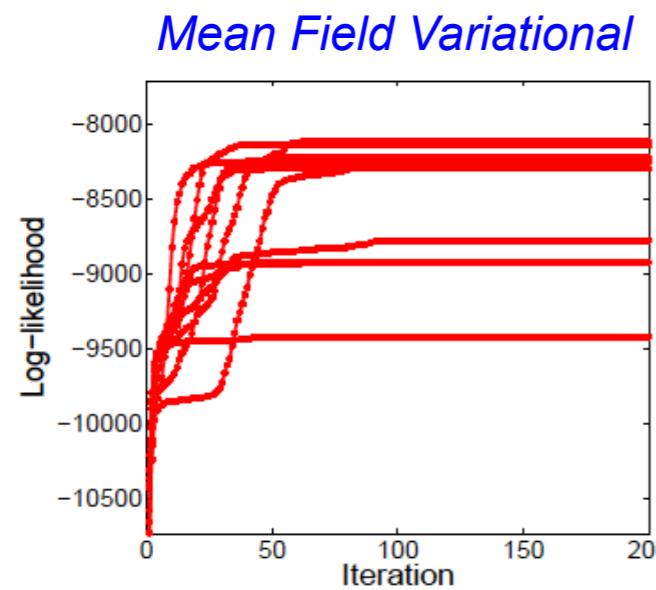


$$p(Z | X, \alpha, \rho) \propto p(X | Z, \rho) \underbrace{p(Z | \alpha)}_{\text{EP}}$$

# *Discrete Search Moves*



# Robust, Reliable Inference

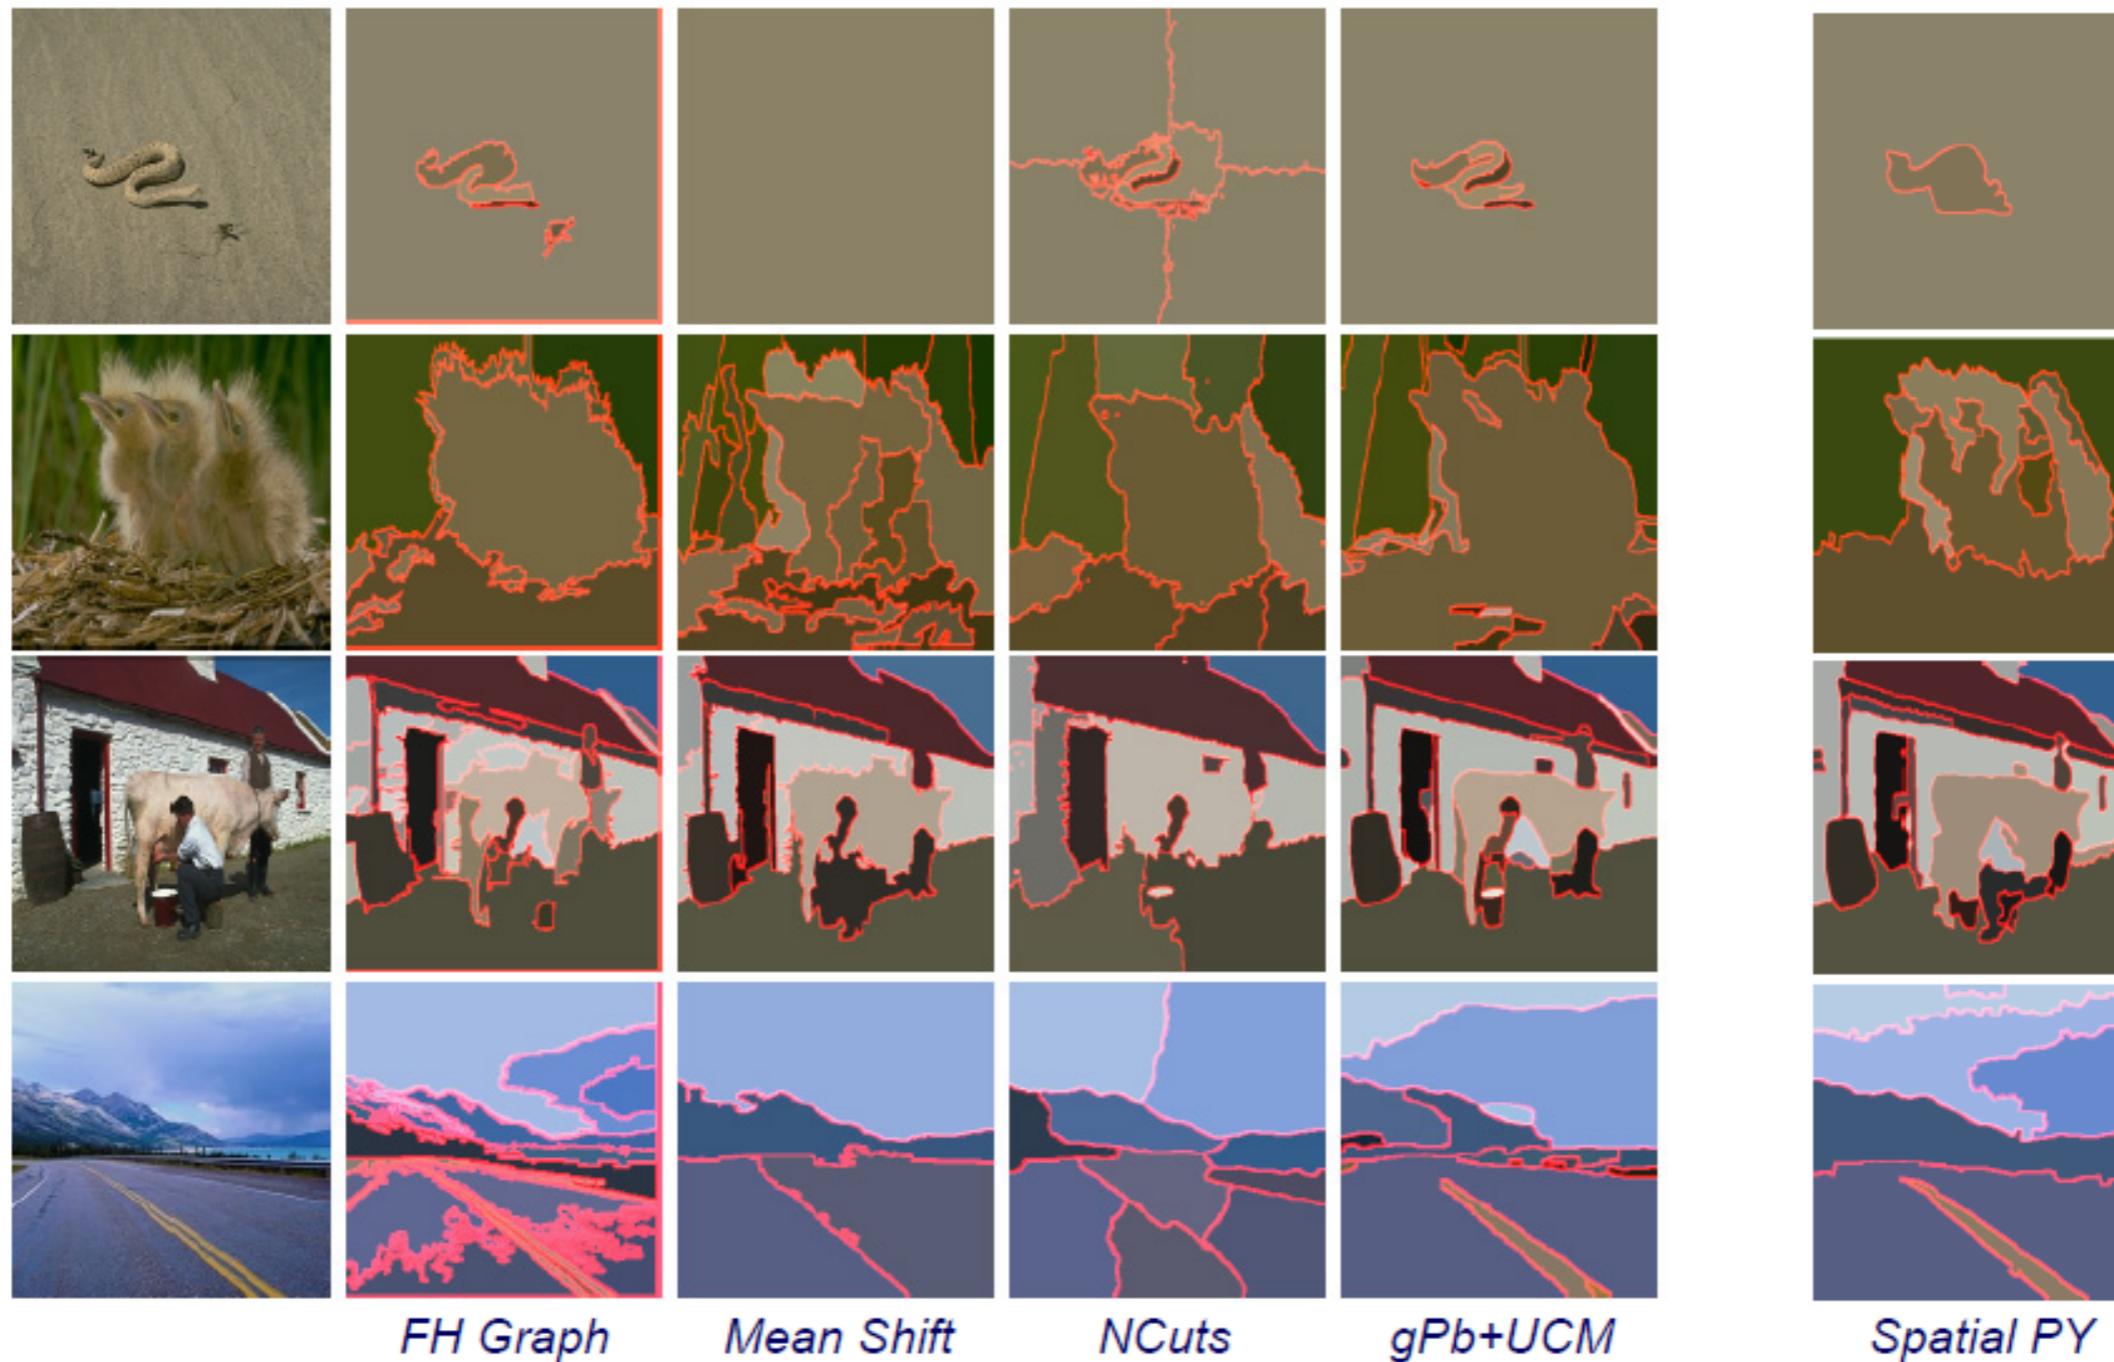


*Best*  
*Worst*  
*Mean Field Variational*

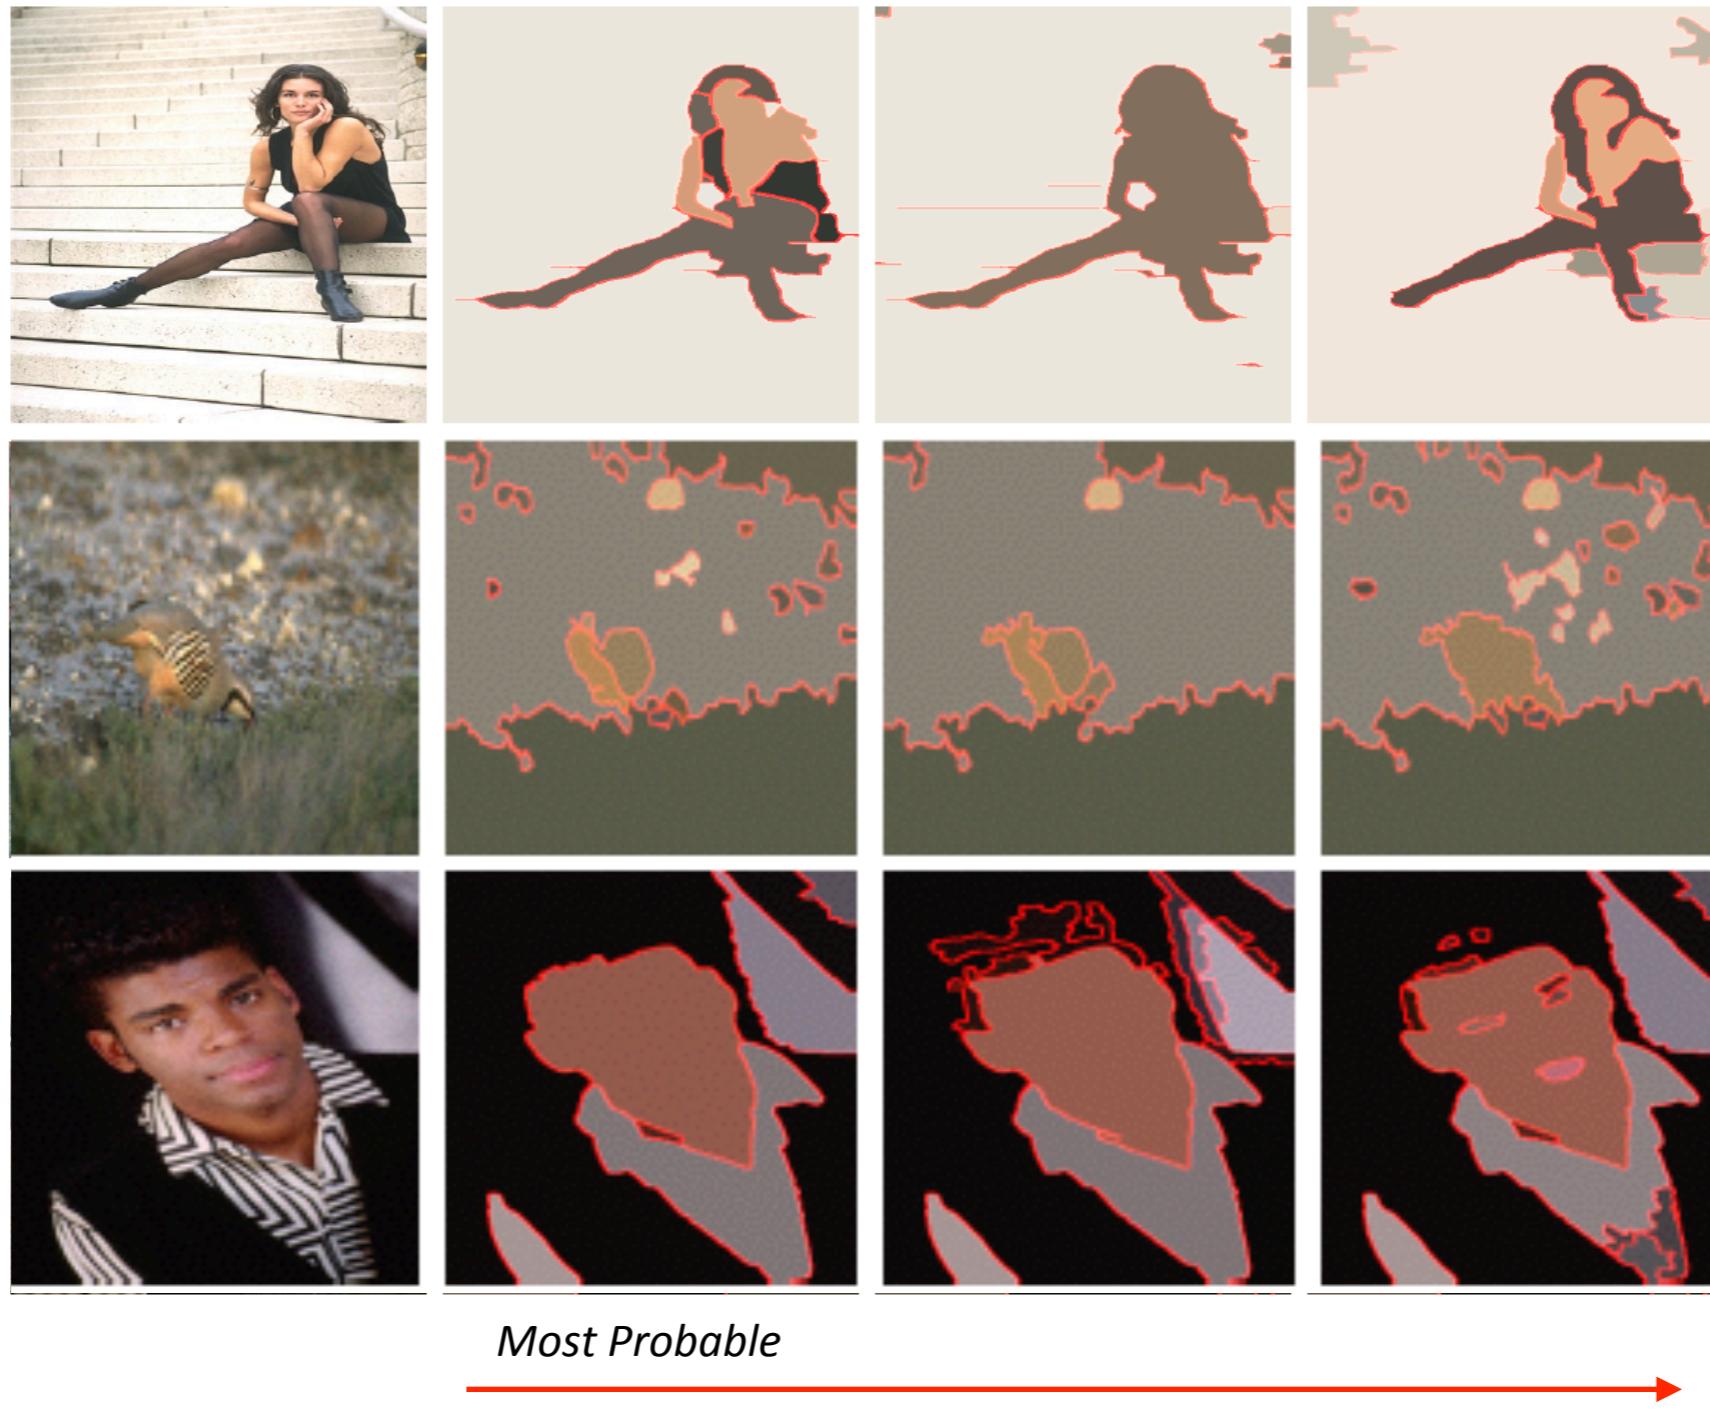
*Best*  
*Worst*  
*Stochastic Search EP*

**External Model Validation**

# Qualitative Results



# Posterior distribution over segmentations



# Quantitative Performance

Algorithms	PRI	VI	SegCover
Ncuts	0.74	2.5	0.38
MS	0.77	2.5	0.44
FH	0.77	2.1	0.52
gPb	0.81	2.0	0.58
PYdist	0.72	2.1	0.51
PYall	0.76	2.1	0.52

Berkeley Segmentation

gPb	0.74	2.1	0.53
PYall	0.73	1.9	0.55

LabelMe Scenes

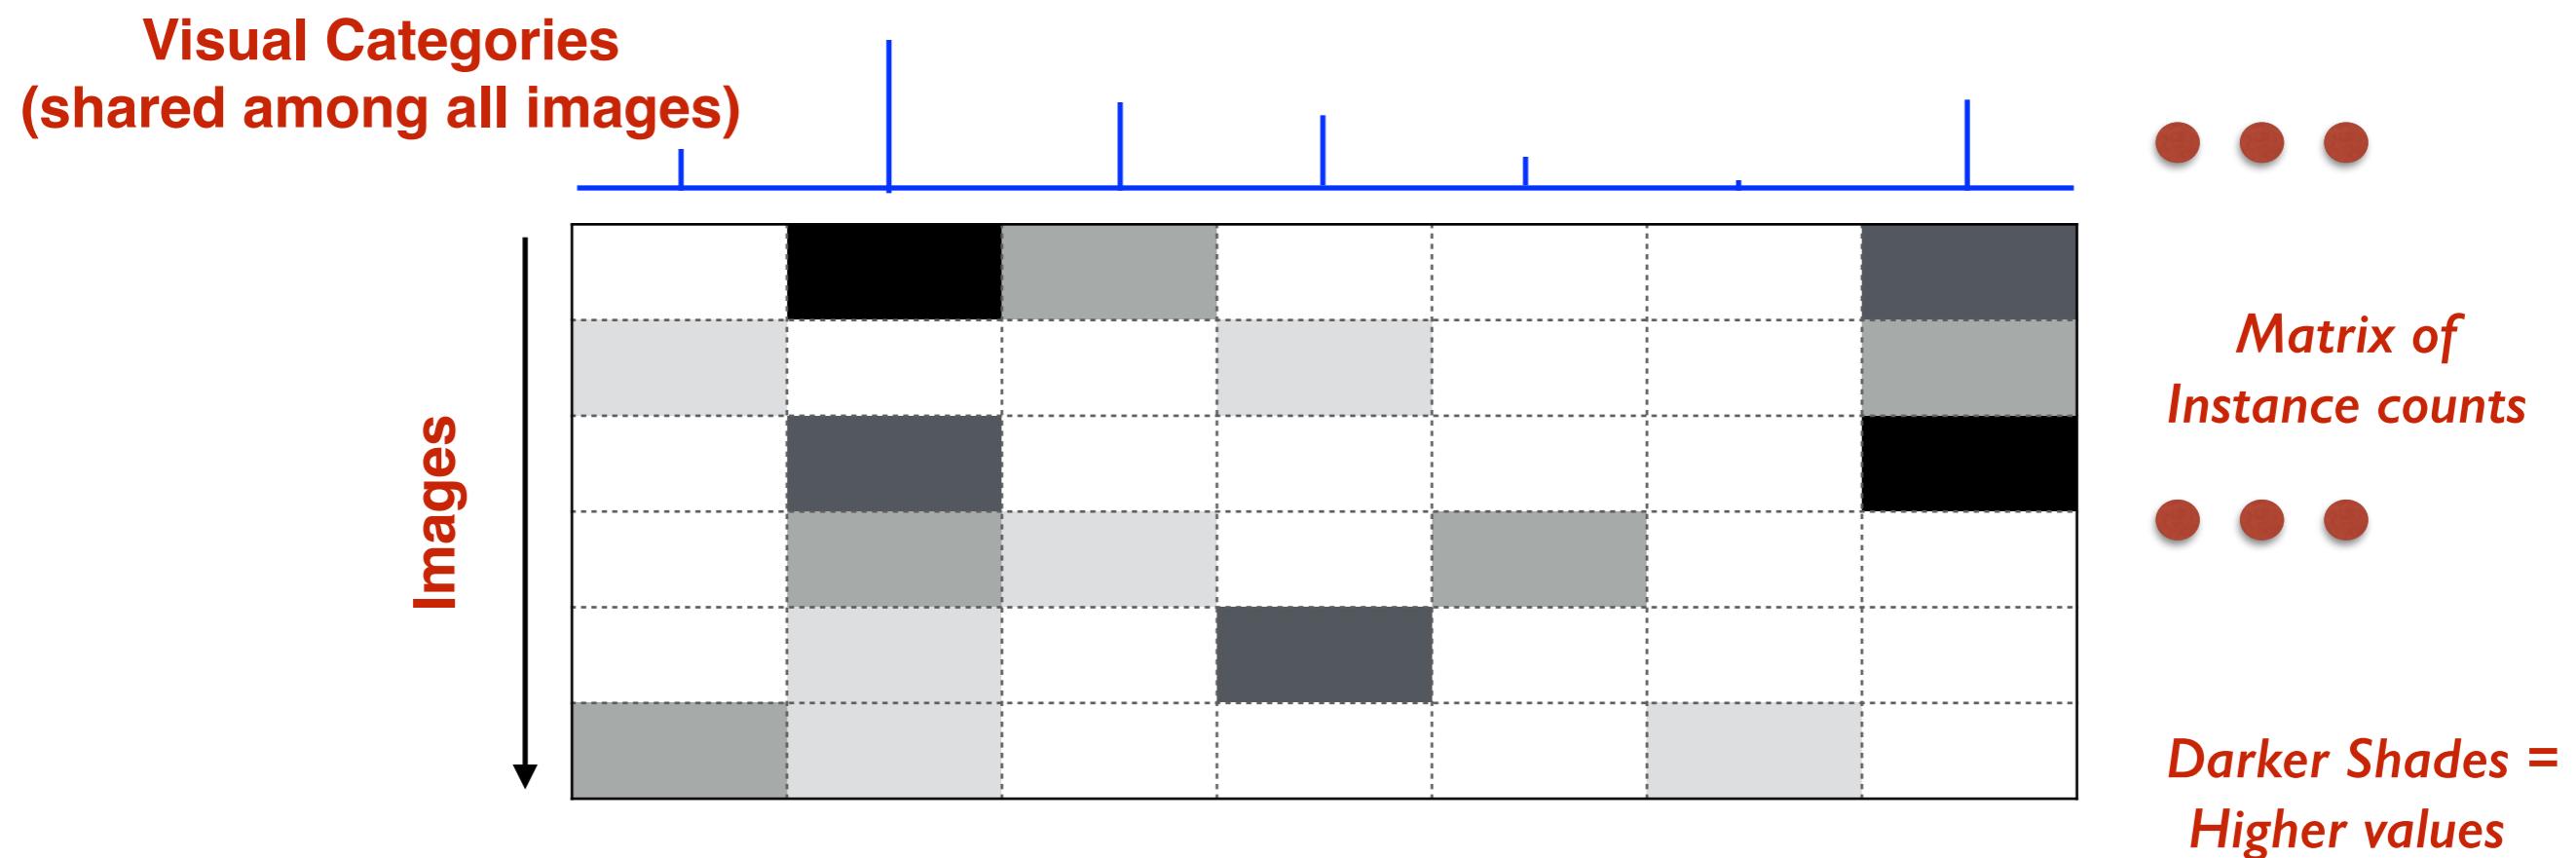
- On BSDS, similar or better than all methods except gPb
- On LabelMe, performance of Spatial PY is better than gPb

# Talk Outline

- Deformable 3D object segmentation
- Image and Video segmentation
- Layered decomposition of natural images
- Proposed Work

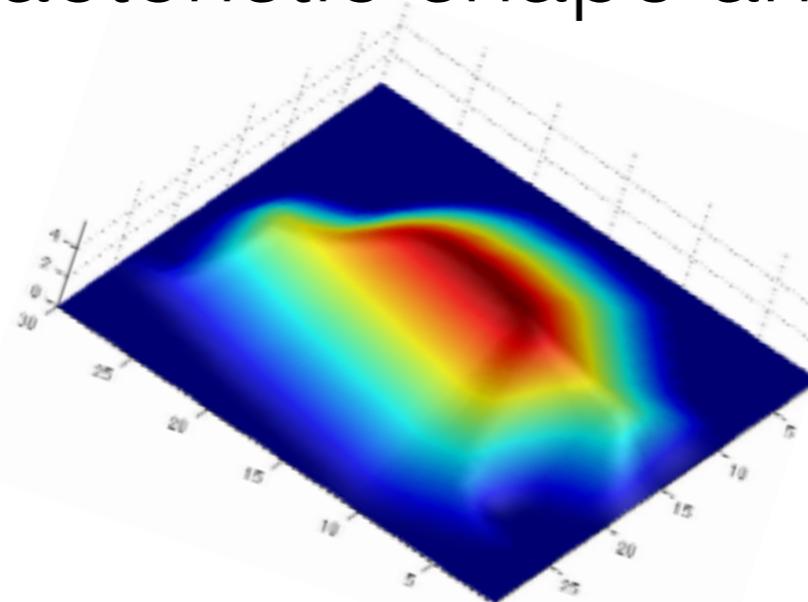
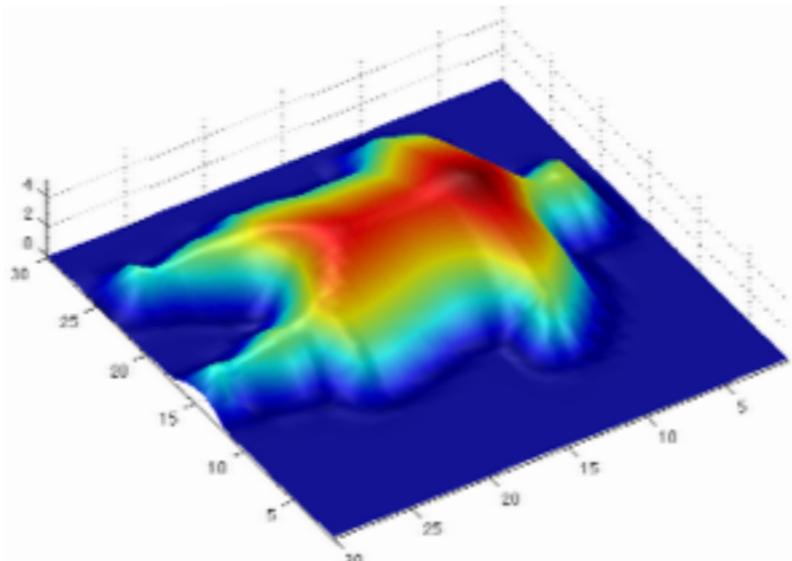
# Shared Analysis of Image Corporuses

- Joint image parsing and labeling
- Simultaneous inference of coarse 3D structure — placement of layers in the 3D scene
- Explicit modeling of object and instance occurrence through the Beta-Geometric process:

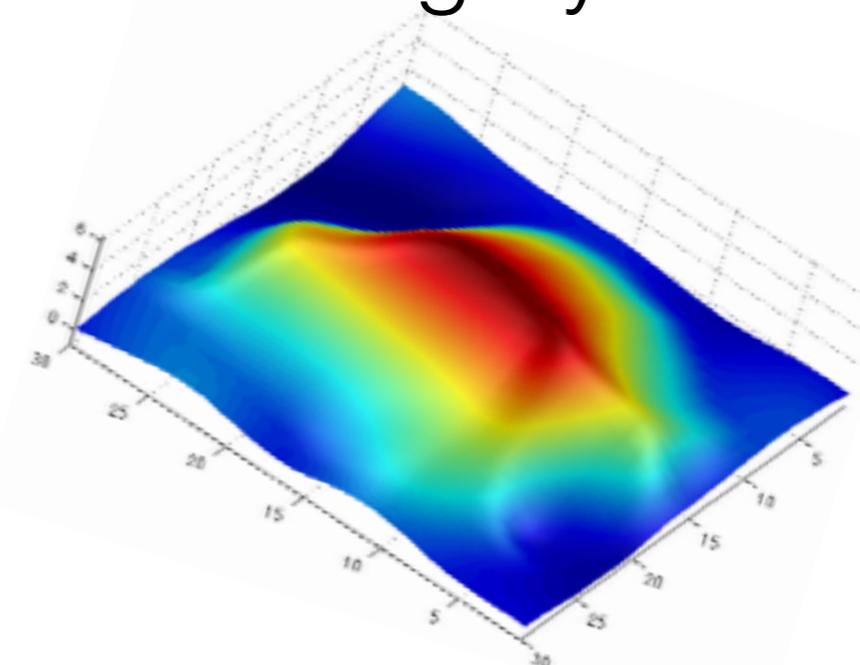
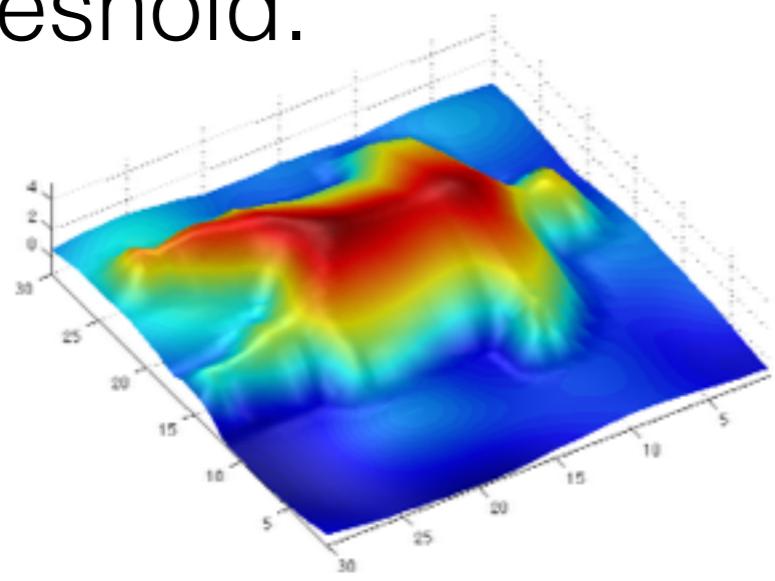


# Category specific shapes and sizes

- Each visual category is embellished with a mean function modeling its characteristic shape and size

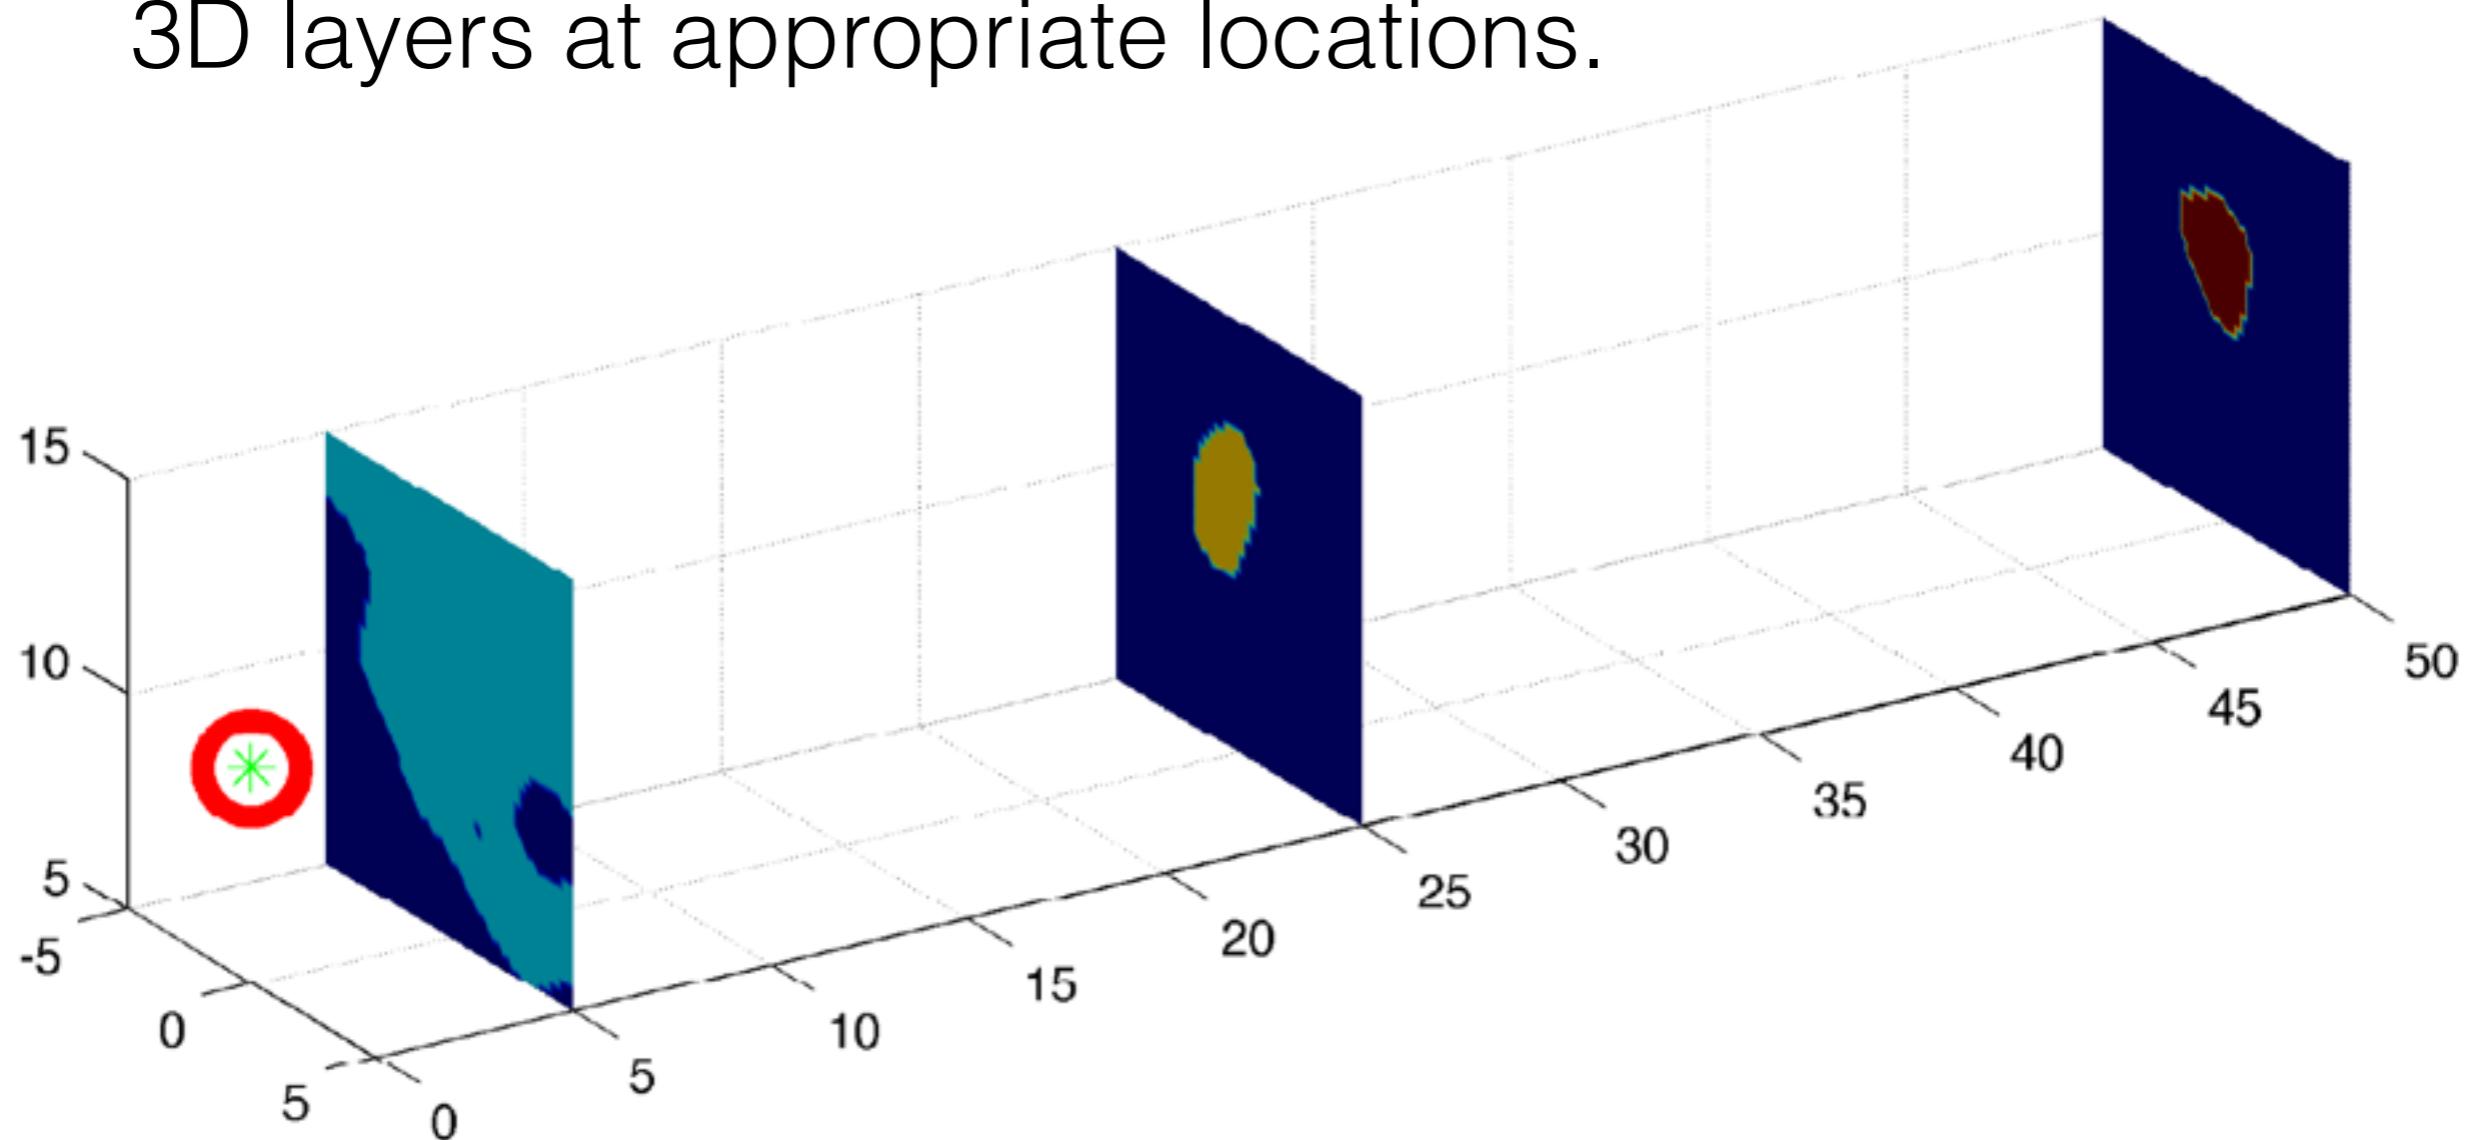


- Sample an instance specific realization of a GP parametrized with the appropriate category mean and threshold.



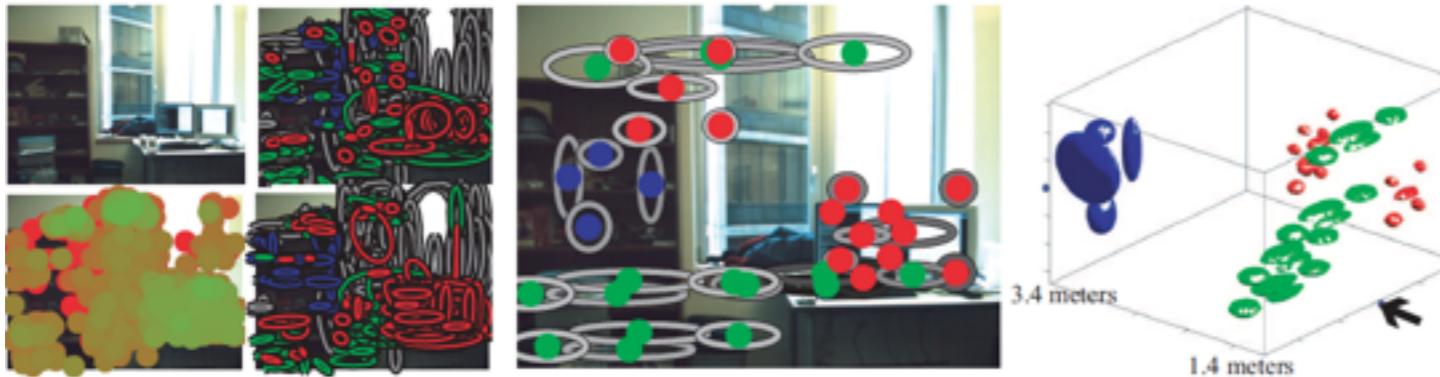
# 2D Images from 3D structure

- Sample a 3D location and place instance specific 3D layers at appropriate locations.



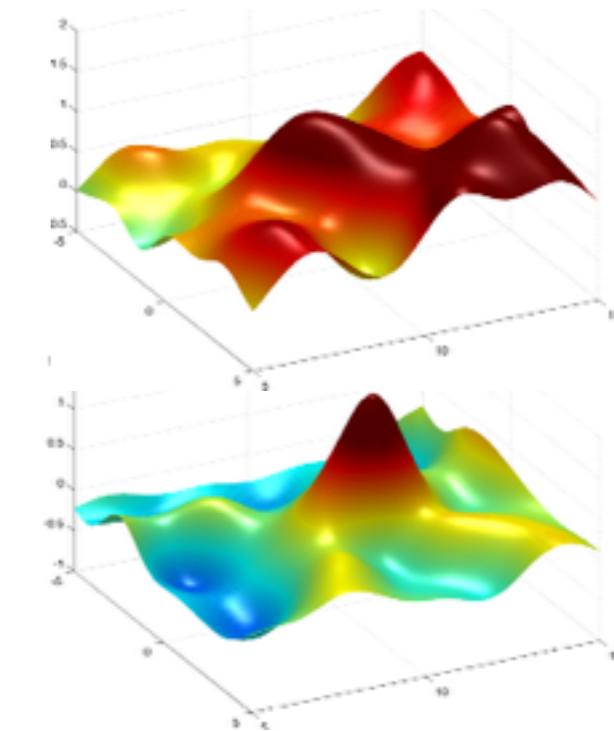
- Project 3D layers into the image through perspective projection.

# Building Blocks

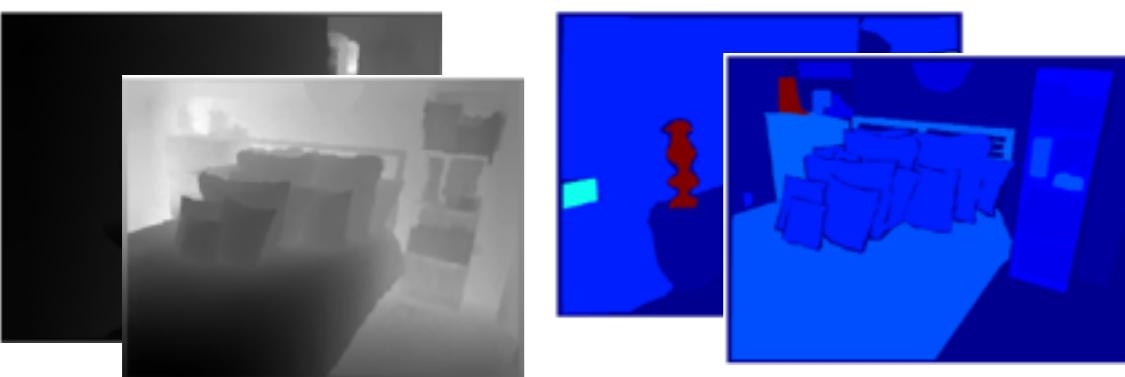


**Depth from familiar objects**

Sudderth, Torralba, Freeman & Willsky. CVPR'06

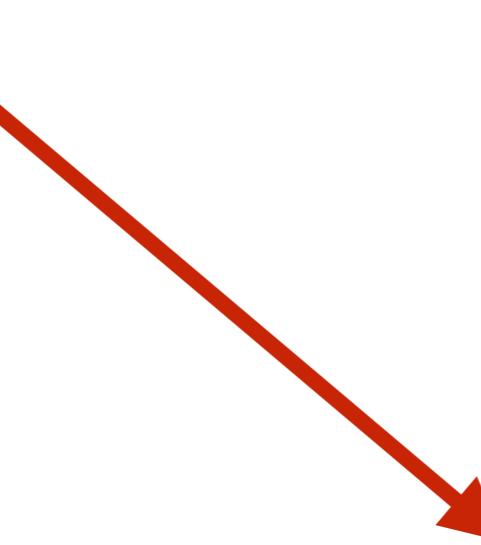


**BNP Layered models**  
Ghosh and Sudderth. CVPR'12  
Sudderth and Jordan. NIPS'08



**Color-depth datasets**

Silberman, Kohli, Hoiem & Fergus. ECCV'12



# Timeline

## Publications

- Ghosh et al., *Spatial distance dependent Chinese restaurant processes for image segmentation*, **NIPS'11**.
- Ghosh & Sudderth.,  
*Nonparametric Learning for Layered Segmentation of Natural Images*, **CVPR'12**.
- Ghosh et al., *From Deformations to Parts: Motion-based Segmentation of 3D Objects*, **NIPS'12**.

## TO DO

- Clean up the work on Generalized Hierarchical ddCRP and submit to ICML 2014
- A parametric version of the shared segmentation model to ECCV 2014
- The Beta-Geometric nonparametric version will go to NIPS 2014
- Summer'2014 tie up loose ends, defend and graduate at the end of summer.

# Questions?

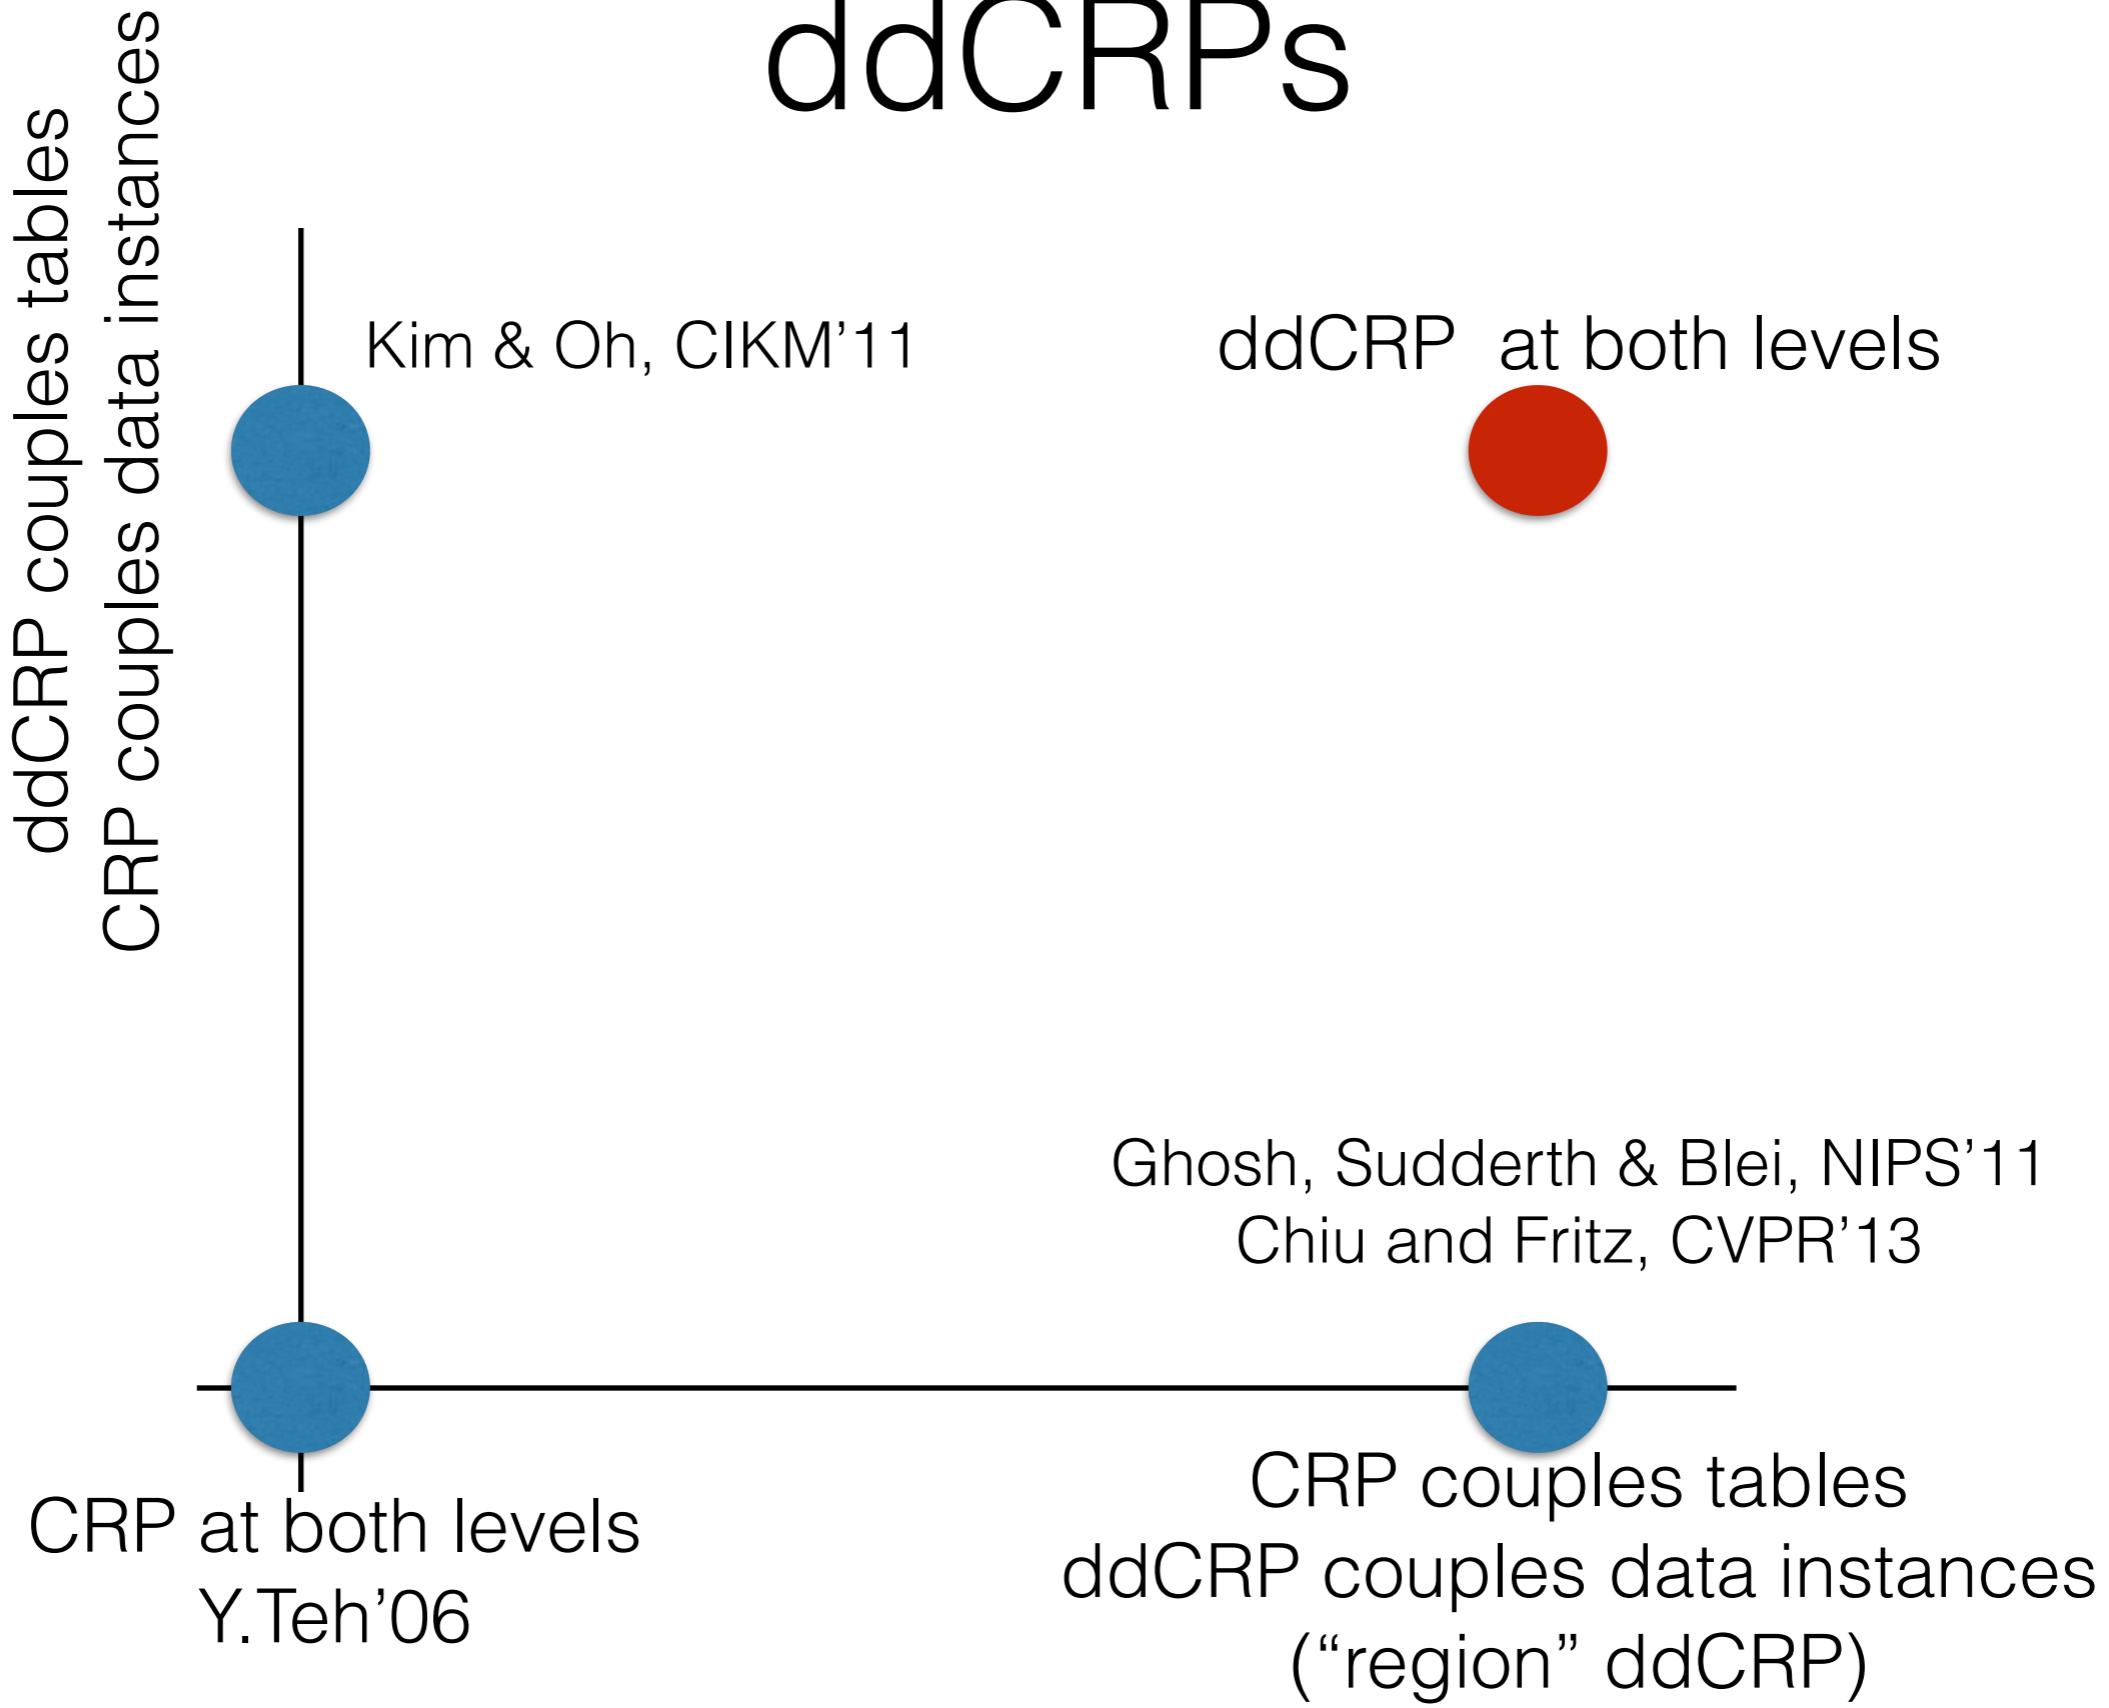
# Mixture of Regressions

- Model is more general.
- Groups (feature,response) pairs with similar relationships.
- Show examples.

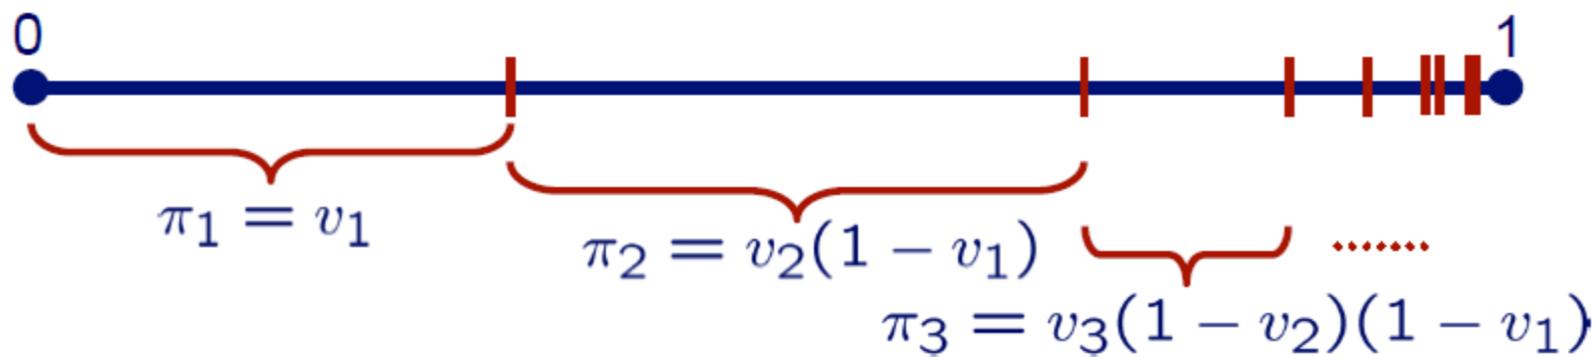
# Moderate robustness to alignment errors



# Generalized Hierarchical ddCRPs



# Stick Breaking to Layers



$$v_k = \mathbb{P}(z_i = k \mid z_i \neq k-1, \dots, 1)$$

## Sequential Binary Sampler:

$$b_{ki} \sim \text{Bernoulli}(v_k)$$

$$z_i = \min\{k \mid b_{ki} = 1\}$$

- For each data instance  $i$ , go through the bins in order 1 through infinity.
- Toss a biased coin (with the probability of heads =  $v_k$ ) for each bin .
- Pick the bin if the coin turns up heads

# Stick Breaking to Layers

**Sequential Binary Sampler:**

$$b_{ki} \sim \text{Bernoulli}(v_k)$$

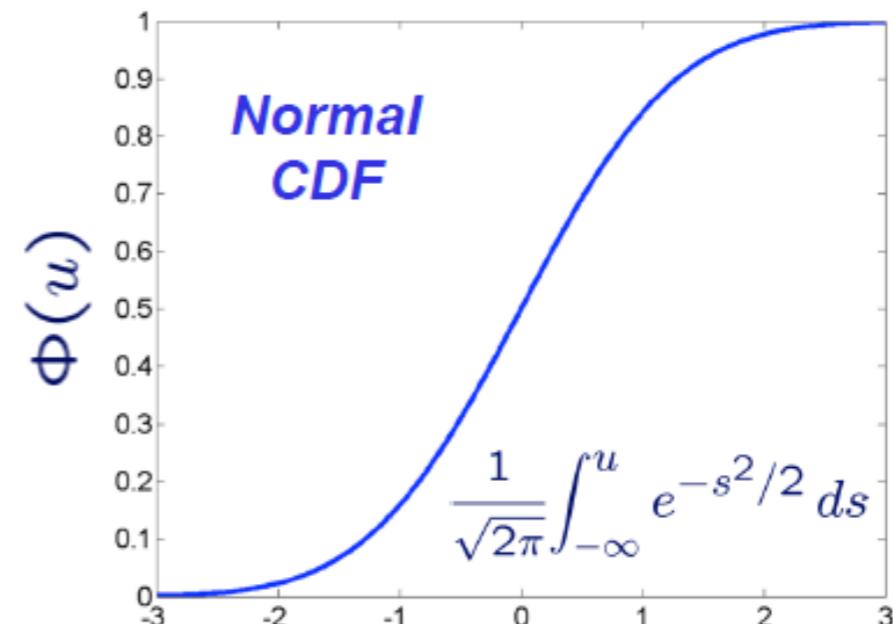
$$z_i = \min\{k \mid b_{ki} = 1\}$$

- For each data instance, go through the bins in order 1 through infinity.
- Sample a unit normal distribution for each bin.
- Pick the bin if the sample is less than the threshold  $\Phi^{-1}(v_k)$

**Gaussian Sampler:**

$$u_{ki} \sim \mathcal{N}(0, 1)$$

$$z_i = \min\{k \mid u_{ki} < \Phi^{-1}(v_k)\}$$



$$\mathbb{P}[\Phi(u_{ki}) < v_k] = v_k$$

because

$$\Phi(u_{ki}) \sim \text{Unif}(0, 1)$$

# Calibrating Covariance Kernels

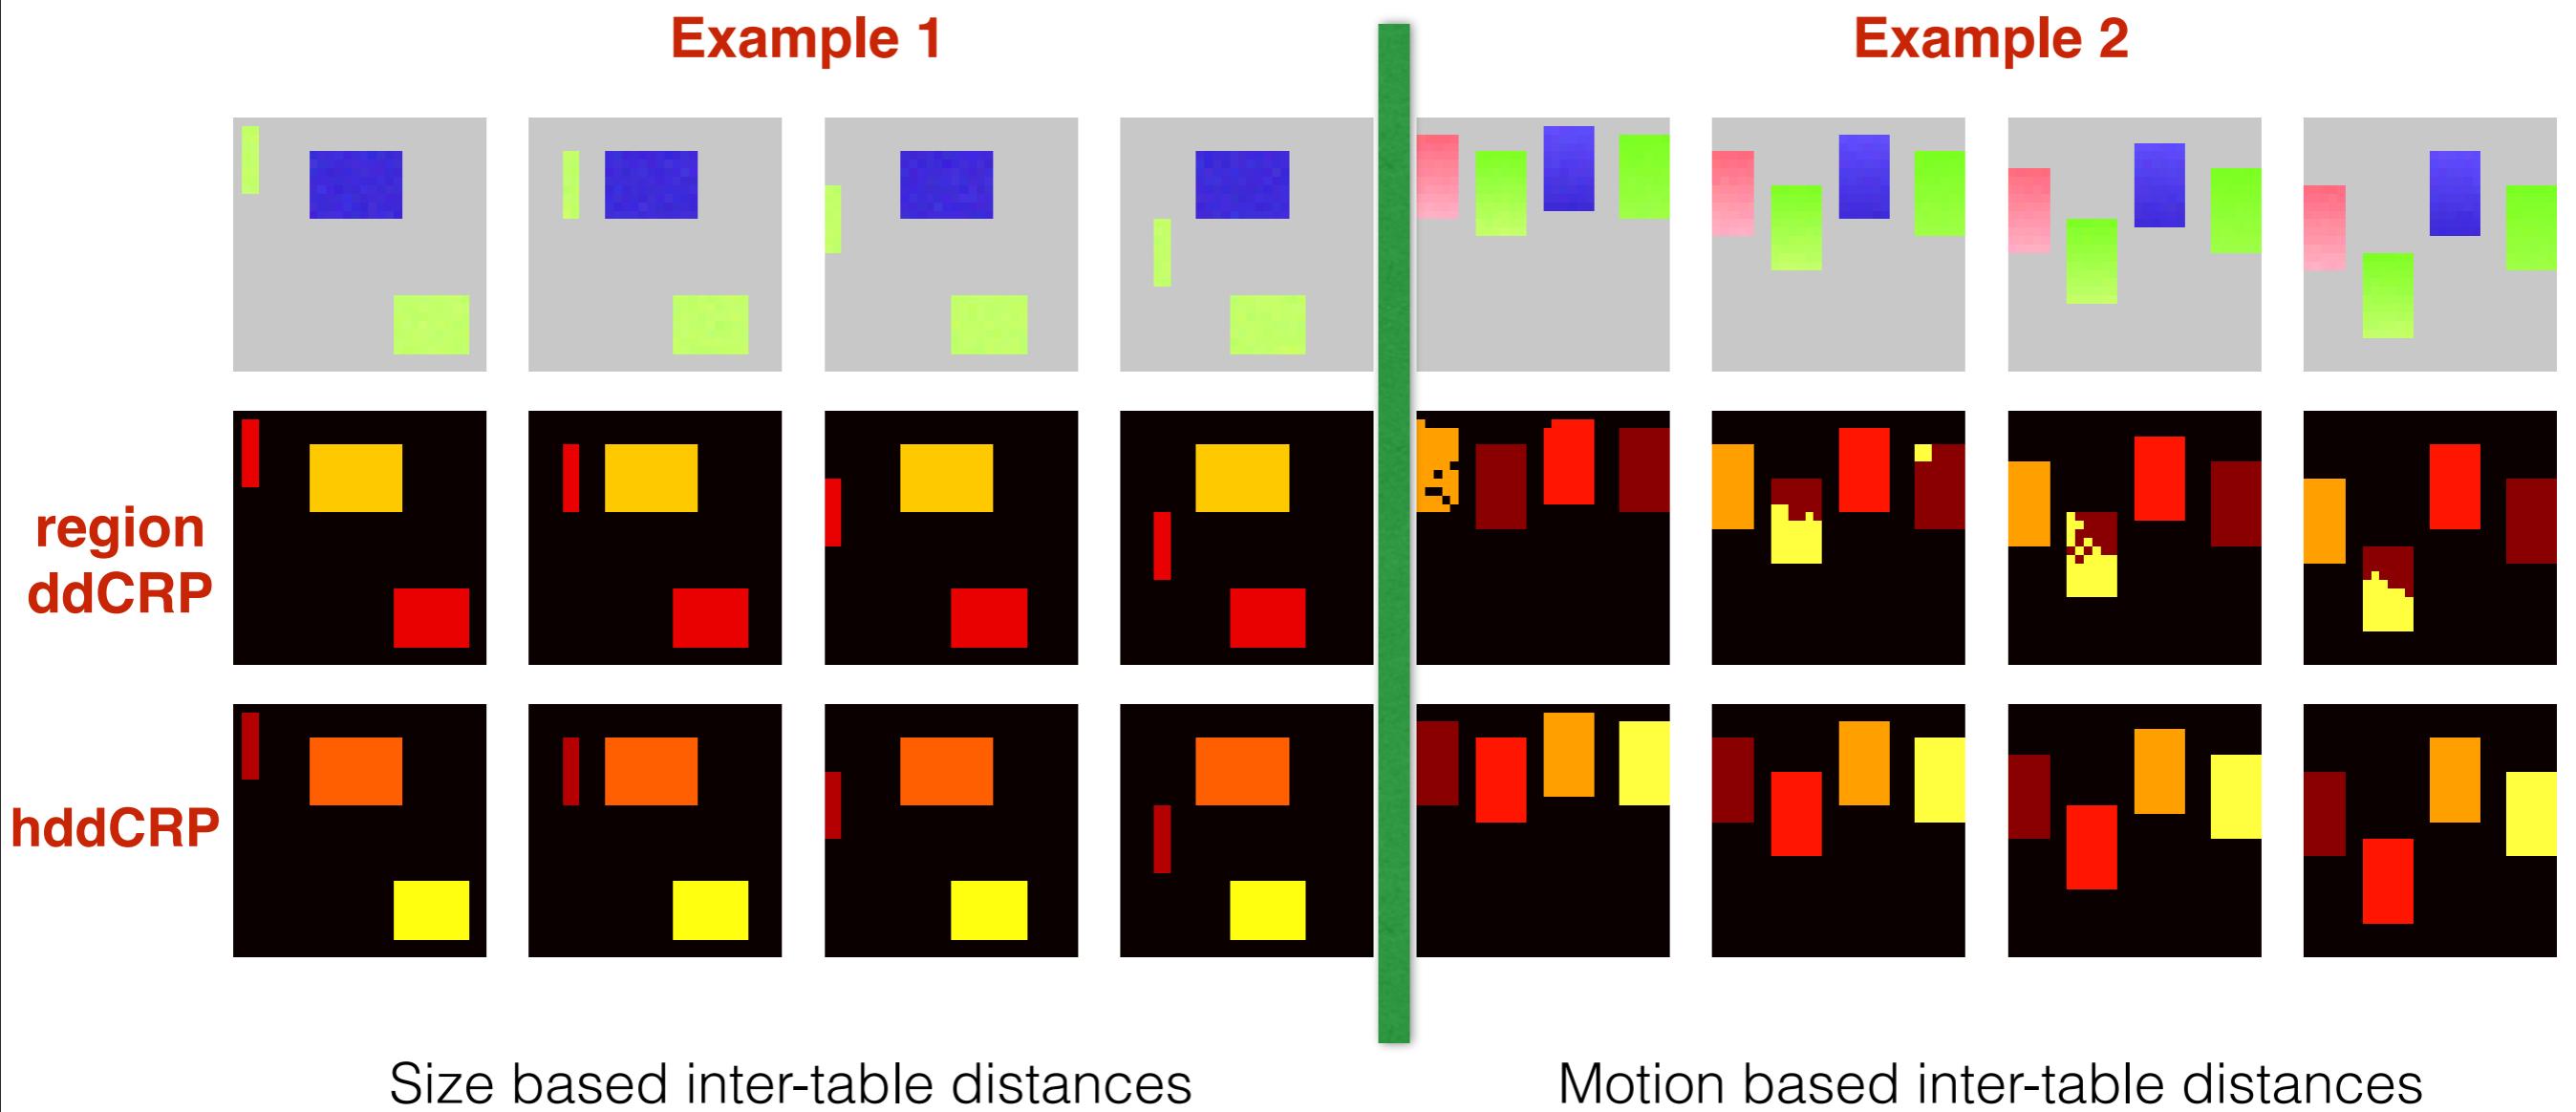
- Gaussian Process covariance determine the layer shape.

$$K(i, j)$$

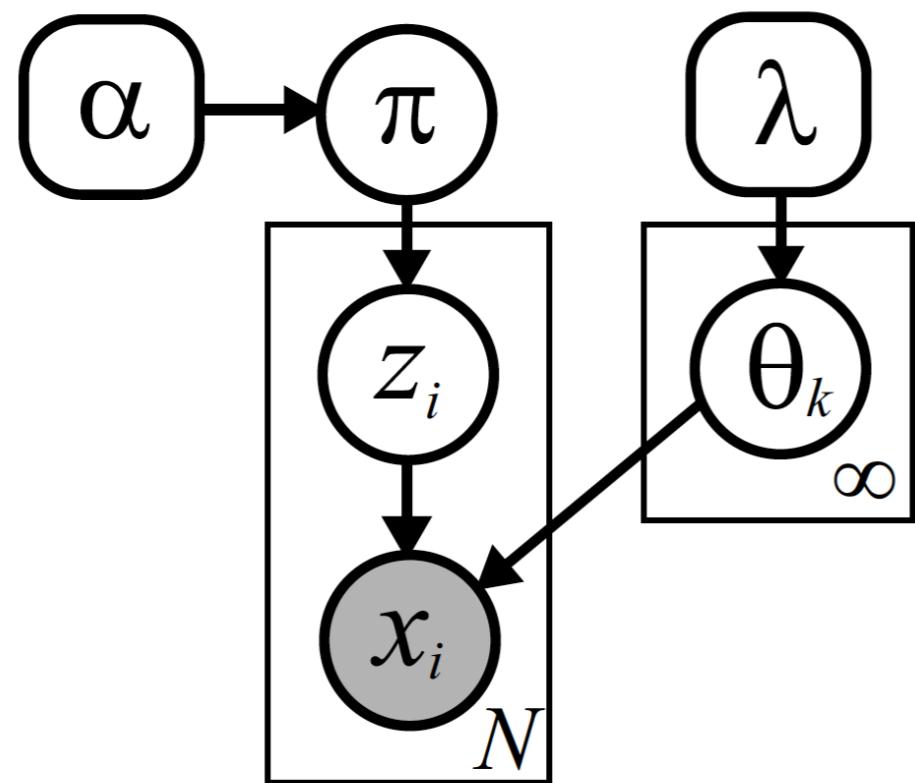


Probability that  
features at i and j are  
in the same layer

# Video Segmentation-Toy Data



# Bayesian Nonparametric Priors



$$G(\theta) = \sum_{k=1}^{\infty} \pi_k \delta_{\theta_k}(\theta)$$

$$\sum_{k=1}^{\infty} \pi_k = 1 \quad 0 \leq \pi_k \leq 1$$

$$\theta_k \sim H(\lambda)$$

# Pitman-Yor Process

**Power Law Behavior**

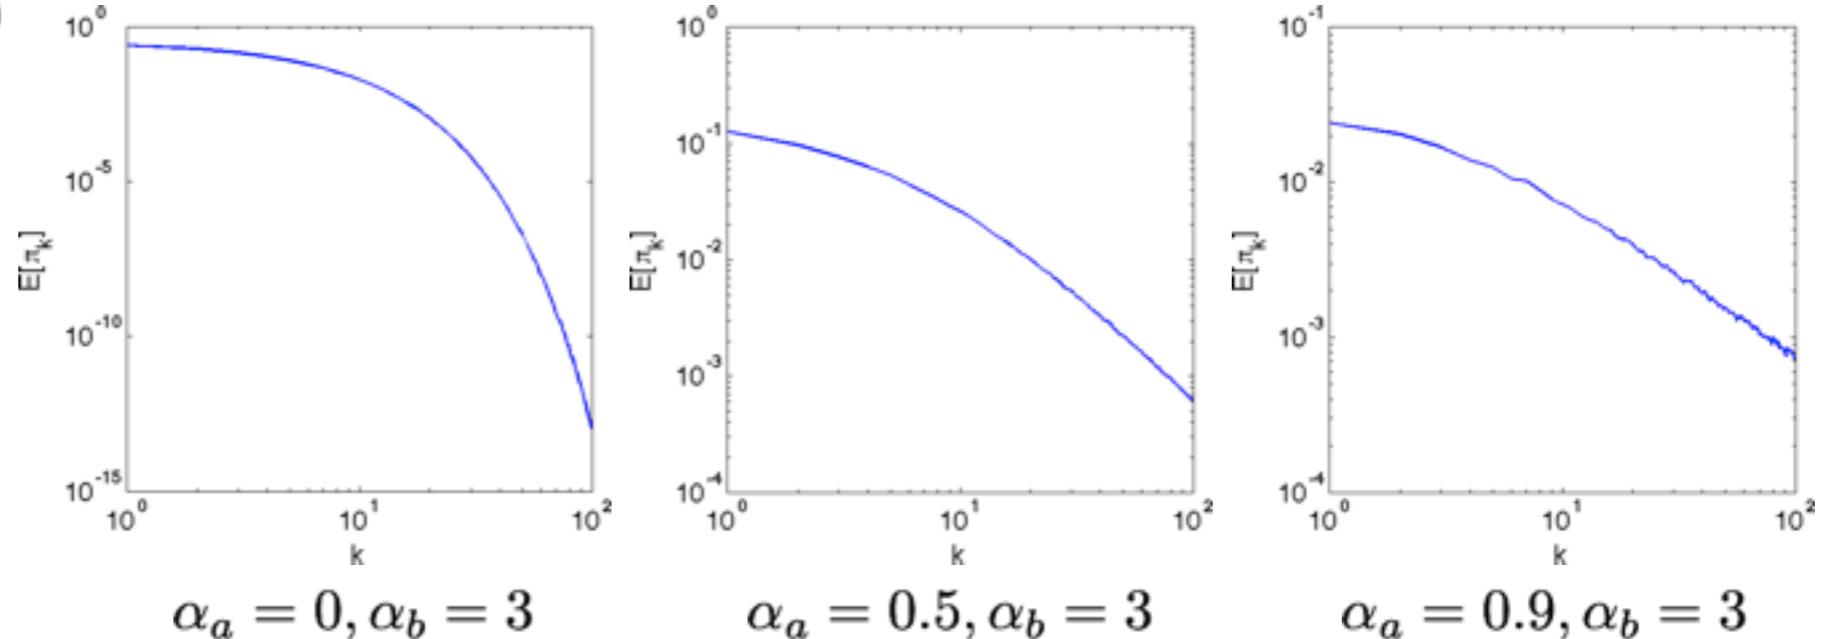
$$E[w_k] = \frac{1 - \alpha_a}{(1 + \alpha_b + (k - 1)\alpha_a)}$$

$$\pi_k = w_k \prod_{l=1}^{k-1} (1 - w_l)$$

$$w_k \sim \text{Beta}(a_k, b_k)$$

$$a_k = 1 - \alpha_a$$

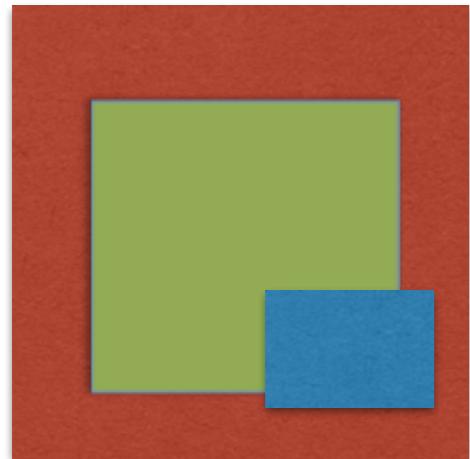
$$b_k = \alpha_b + k\alpha_a$$



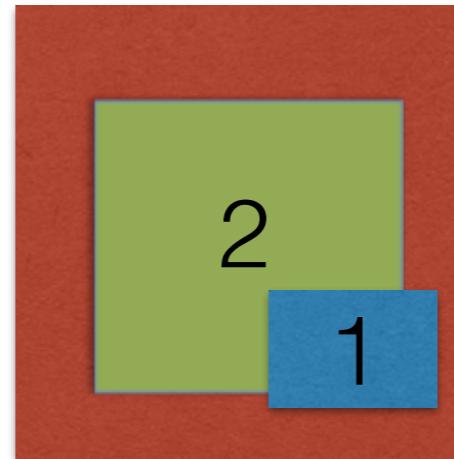
**Number of unique clusters in  $N$  observations:**  $O(\alpha_b N^{\alpha_a})$

**Expected size of sorted component  $k$ :**  $O(k^{-\frac{1}{\alpha_a}})$

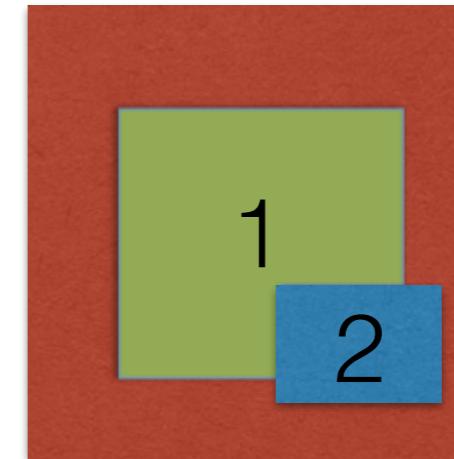
# Depth-ordered Layers



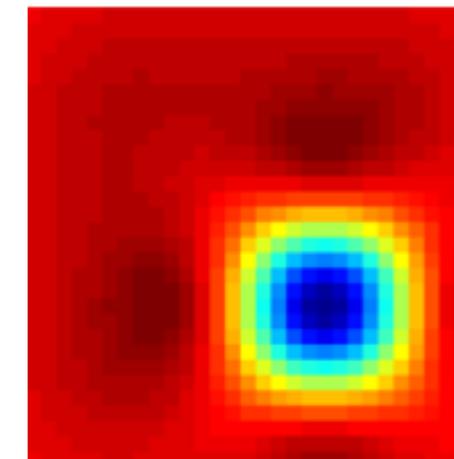
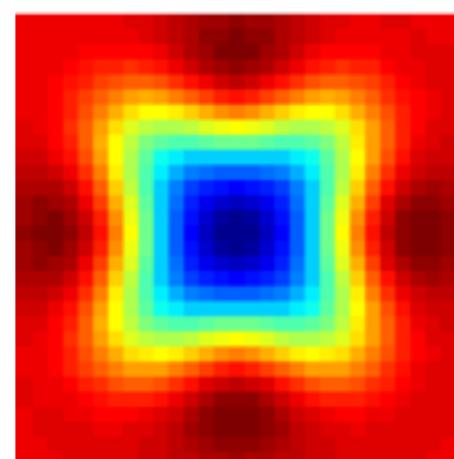
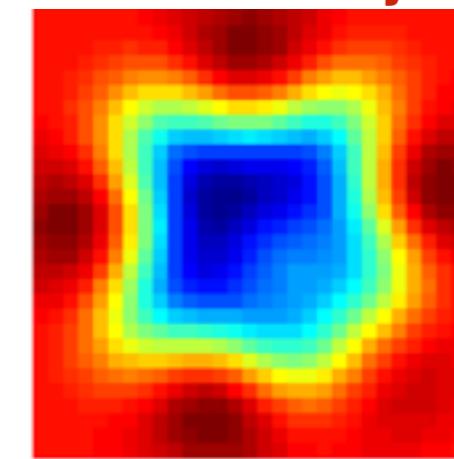
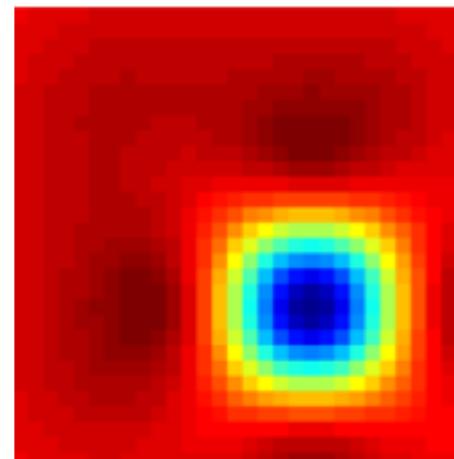
Layer 1



Layer 2



Lower Probability Partition

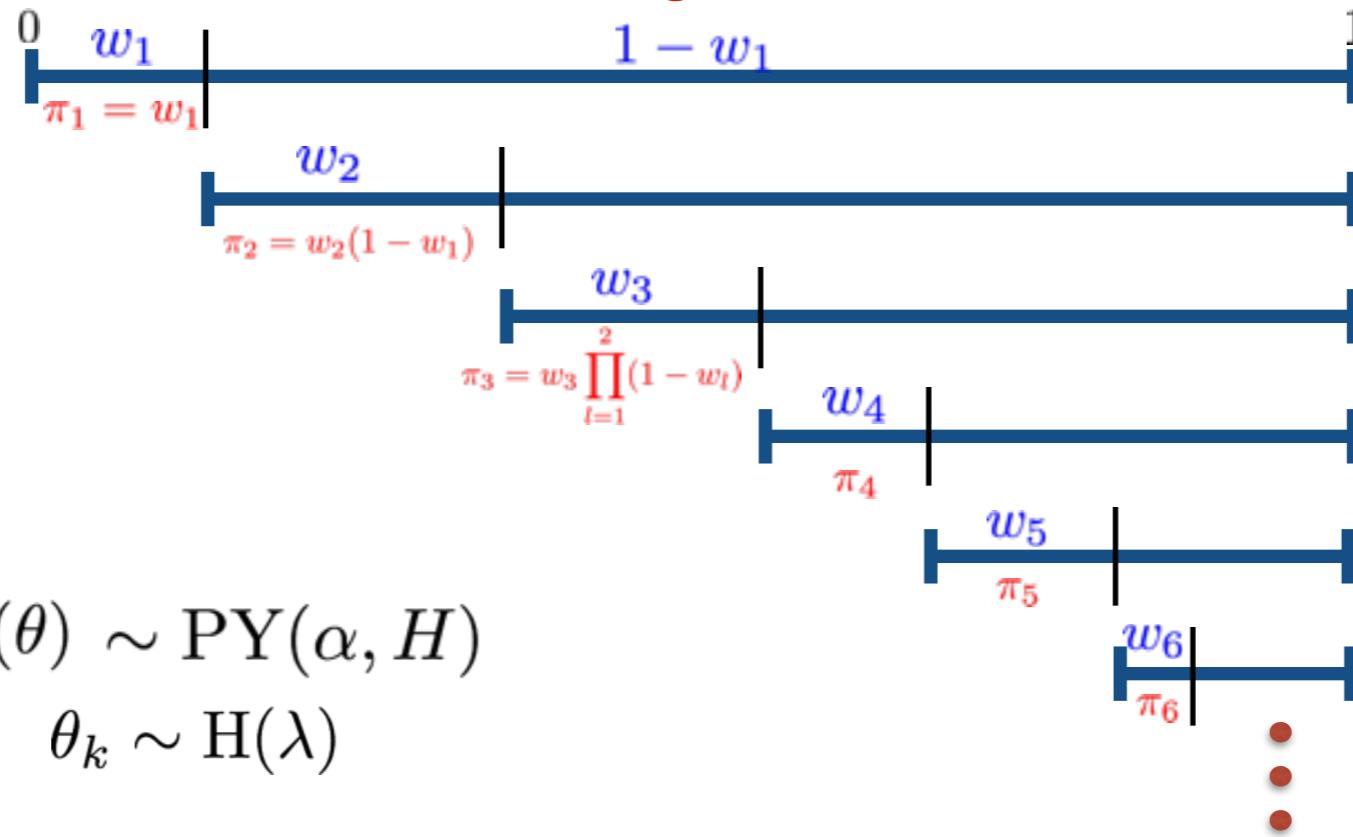


# Pitman-Yor Process

- The Pitman-Yor process defines a distribution on infinite discrete measures, or partitions

$$\pi_k = w_k \prod_{l=1}^{k-1} (1 - w_l) \quad w_k \sim \text{Beta}(1 - \alpha_a, \alpha_b + k\alpha_a)$$

## Stick Breaking Construction:



$$G(\theta) = \sum_{k=1}^{\infty} \pi_k \delta_{\theta_k}(\theta) \sim \text{PY}(\alpha, H)$$
$$\theta_k \sim H(\lambda)$$

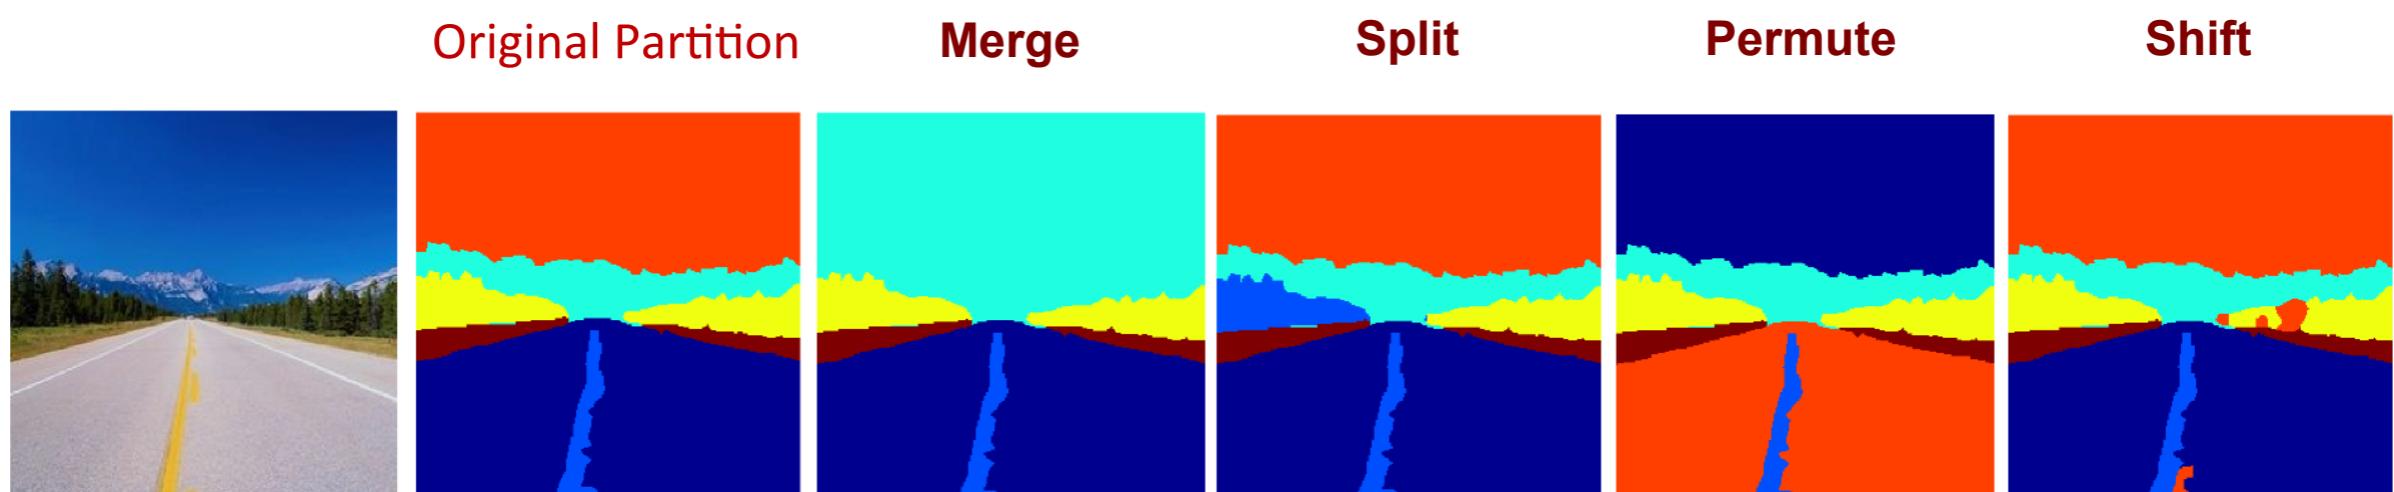
*Sethuraman, 1994  
Ishwaran and James,  
2001*

# Proposed Work

- Shared segmentation of image corpuses.
- Recover coarse scene geometry.
- Discover visual categories.
- Novel, scalable inference.

# Discrete Search Moves

- **Merge**: Combine a pair of regions into a single region
- **Split**: Break a single region into a pair of regions.
- **Shift**: Sequentially move super-pixels to most probable layers.
- **Permute**: Swap depth-order of two layers.



*Marginalization of continuous variables simplifies these moves...*

Ghosh & Sudderth, 2012

# Pitman-Yor Process

Power Law Behavior

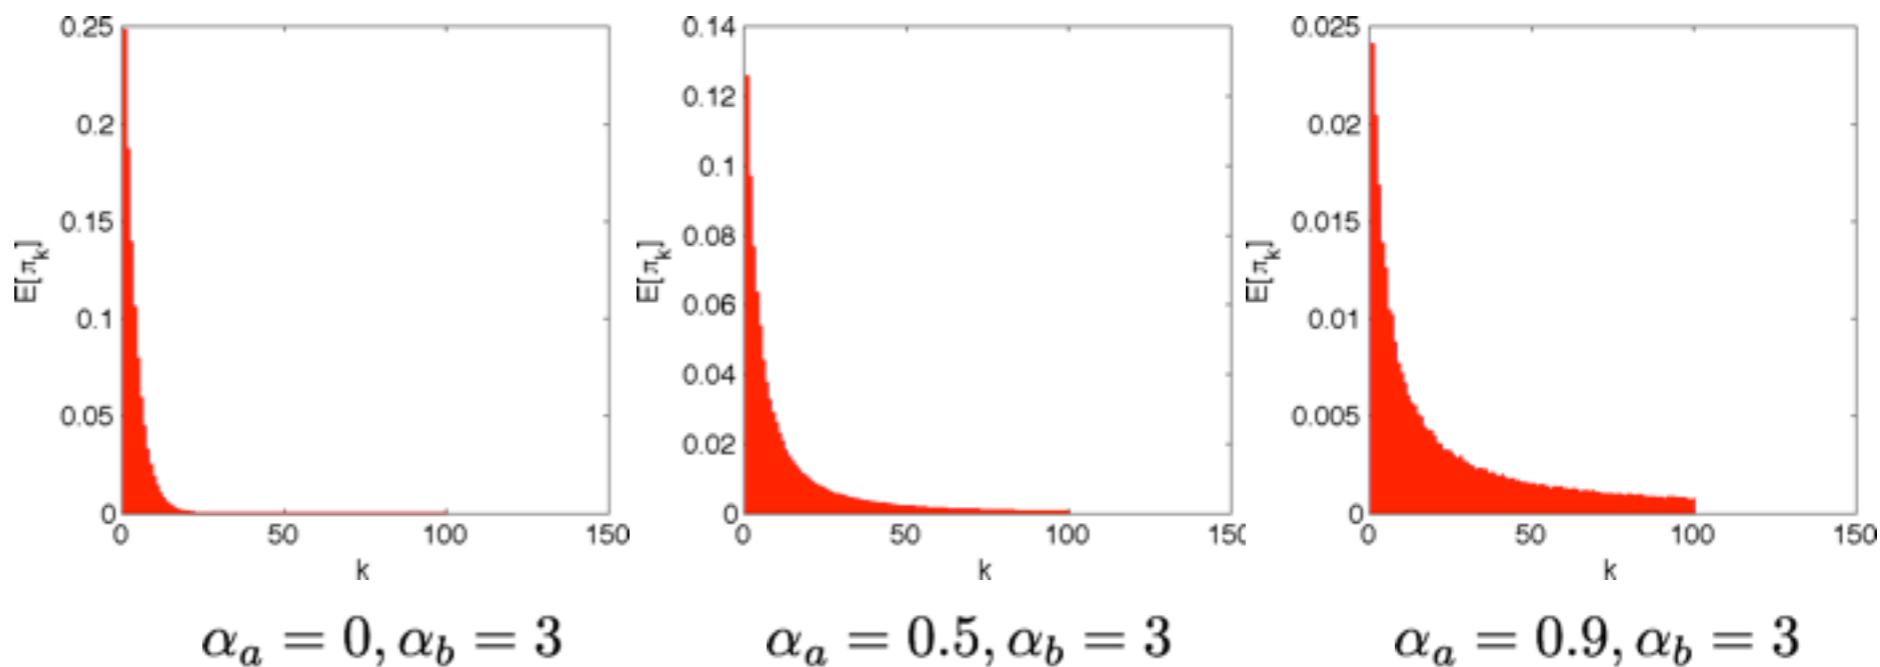
$$E[w_k] = \frac{1 - \alpha_a}{(1 + \alpha_b + (k - 1)\alpha_a)}$$

$$\pi_k = w_k \prod_{l=1}^{k-1} (1 - w_l)$$

$$w_k \sim \text{Beta}(a_k, b_k)$$

$$a_k = 1 - \alpha_a$$

$$b_k = \alpha_b + k\alpha_a$$



Dirichlet Process:

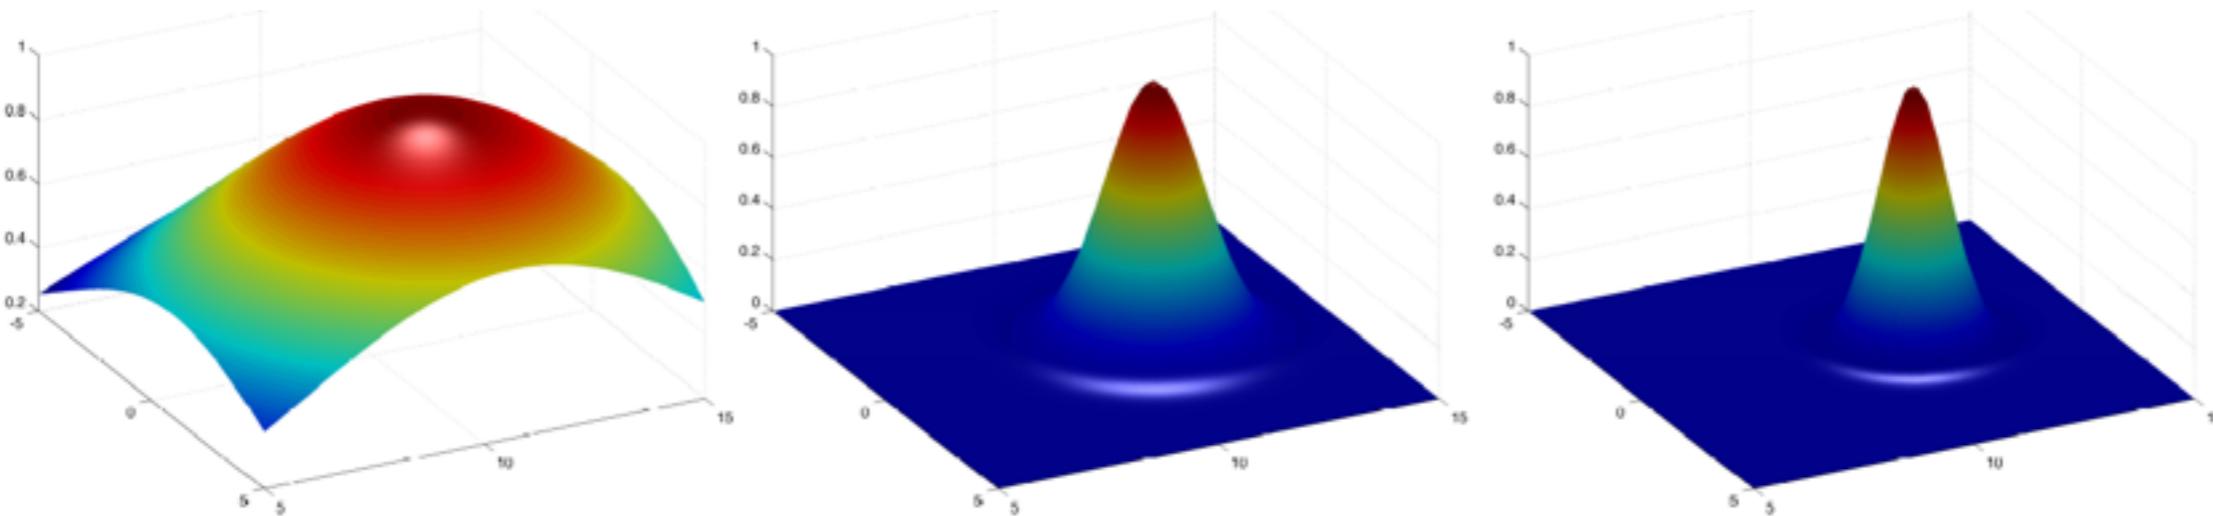
$$\alpha_a = 0$$

Number of unique clusters in  $N$  observations:  $O(\alpha_b N^{\alpha_a})$

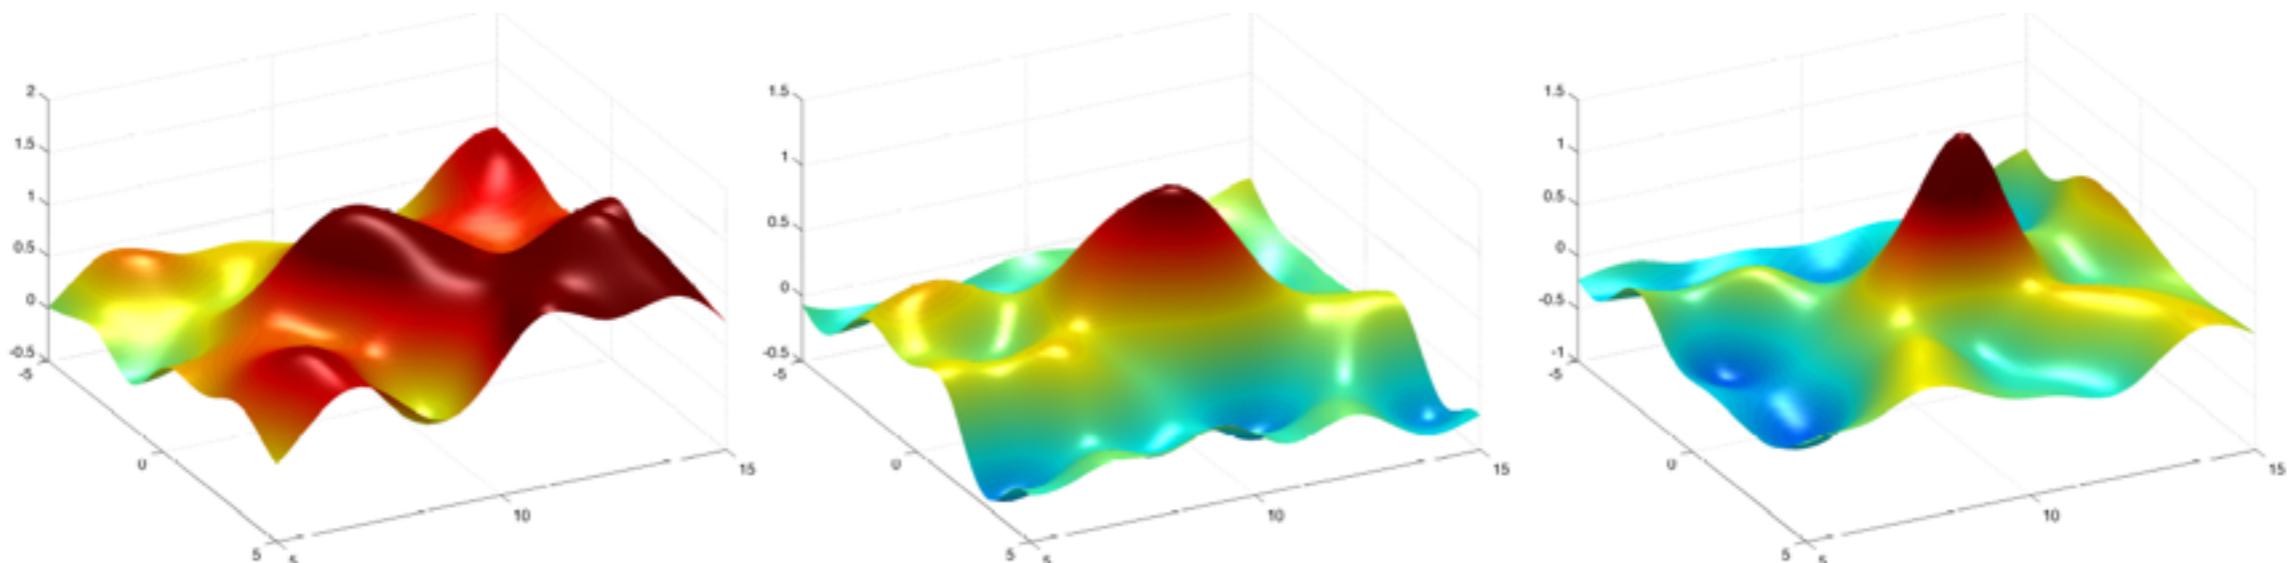
Expected size of sorted component  $k$ :  $O(k^{-\frac{1}{\alpha_a}})$

# Category specific shapes and sizes

- Each visual category is embellished with a mean function modeling its characteristic shape and size



- Sample an instance specific realization of a GP parametrized with the appropriate category mean and threshold.



# Probabilities to Correlations

- Two super-pixels belong to the same layer if they are both below the threshold. Thus w.p.:

$$p_- | \delta_k = \int_{-\infty}^{\delta_k} \int_{-\infty}^{\delta_k} \mathcal{N} \left( \begin{bmatrix} u_i \\ u_j \end{bmatrix} \mid \begin{bmatrix} 0 \\ 0 \end{bmatrix}, \begin{bmatrix} 1 & c \\ c & 1 \end{bmatrix} \right) du_i du_j$$

- Marginalizing over random threshold we have:

$$q_-^k(\alpha, \rho) = \int_{-\infty}^{\infty} \int_{-\infty}^{\delta_k} \int_{-\infty}^{\delta_k} \mathcal{N} \left( \begin{bmatrix} u_i \\ u_j \end{bmatrix} \mid \begin{bmatrix} 0 \\ 0 \end{bmatrix}, \begin{bmatrix} 1 & c \\ c & 1 \end{bmatrix} \right) p(\delta_k | \alpha) du_i du_j d\delta_k$$

$$q_+^k(\alpha, \rho) = \int_{-\infty}^{\infty} \int_{\delta_k}^{\infty} \int_{\delta_k}^{\infty} \mathcal{N} \left( \begin{bmatrix} u_i \\ u_j \end{bmatrix} \mid \begin{bmatrix} 0 \\ 0 \end{bmatrix}, \begin{bmatrix} 1 & c \\ c & 1 \end{bmatrix} \right) p(\delta_k | \alpha) du_i du_j d\delta_k$$

- Probability of belonging to the same layer then is:

$$\begin{aligned} q_{ij} &= q_-^1(\alpha, \rho) + q_-^2(\alpha, \rho) q_+^1(\alpha, \rho) + q_-^3(\alpha, \rho) q_+^1(\alpha, \rho) q_+^2(\alpha, \rho) + \dots \\ &\approx \sum_{k=1}^K q_-^k(\alpha, \rho) \prod_{l=1}^{K-1} q_+^l(\alpha, \rho) \end{aligned}$$